

CASSANDRA DATA MODELING STRATEGIES

Rohit Bhardwaj
Principal Cloud Engineer
rbhardwaj@kronos.com
Twitter: rbhardwaj1

AGENDA

- What is Cassandra?
- Cassandra data Model
- Cassandra Query Language
- TTLs,Tomstones and counters
- Designing multiple dimensions
- Cassandra Modeling strategies

ACID

- Atomicity: All operations are performed or none of them are. If one part of the transaction fails, then all fail.
- Consistency: The transaction must meet all rules defined by the system at all times; there are never any half-completed transactions.
- Isolation: Each transaction is independent unto itself.
- Durability: Once complete, the transaction cannot be undone.

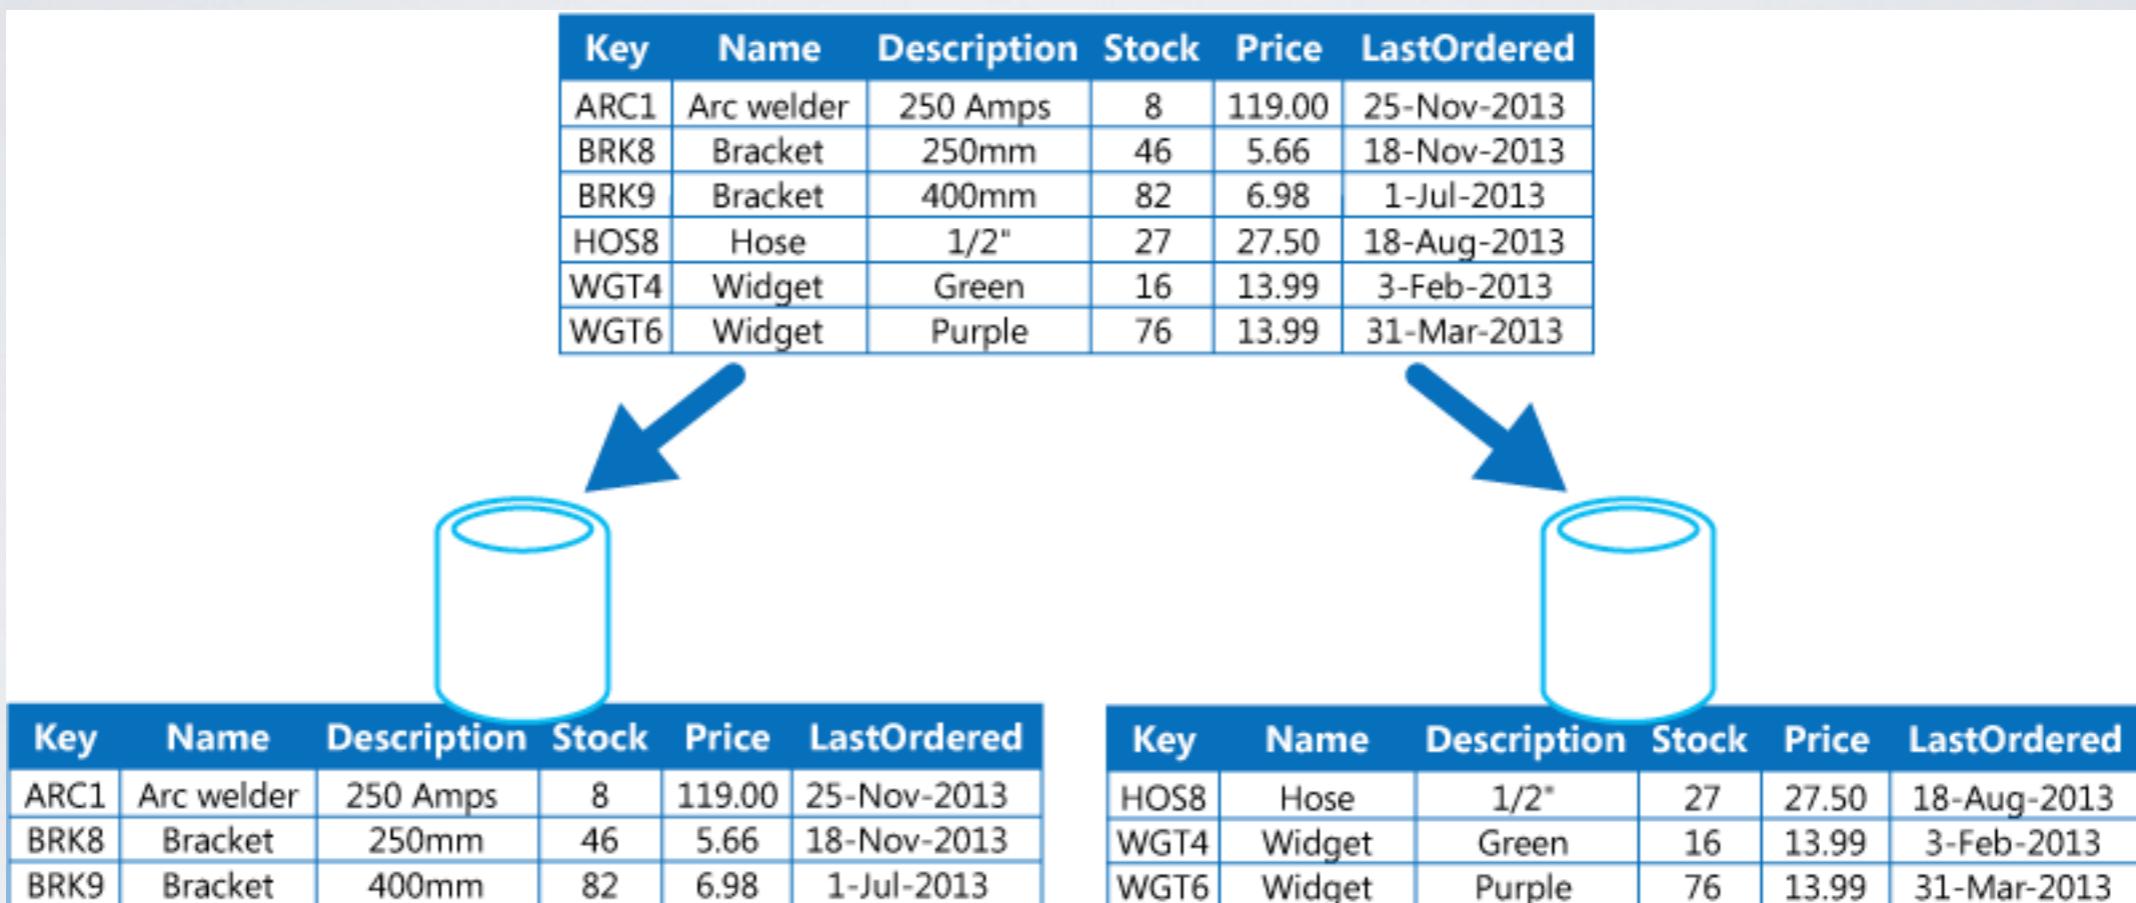
- Would ACID work with Bigdata ?

REPLICATION



```
SELECT COUNT(ArtifactID) FROM Document WHERE AccessControlListID_D IN (1,1000062) AND
(ArtifactID IN
(SELECT ArtifactID FROM Document WHERE AccessControlListID_D IN (1,1000062)
AND EXISTS
(SELECT CodeArtifactID FROM CodeArtifact WHERE AssociatedArtifactID = Document.ArtifactID
AND CodeArtifactID IN (17375543,17375544)
))
OR ArtifactID IN
(SELECT ArtifactID FROM Document WHERE AccessControlListID_D IN (1,1000062) AND
(EXISTS
(SELECT CodeArtifactID FROM CodeArtifact
WHERE AssociatedArtifactID = Document.ArtifactID AND CodeArtifactID IN (13002091,13002080,17018689,13002017)
)
AND NOT EXISTS
(SELECT CodeArtifactID FROM CodeArtifact WHERE AssociatedArtifactID = Document.ArtifactID
AND CodeArtifactID IN (16851390,17018659)
)
)
)
))
```

DB SHARDING



BASE

- Basically Available – the system guarantees some level of availability to the data even in regards to node failures. The data may be stale, but will still give and accept responses.
- Soft State – the data is in a constant state of flux; so, while a response may be given, the freshness or consistency of the data is not guaranteed to be the most current.
- Eventual Consistency – the data will eventually be consistent through all nodes and in all databases, but not every transaction at every moment. It will reach some guaranteed state eventually.

ACID

- Strong Consistency
- Isolation
- Focus on “commit”
- Nested transactions
- Availability?
- Conservative (pessimistic)
- Difficult evolution (e.g. schema)

BASE

- Weak Consistency – stale data OK
- Availability first
- Best effort
- Approximate answers OK
- Aggressive (optimistic)
- Simpler!
- Faster
- Easier evolution

AVAILABILITY LEVEL¹¹ **ESTIMATED DOWNTIME PER YEAR**

99.999%

5 minutes

99.99%

52 minutes

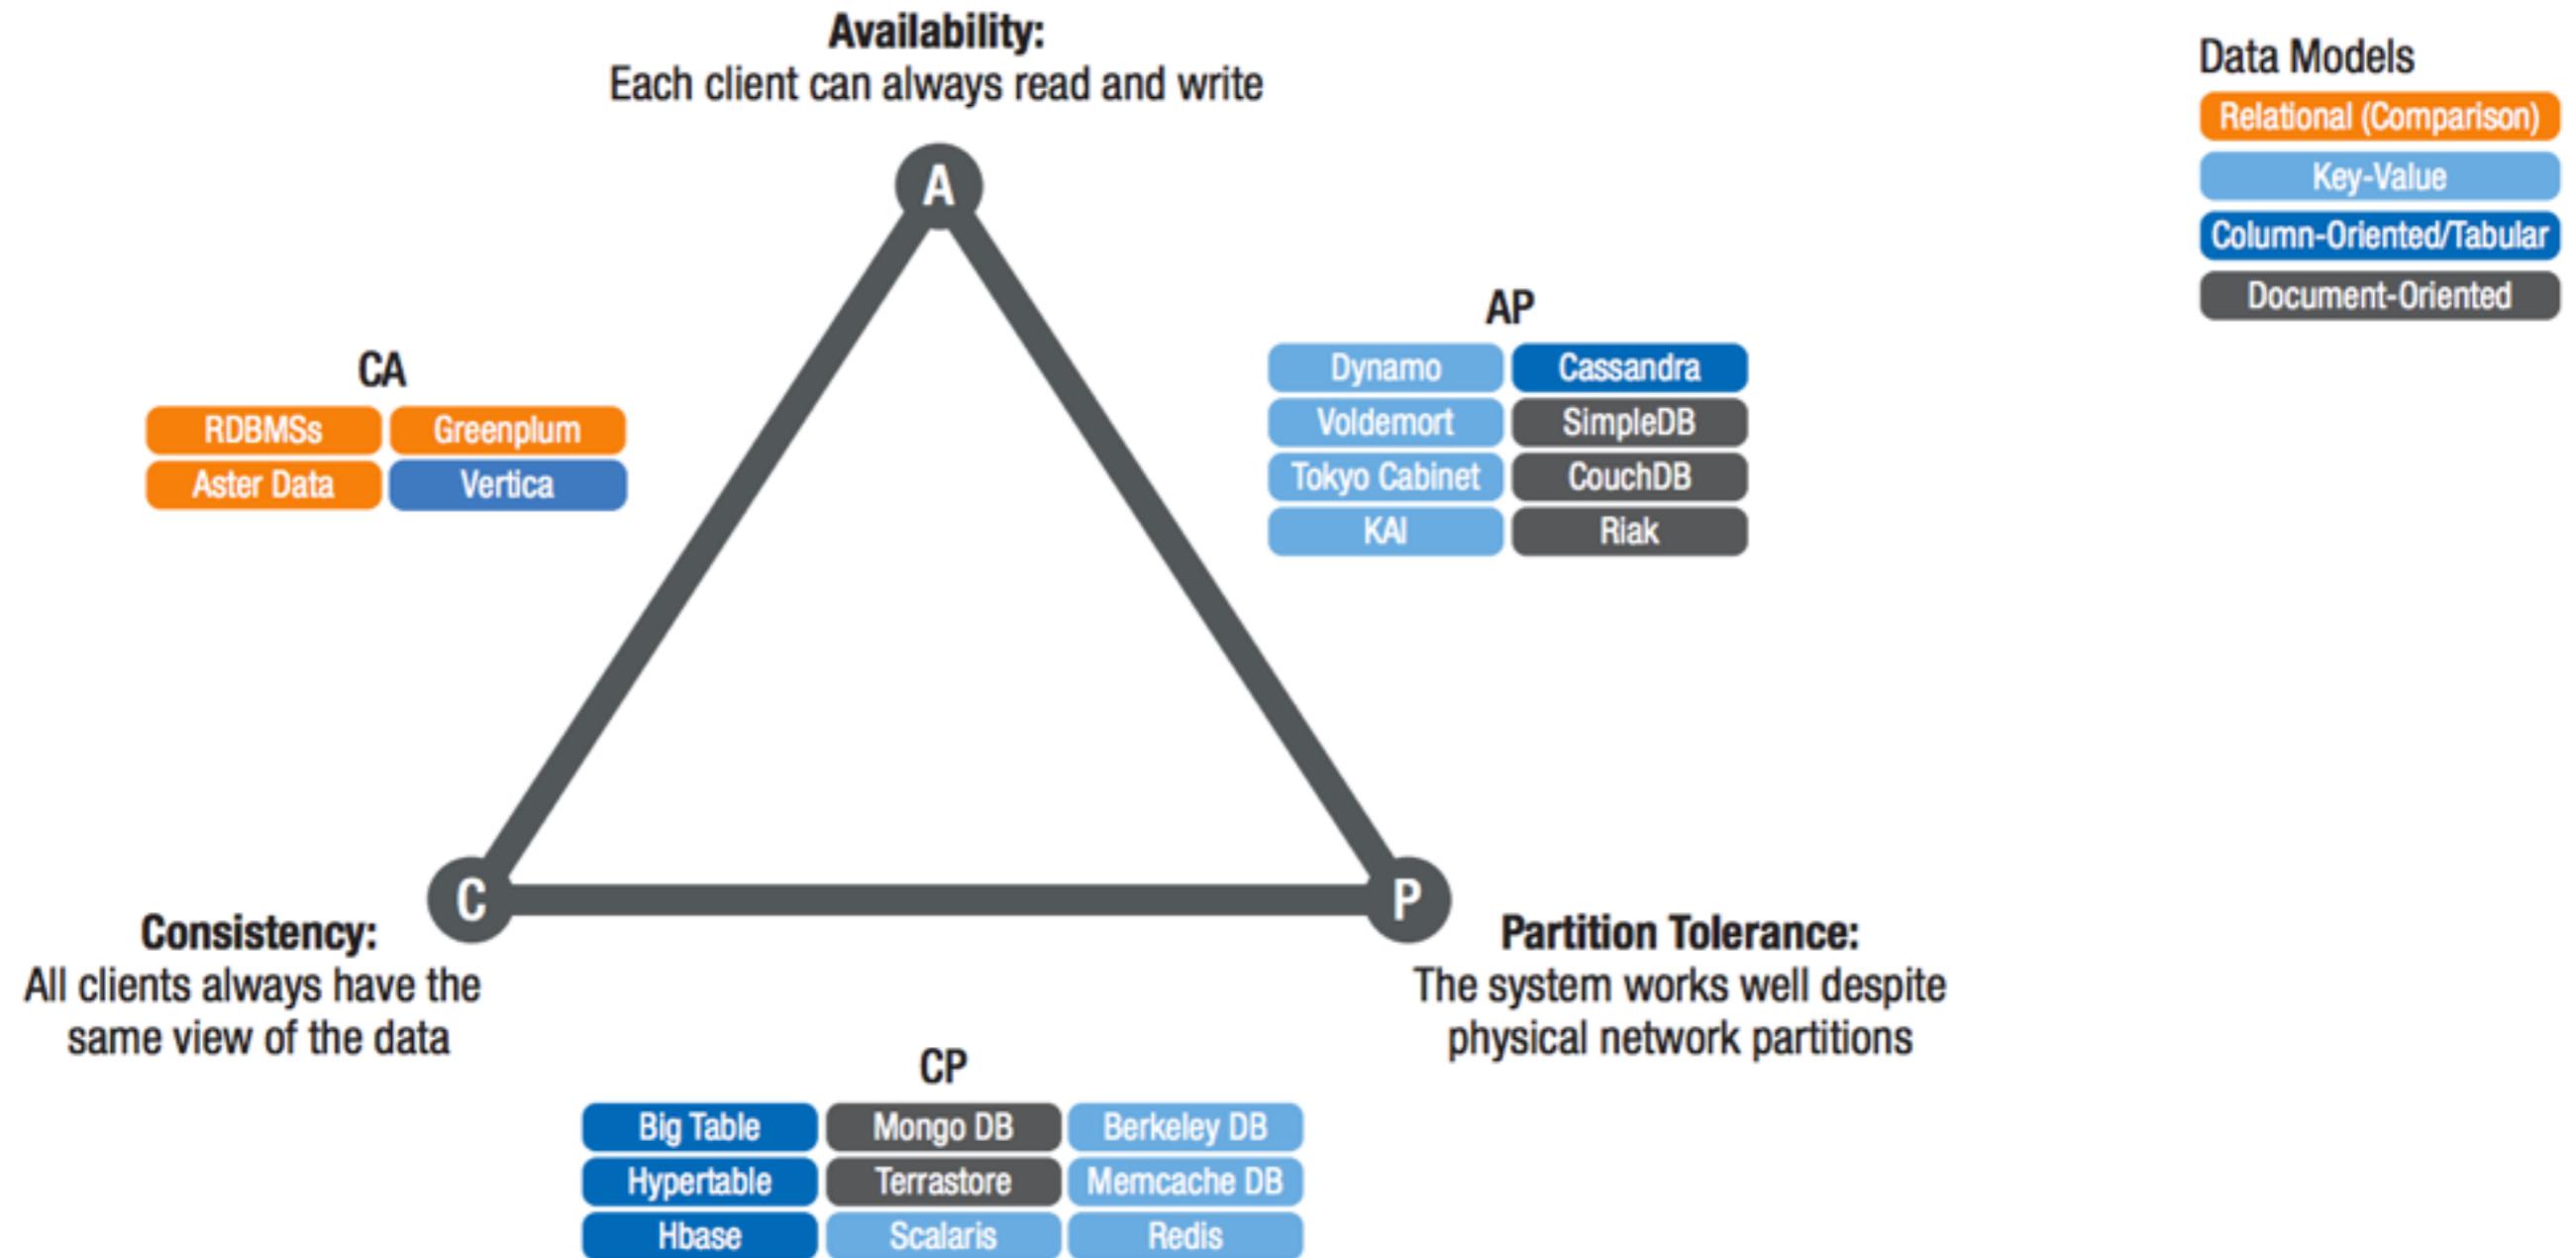
99.9%

8.5 hours

99%

3.5 days or about 87 hours

CAP THEOREM



BASE (Basically Available, Soft-State, Eventual Consistency) data store

SELECTION	CHARACTERISTICS	EXAMPLES
C + A (No P)	2-phase commits Cache validation protocols	Single-site databases Cluster databases LDAP xFS file system
C + P (no A)	Pessimistic locking minority partitions unavailable	Distributed databases Distributed locking Majority protocols
A + P (no C)	Expirations/leases Conflict resolution Optimistic	Coda Web caching DNS



Cassandra

A highly scalable, eventually consistent, distributed, structured key-value store.

Bigtable



Dynamo



Cassandra

Cassandra – A structured storage system on a P2P Network

By [Avinash Lakshman](#) on Monday, August 25, 2008 at 4:31pm 

When I joined Facebook I was eagerly looking forward to a new challenge. Fortunately, Facebook cannot be accused of a lack of challenging assignments. Prashant Malik, a colleague in Facebook from the Search team, was thinking about how to solve the Inbox Search problem. This challenge is about storing reverse indices of Facebook messages that Facebook users send and receive while communicating with their friends on the Facebook network. The amount of data to be stored, the rate of growth of the data and the requirement to serve it within strict SLAs made it very apparent that a new storage solution was absolutely essential. The solution needed to scale incrementally and in a cost effective fashion. Traditional data storage solutions just wouldn't fit the bill. The aim was to design a solution that not only solved the Inbox Search problem but also provided a system as a storage infrastructure for many problems of the same nature. Hence was born Cassandra. To keep up with Facebook tradition, Prashant and I started the implementation of Cassandra about a year ago in one of our Hackthons.



Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati,
Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall
and Werner Vogels

Amazon.com

ABSTRACT

Reliability at massive scale is one of the biggest challenges we face at Amazon.com, one of the largest e-commerce operations in the world; even the slightest outage has significant financial consequences and impacts customer trust. The Amazon.com platform, which provides services for many web sites worldwide, is implemented on top of an infrastructure of tens of thousands of servers and network components located in many datacenters around the world. At this scale, small and large components fail continuously and the way persistent state is managed in the face of these failures drives the reliability and scalability of the software systems.

One of the lessons our organization has learned from operating Amazon's platform is that the reliability and scalability of a system is dependent on how its application state is managed. Amazon uses a highly decentralized, loosely coupled, service oriented architecture consisting of hundreds of services. In this environment there is a particular need for storage technologies that are always available. For example, customers should be able to view and add items to their shopping cart even if disks are failing, network routes are flapping, or data centers are being destroyed by tornados. Therefore, the service responsible for managing shopping carts requires that it can always write to and read from its data store, and that its data needs to be available across multiple data centers.

Bigtable: A Distributed Storage System for Structured Data

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach
Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber

{fay,jeff,sanjay,wilsonh,kerr,m3b,tushar,fikes,gruber}@google.com

Google, Inc.



Cassandra

[Home](#) [Download](#) [Getting Started](#) [Planet Cassandra](#) [Contribute](#)

[Welcome](#) [Video](#) [Slides](#)

Welcome to Apache Cassandra™

The Apache Cassandra database is the right choice when you need scalability and high availability without compromising performance. [Linear scalability](#) and proven fault-tolerance on commodity hardware or cloud infrastructure make it the perfect platform for mission-critical data. Cassandra's support for replicating across multiple datacenters is best-in-class, providing lower latency for your users and the peace of mind of knowing that you can survive regional outages.

Cassandra's data model offers the convenience of [column indexes](#) with the performance of log-structured updates, strong support for [denormalization](#) and [materialized views](#), and powerful built-in caching.

Download

[Tick-Tock release 3.4 \(Changes\)](#)

[3.0.x release 3.0.4 \(Changes\)](#)

[2.2.x release 2.2.5 \(Changes\)](#)



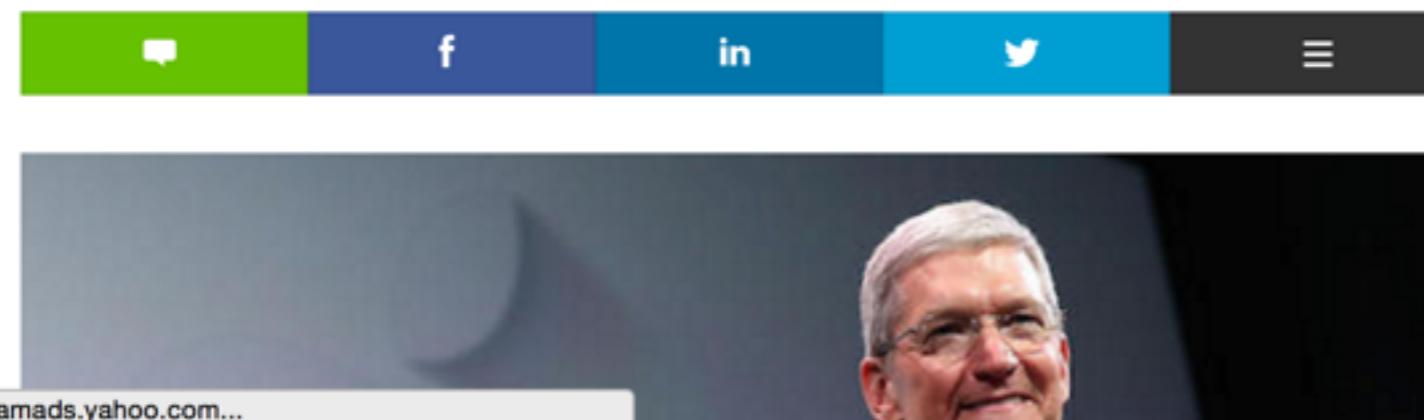
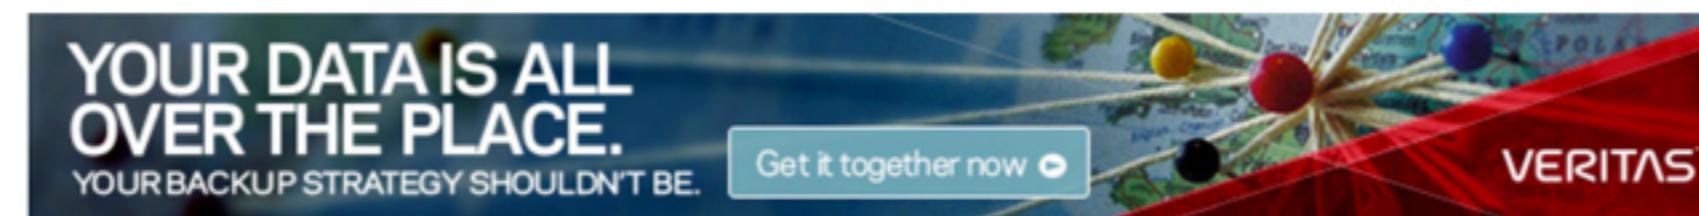
[Download options](#)



Apple's secret NoSQL sauce includes a hefty dose of Cassandra

If you want to scale like Apple, you might need to consider Cassandra. Matt Asay explains.

By Matt Asay | September 16, 2015, 6:49 AM PST

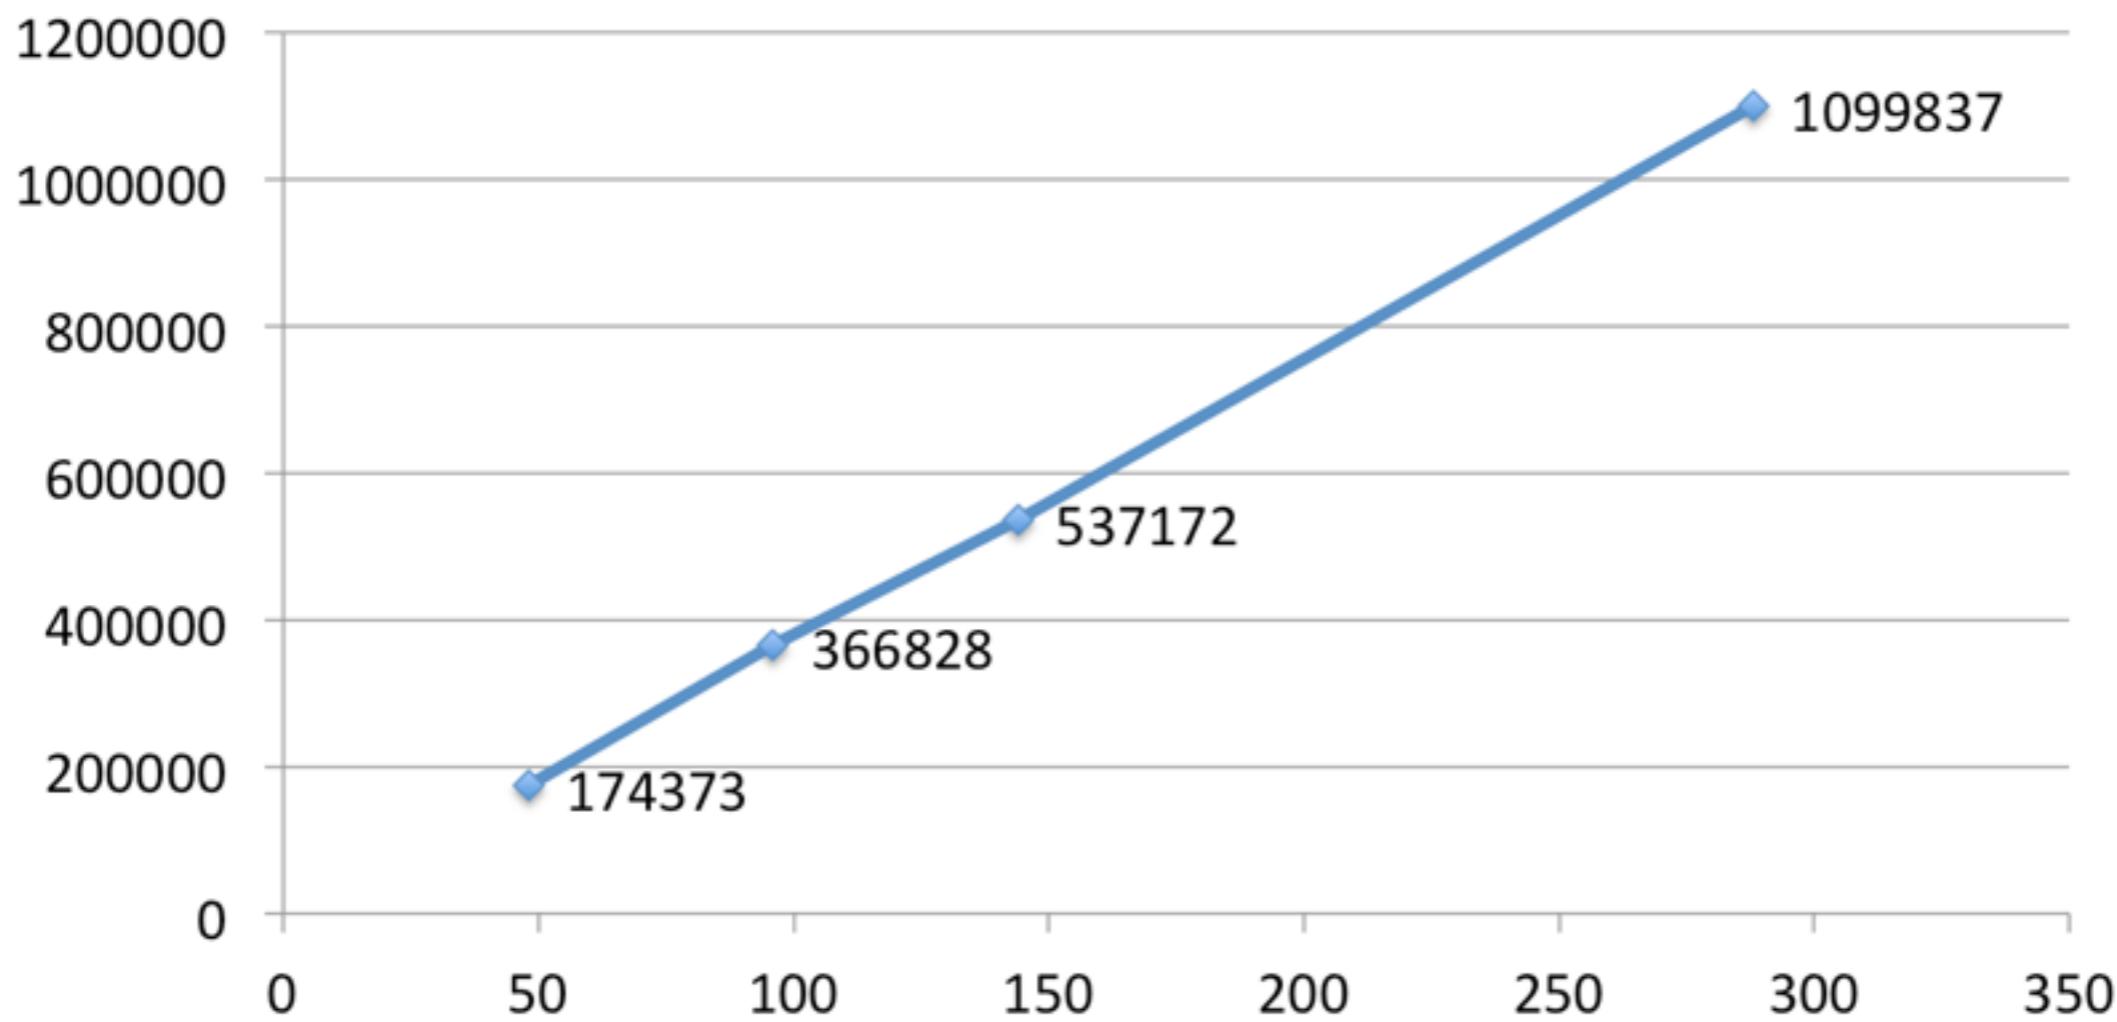


for syndication.streamads.yahoo.com...

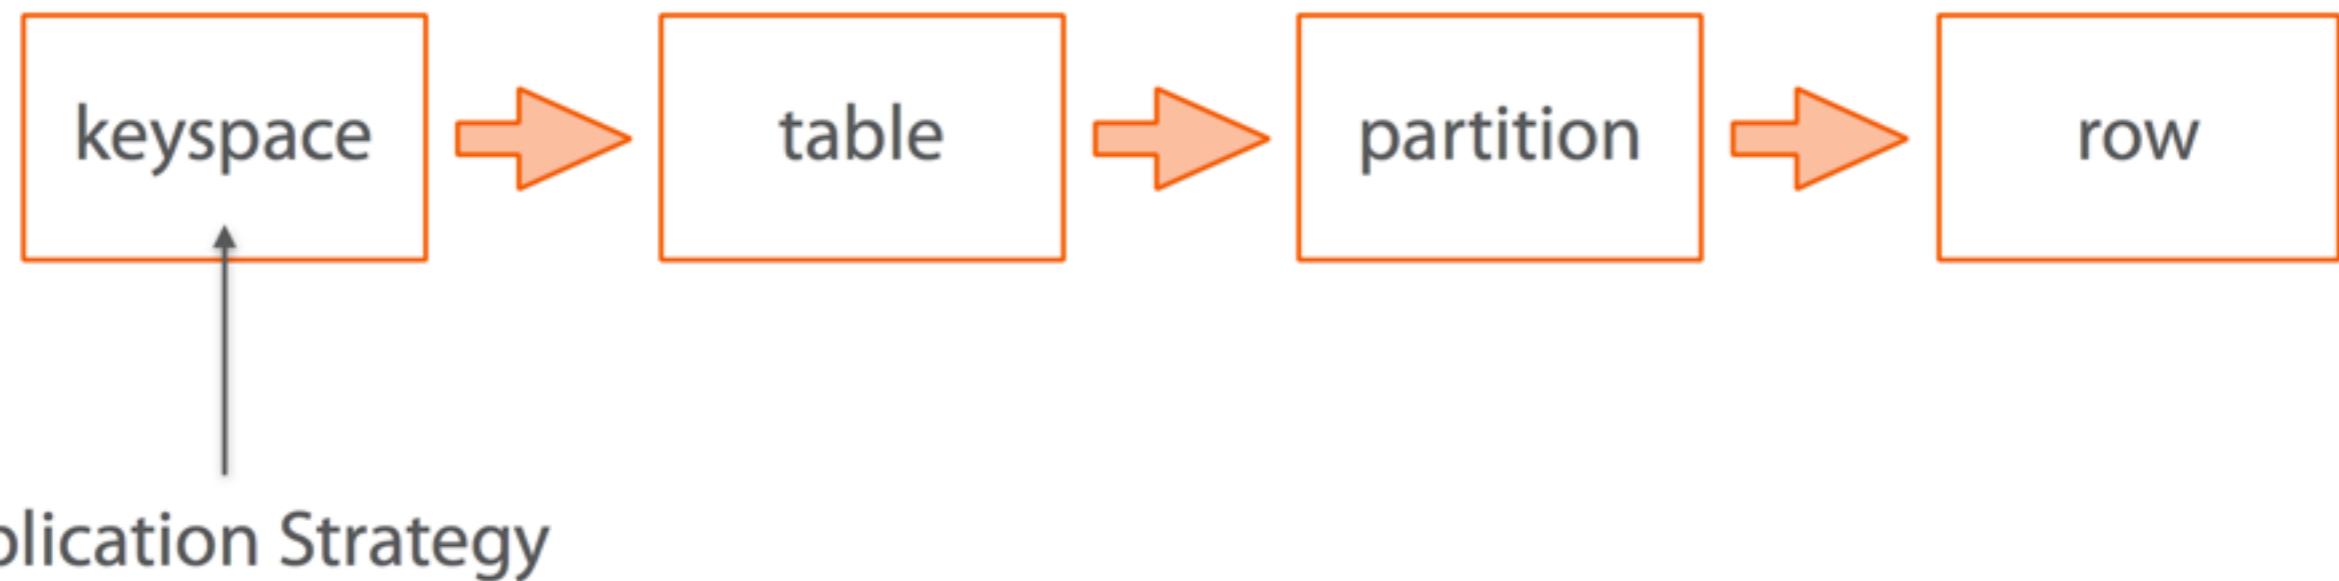


Scale-Up Linearity

Client Writes/s by node count – Replication Factor = 3

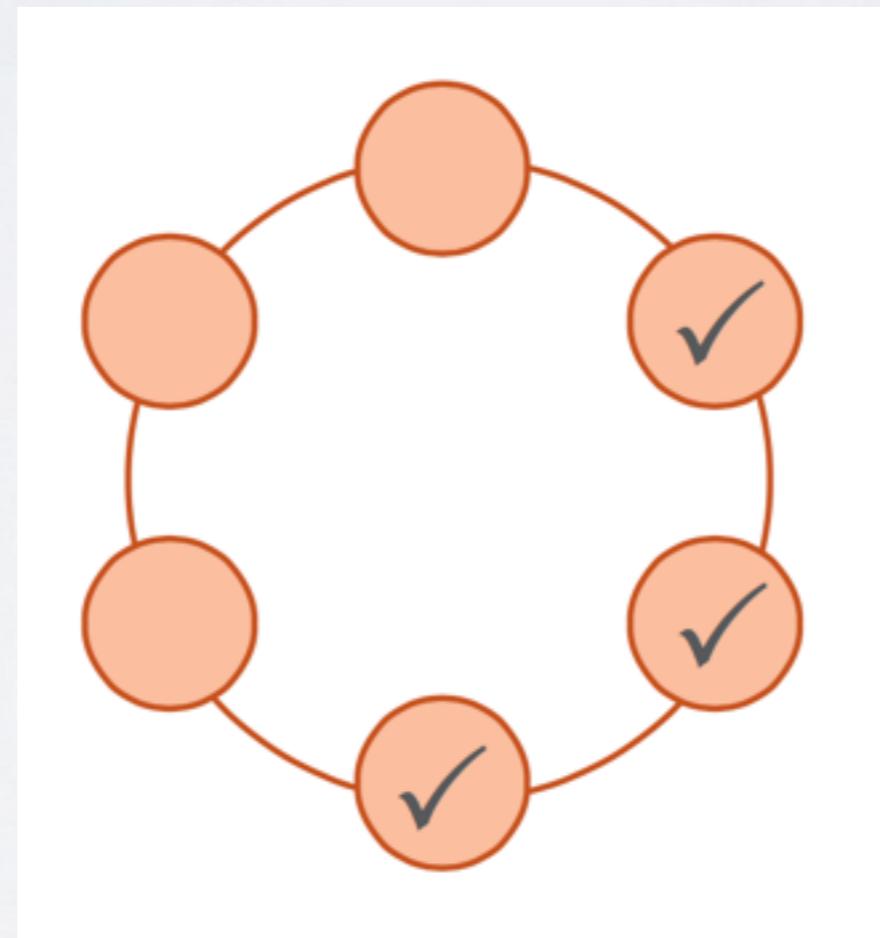


CASSANDRA TERMINOLOGY



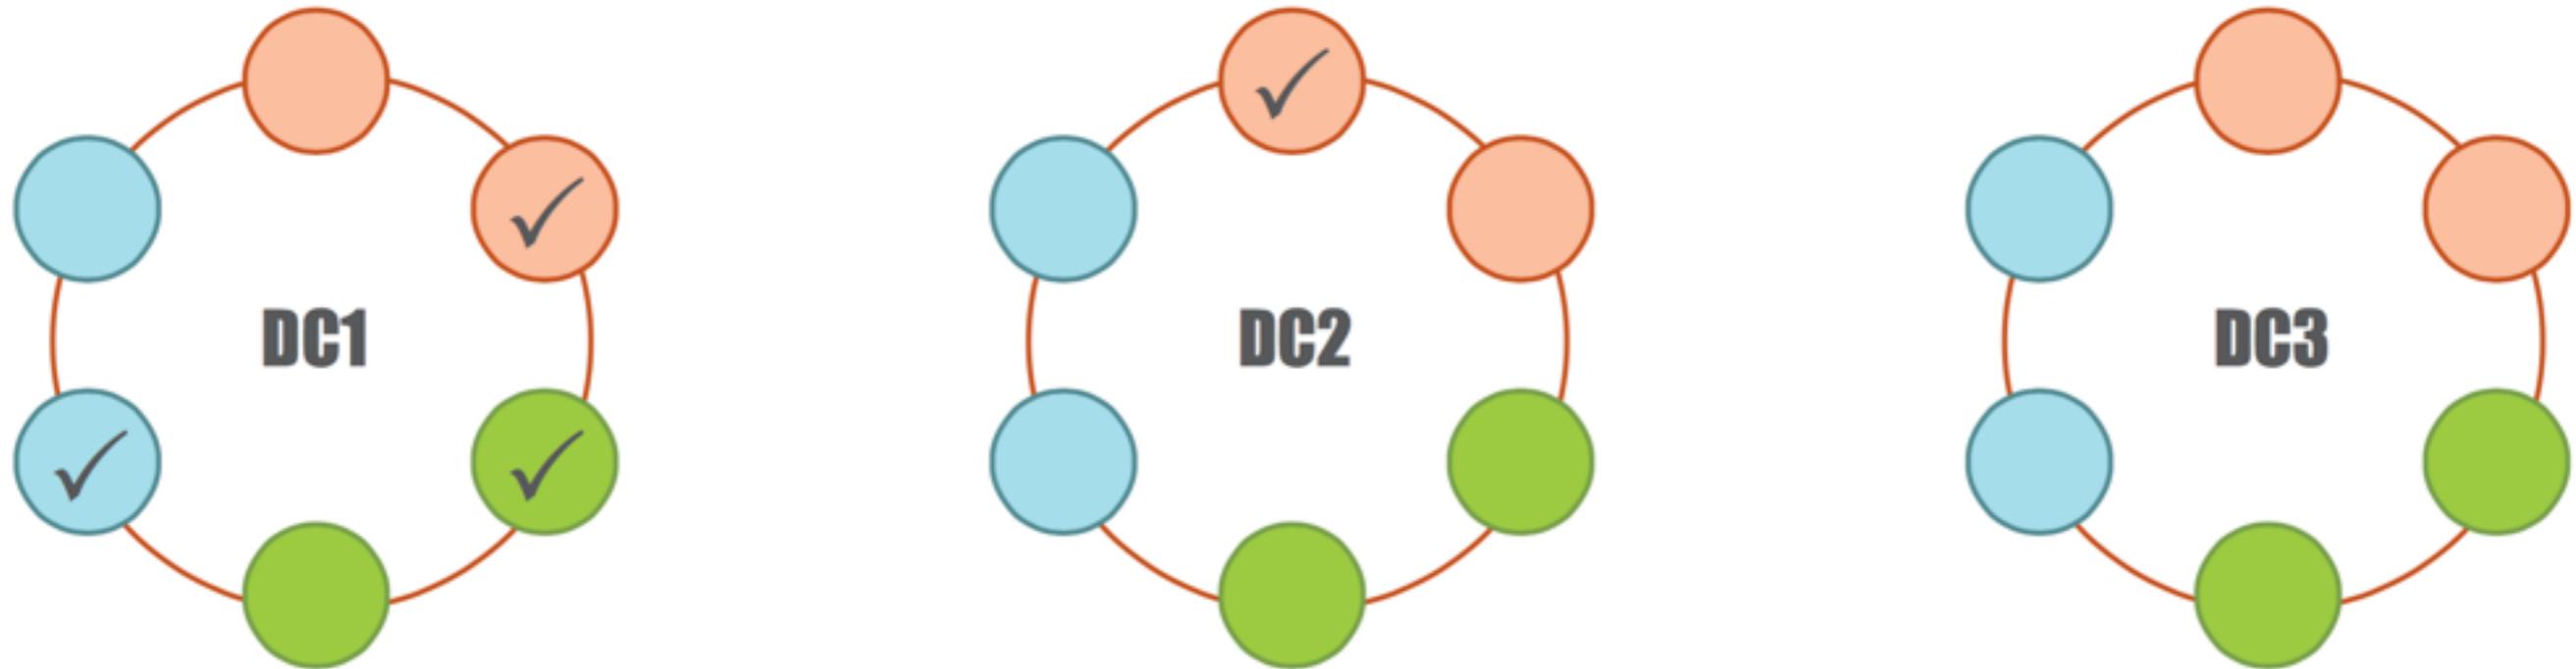
SIMPLE STRATEGY

- CREATE KEYSPACE nfjs WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 3 };

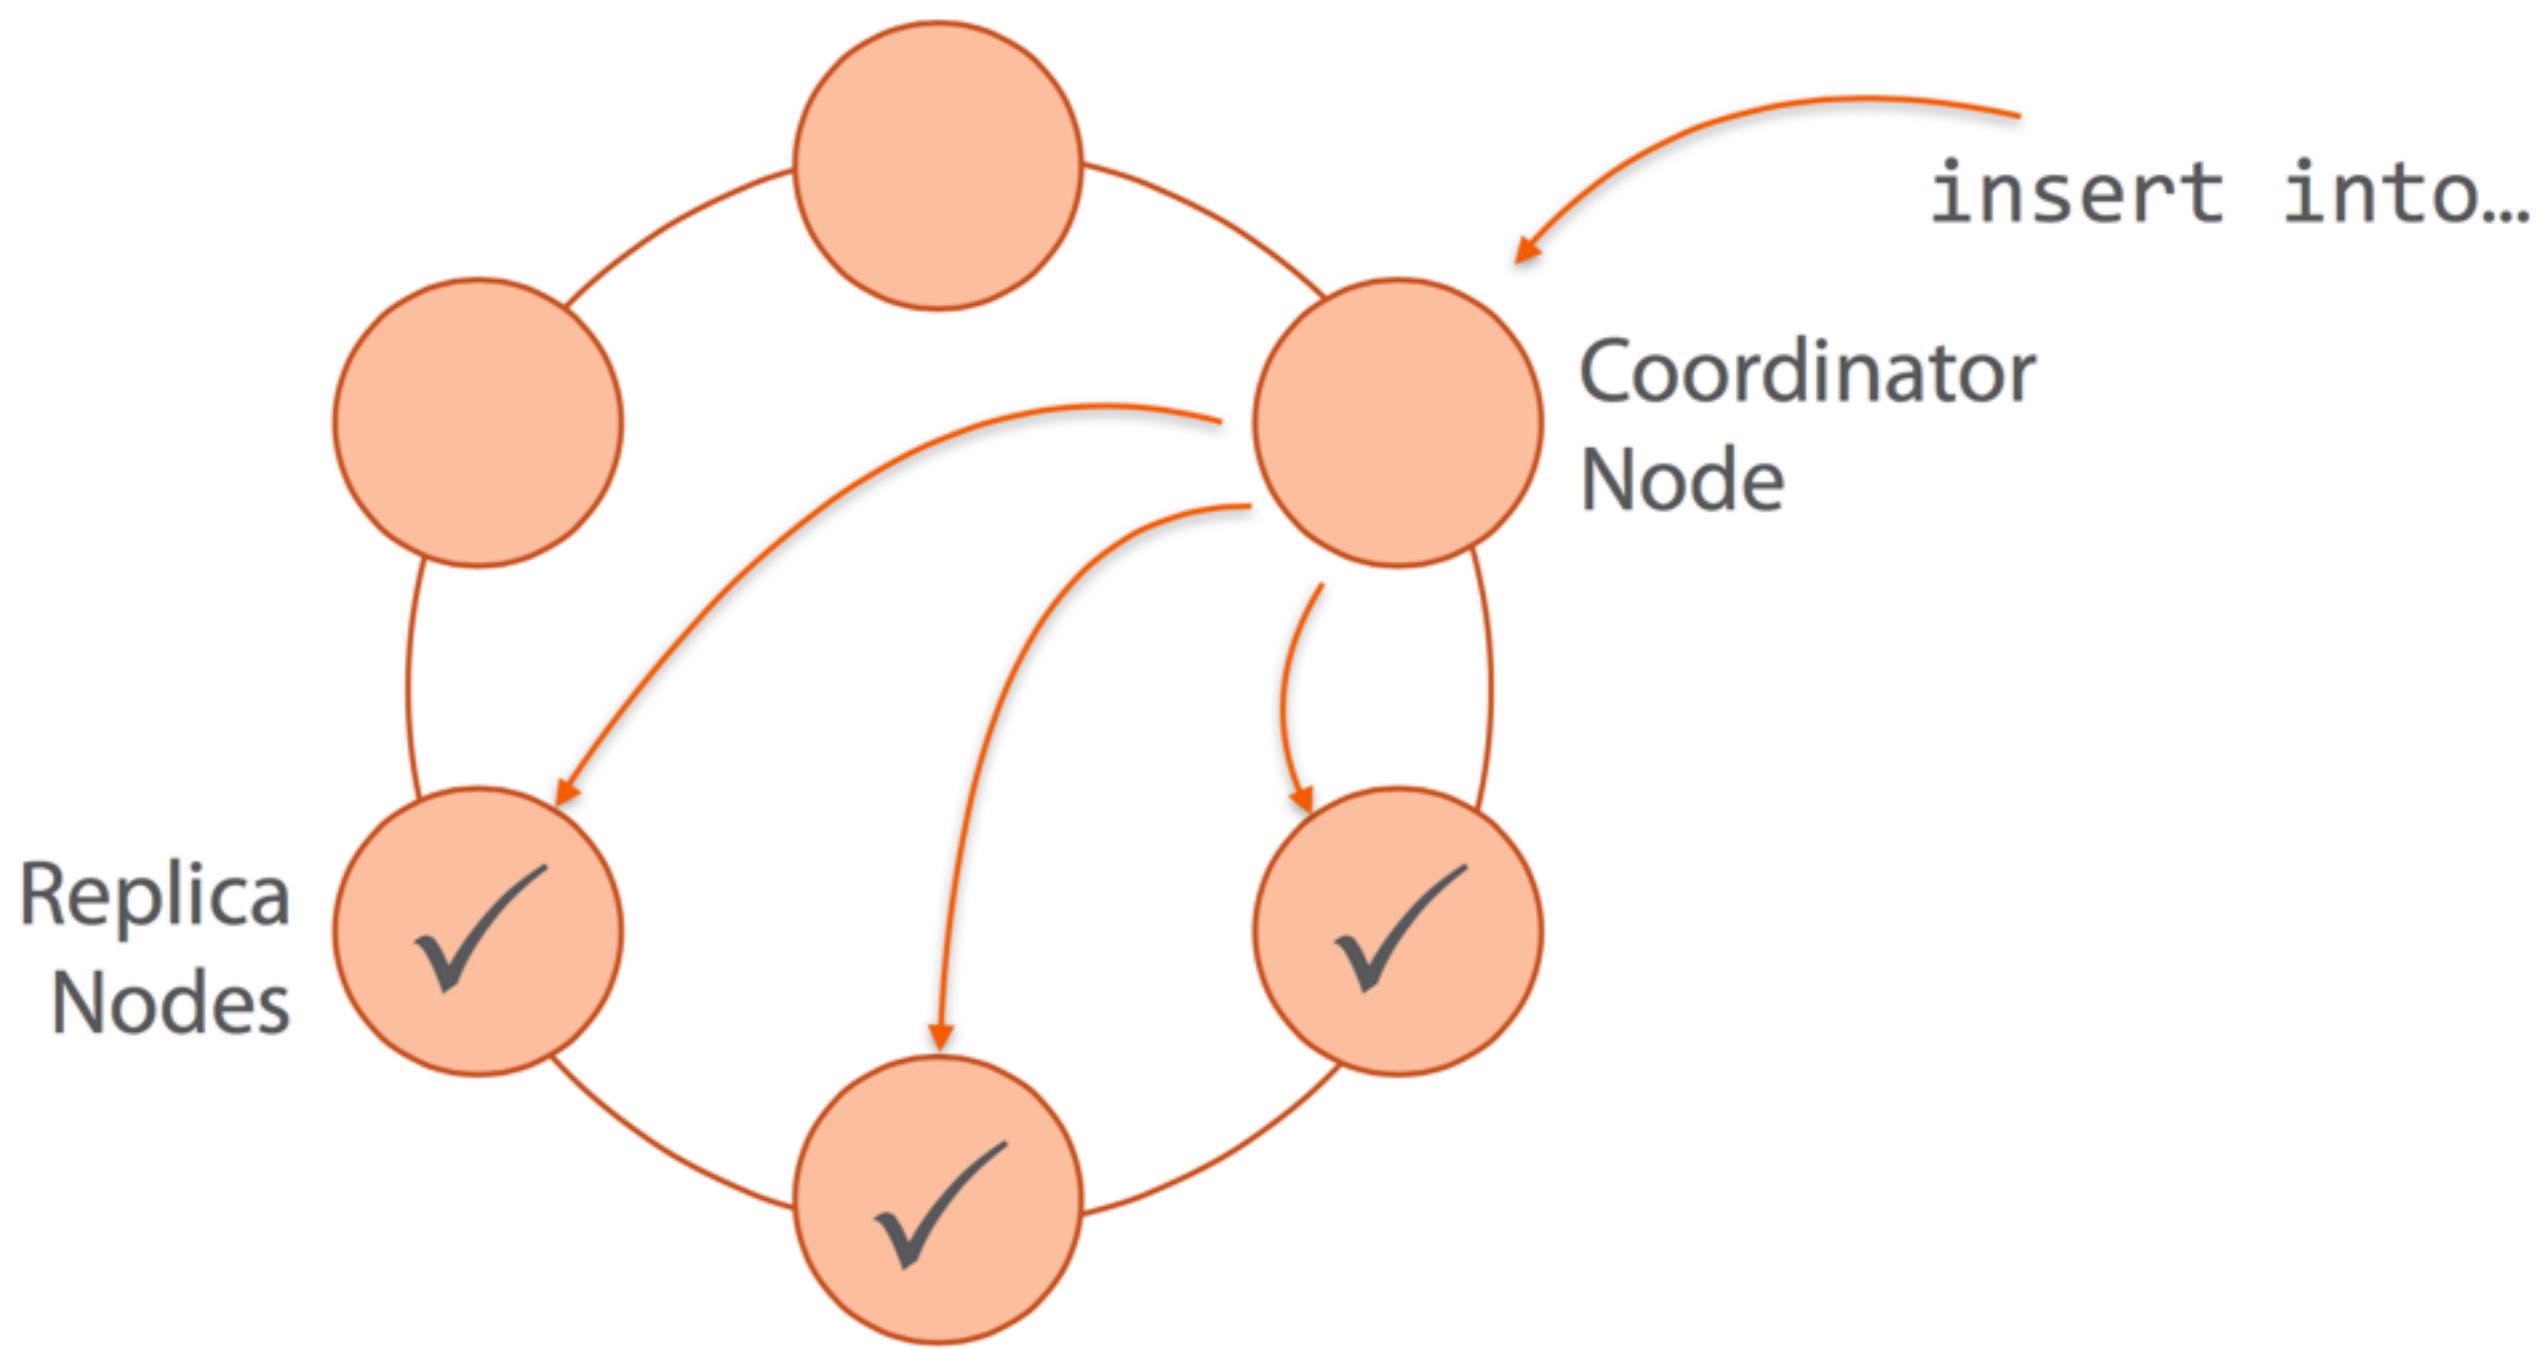


NETWORK TOPOLOGY STRATEGY

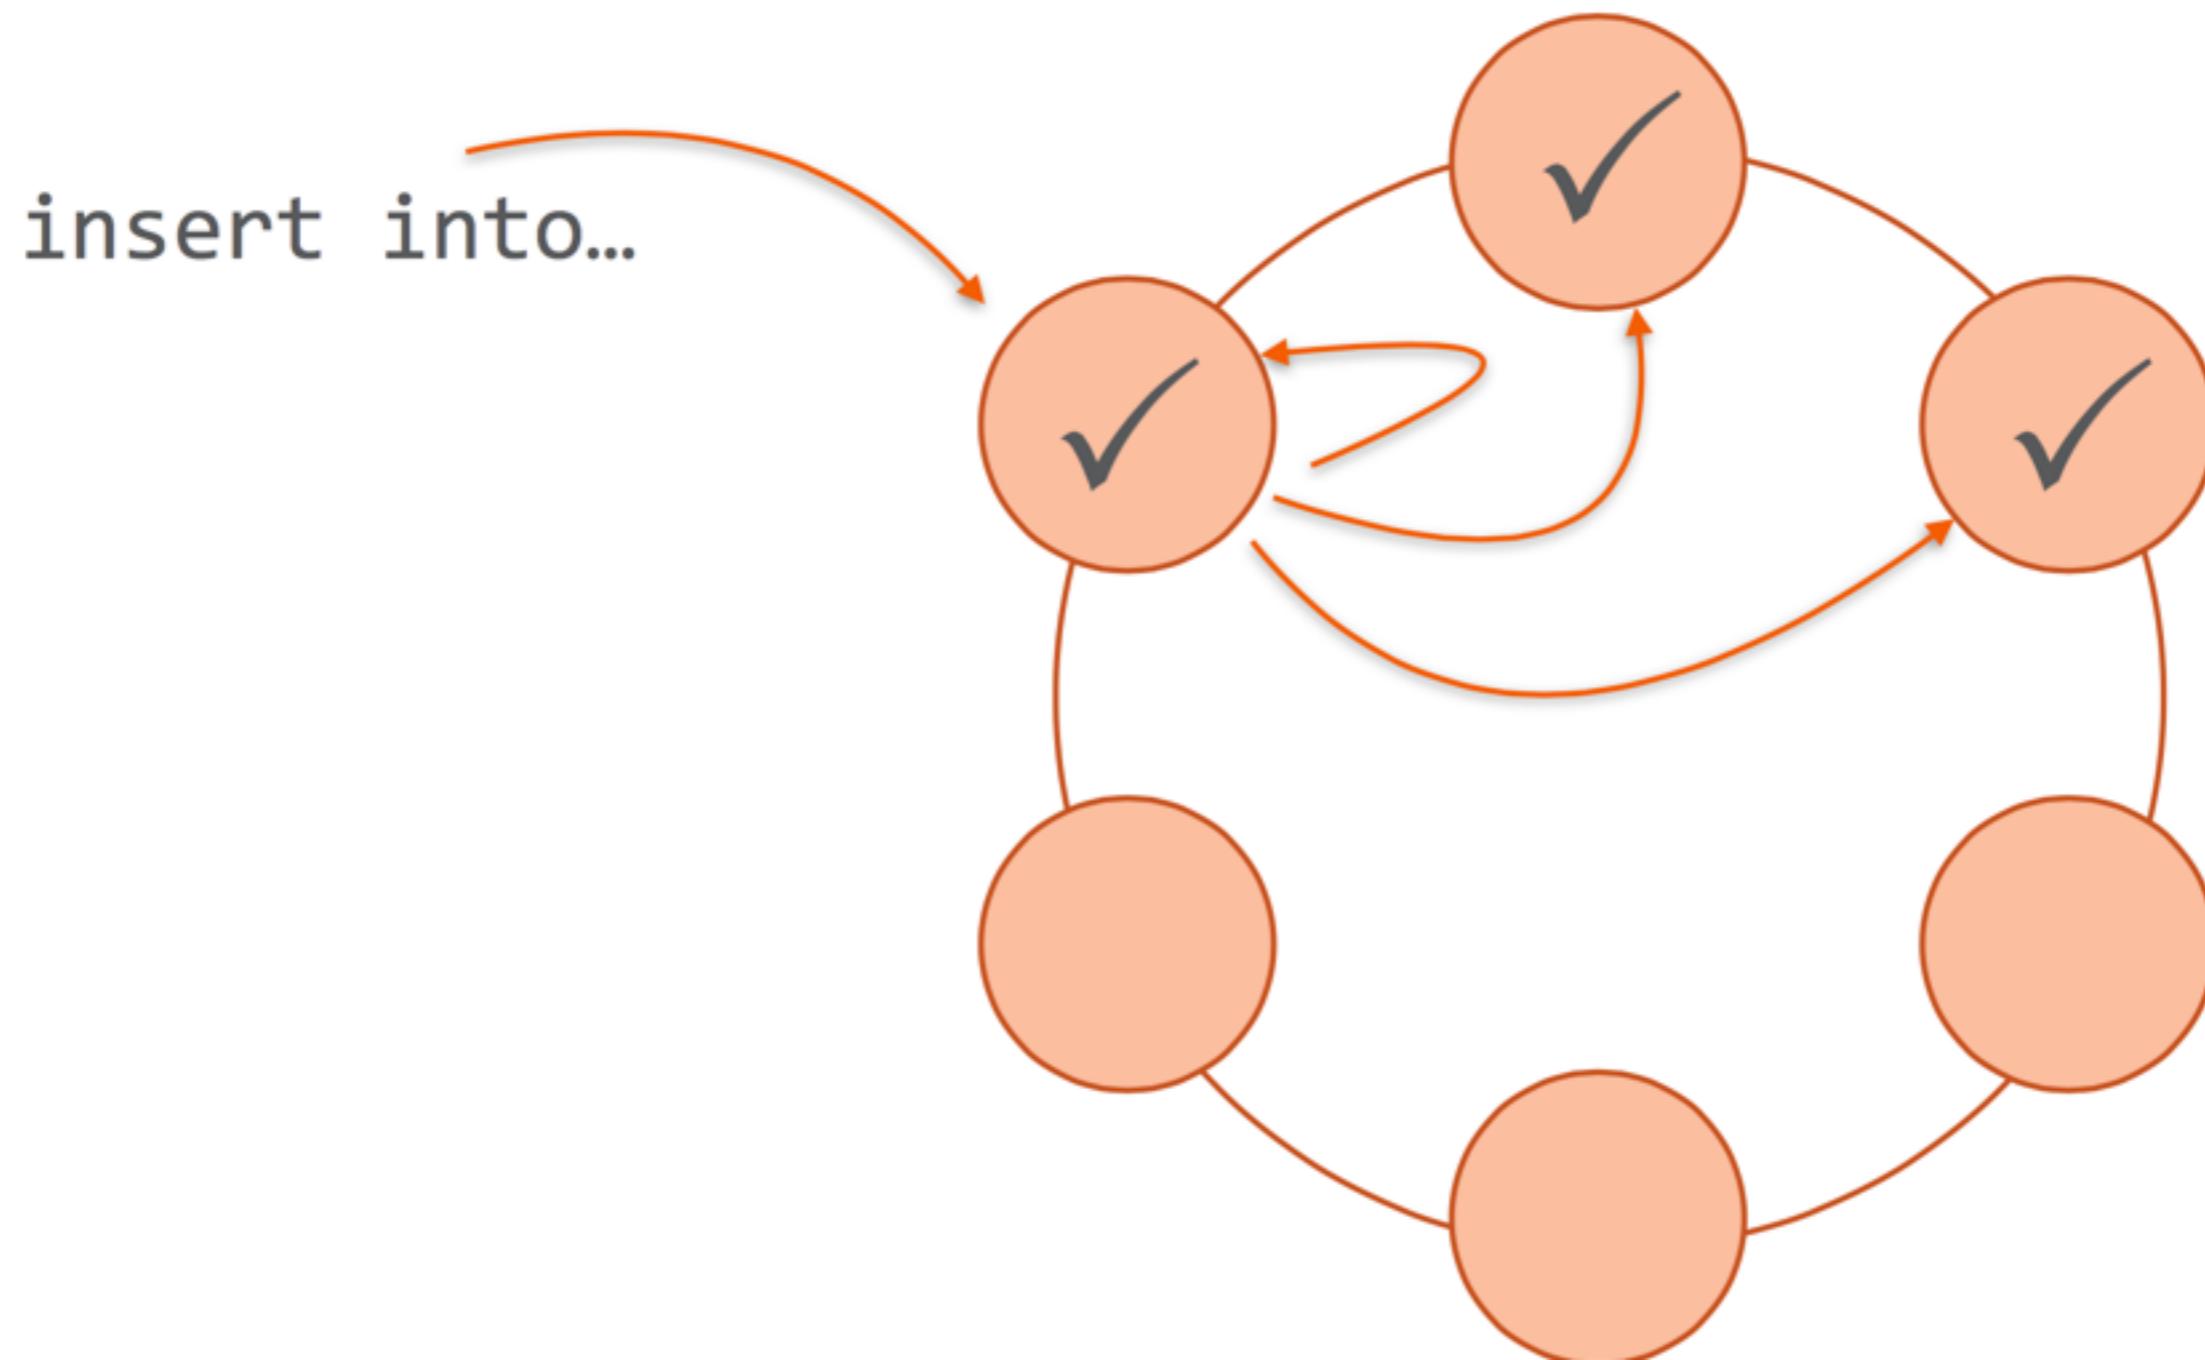
- create KEYSPACE nfjs WITH REPLICATION =
{'class': 'NetworkTopologyStrategy', 'DC1': 3, 'DC2':
1};

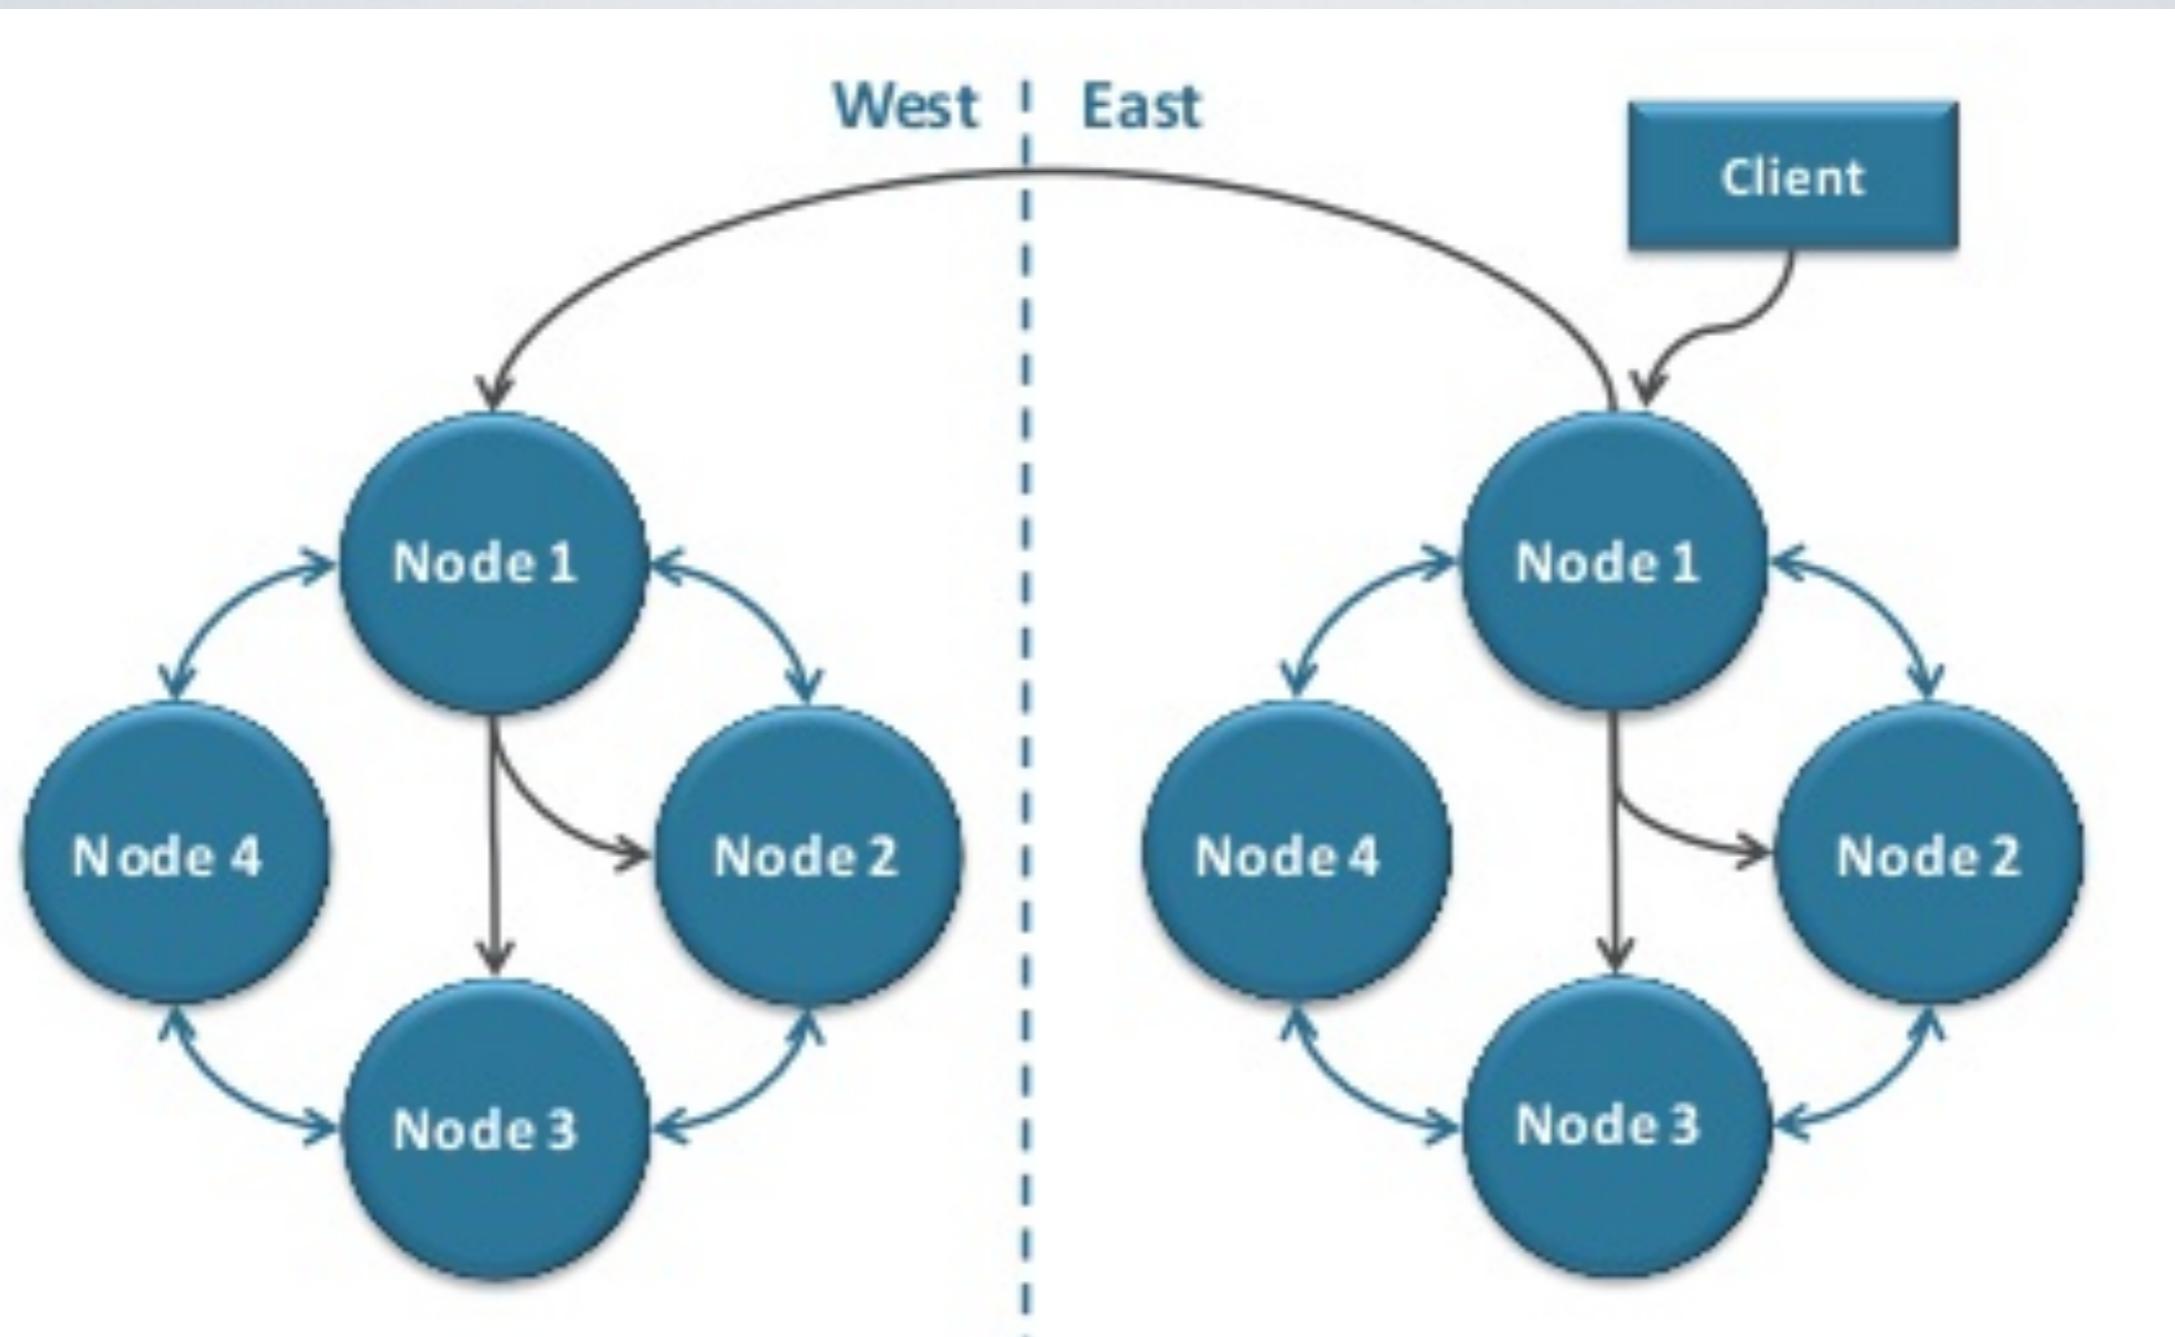


READS AND WRITES IN CASSANDRA



READS AND WRITES IN CASSANDRA

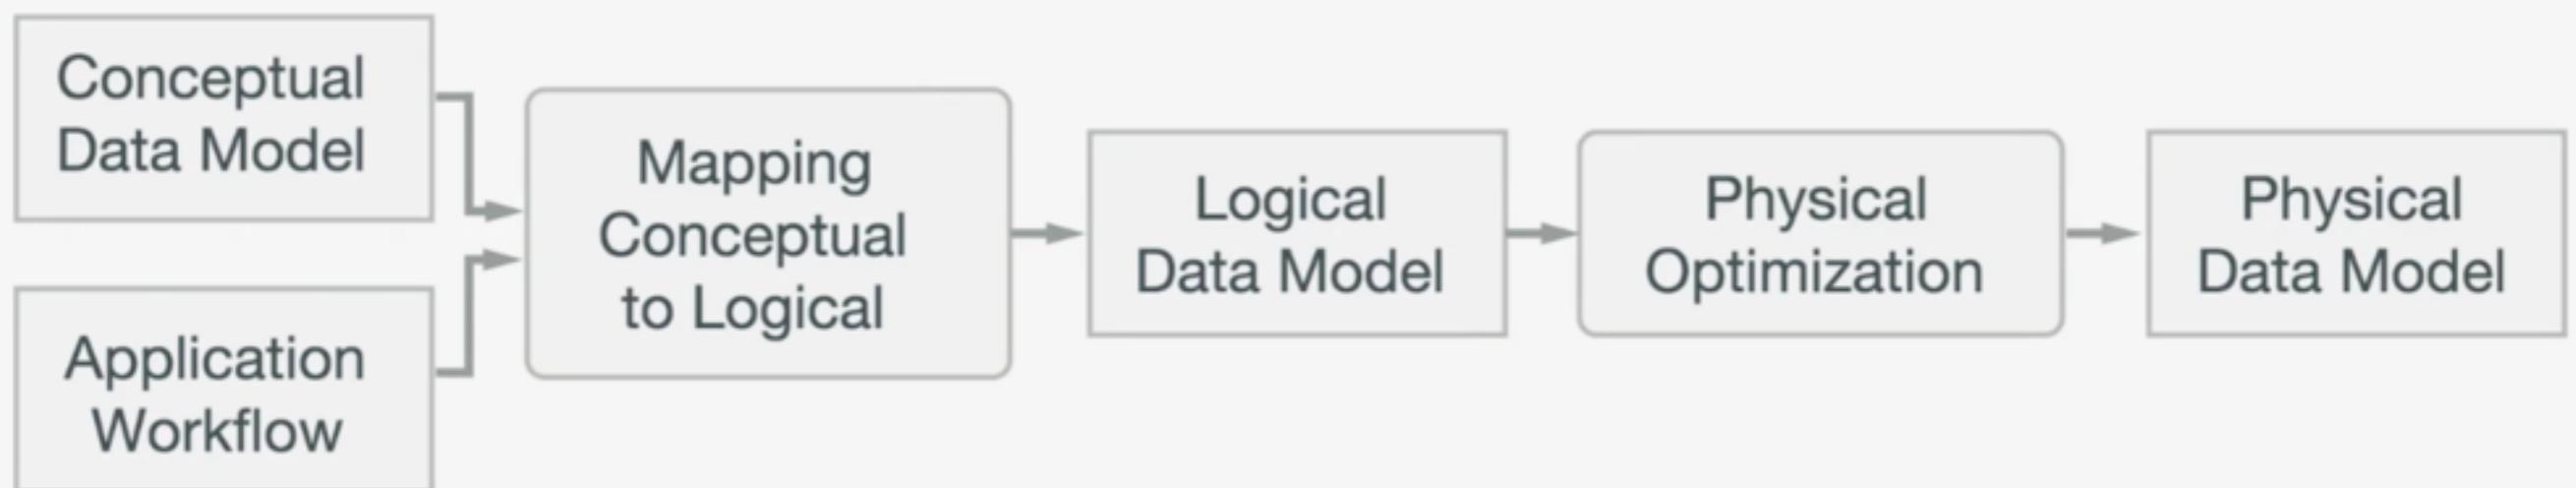




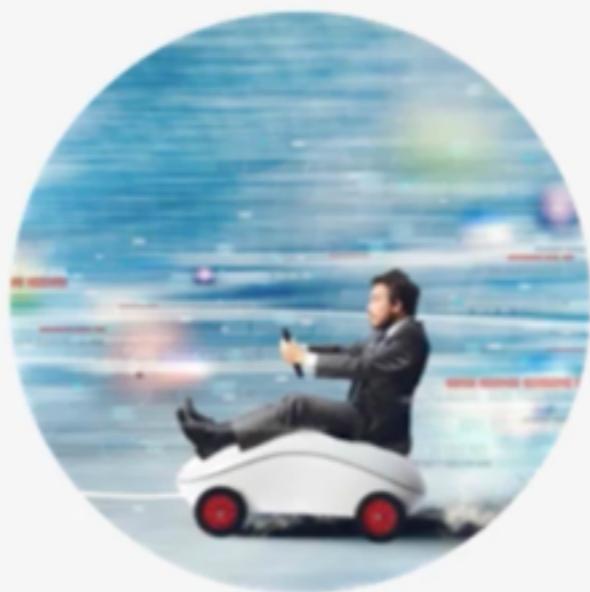
CASSANDRA DEMO

- <http://www.planetcassandra.org/try-cassandra/>

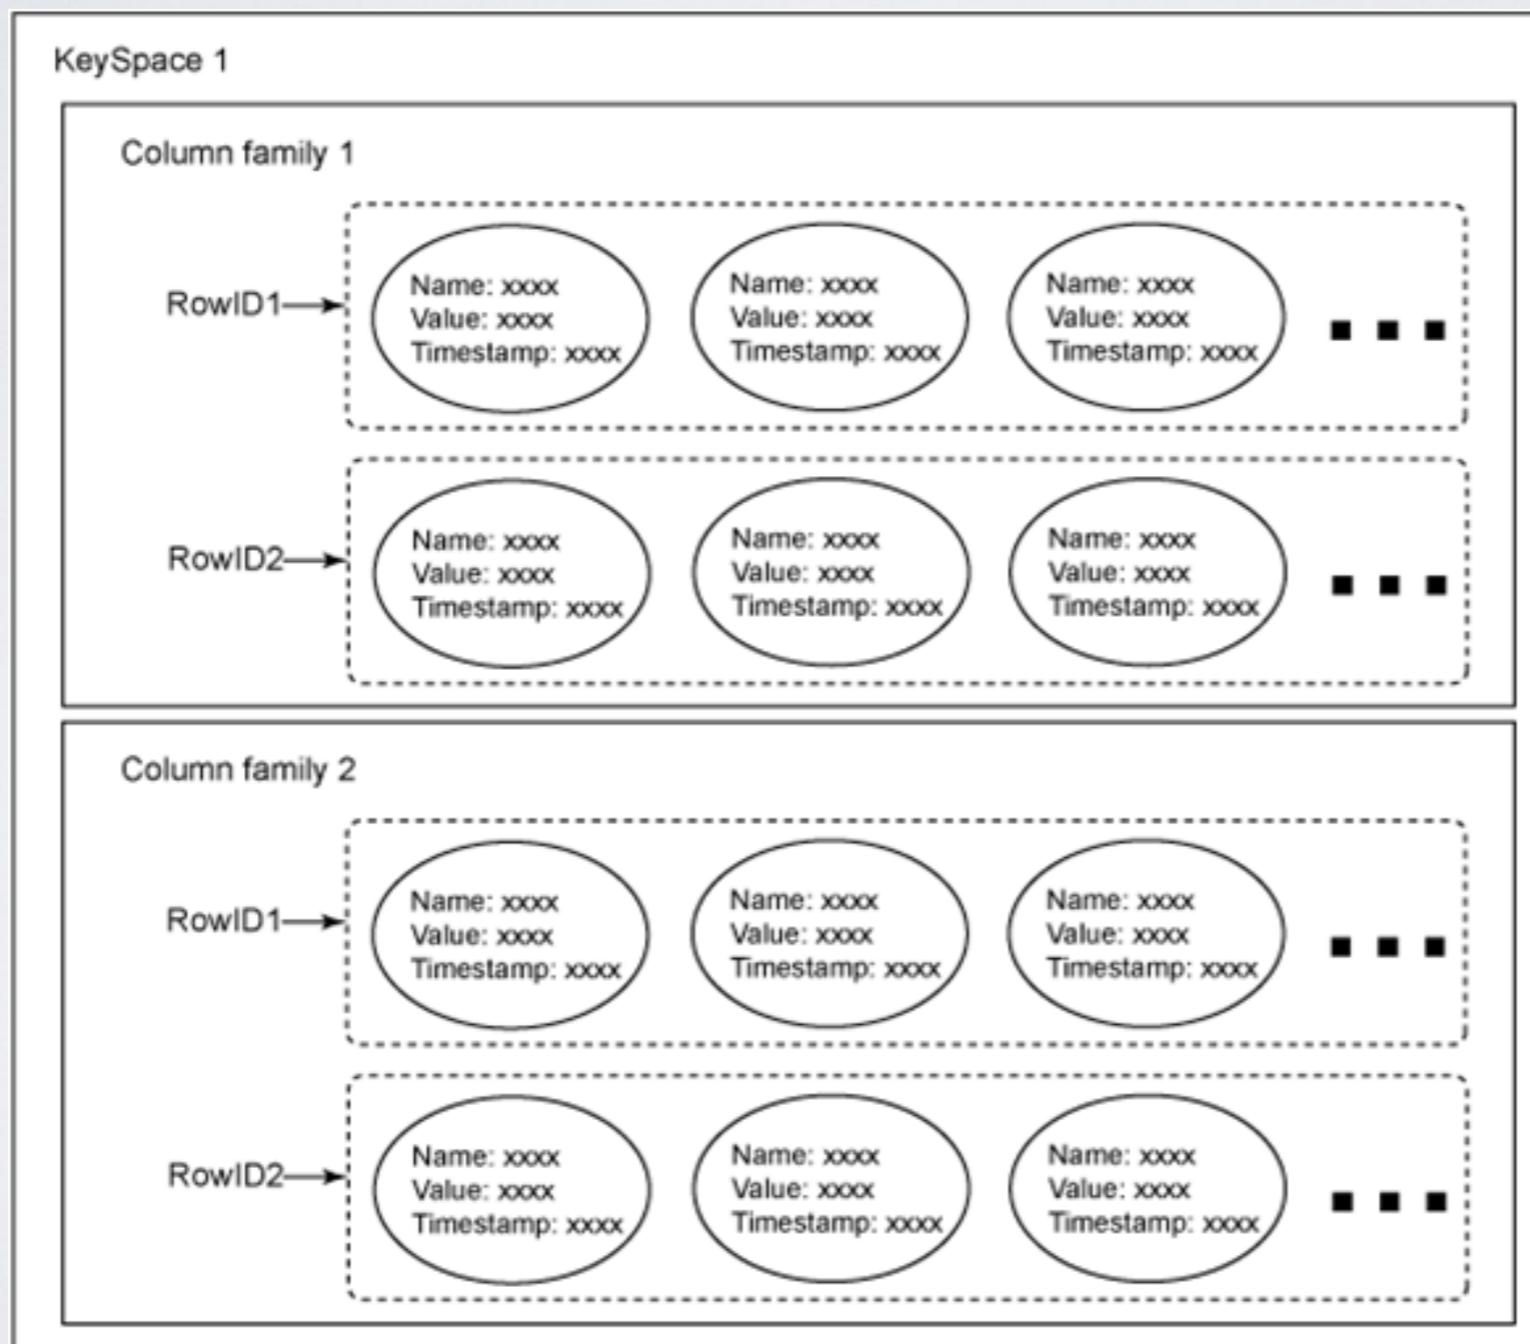
CASSANDRA DATA MODEL



QUERIES



KEY-VALUE MODEL



Column Name

Column Family: Tweets

	Text	User_ID	Date
1234e530-8b82-11df	Hello, World!	39823	2009-03-25T19:20:30

	Text	User_ID	Date
22615e20-8b82-11df	Gooooal!	592	2009-03-25T19:25:43

Key

Column Value

COLUMN NAMES CAN STORE VALUES

Column Family: User_Timelines

39823	cef7be80-8b88-11df	1234e530-8b82-11df	...
	-	-	...

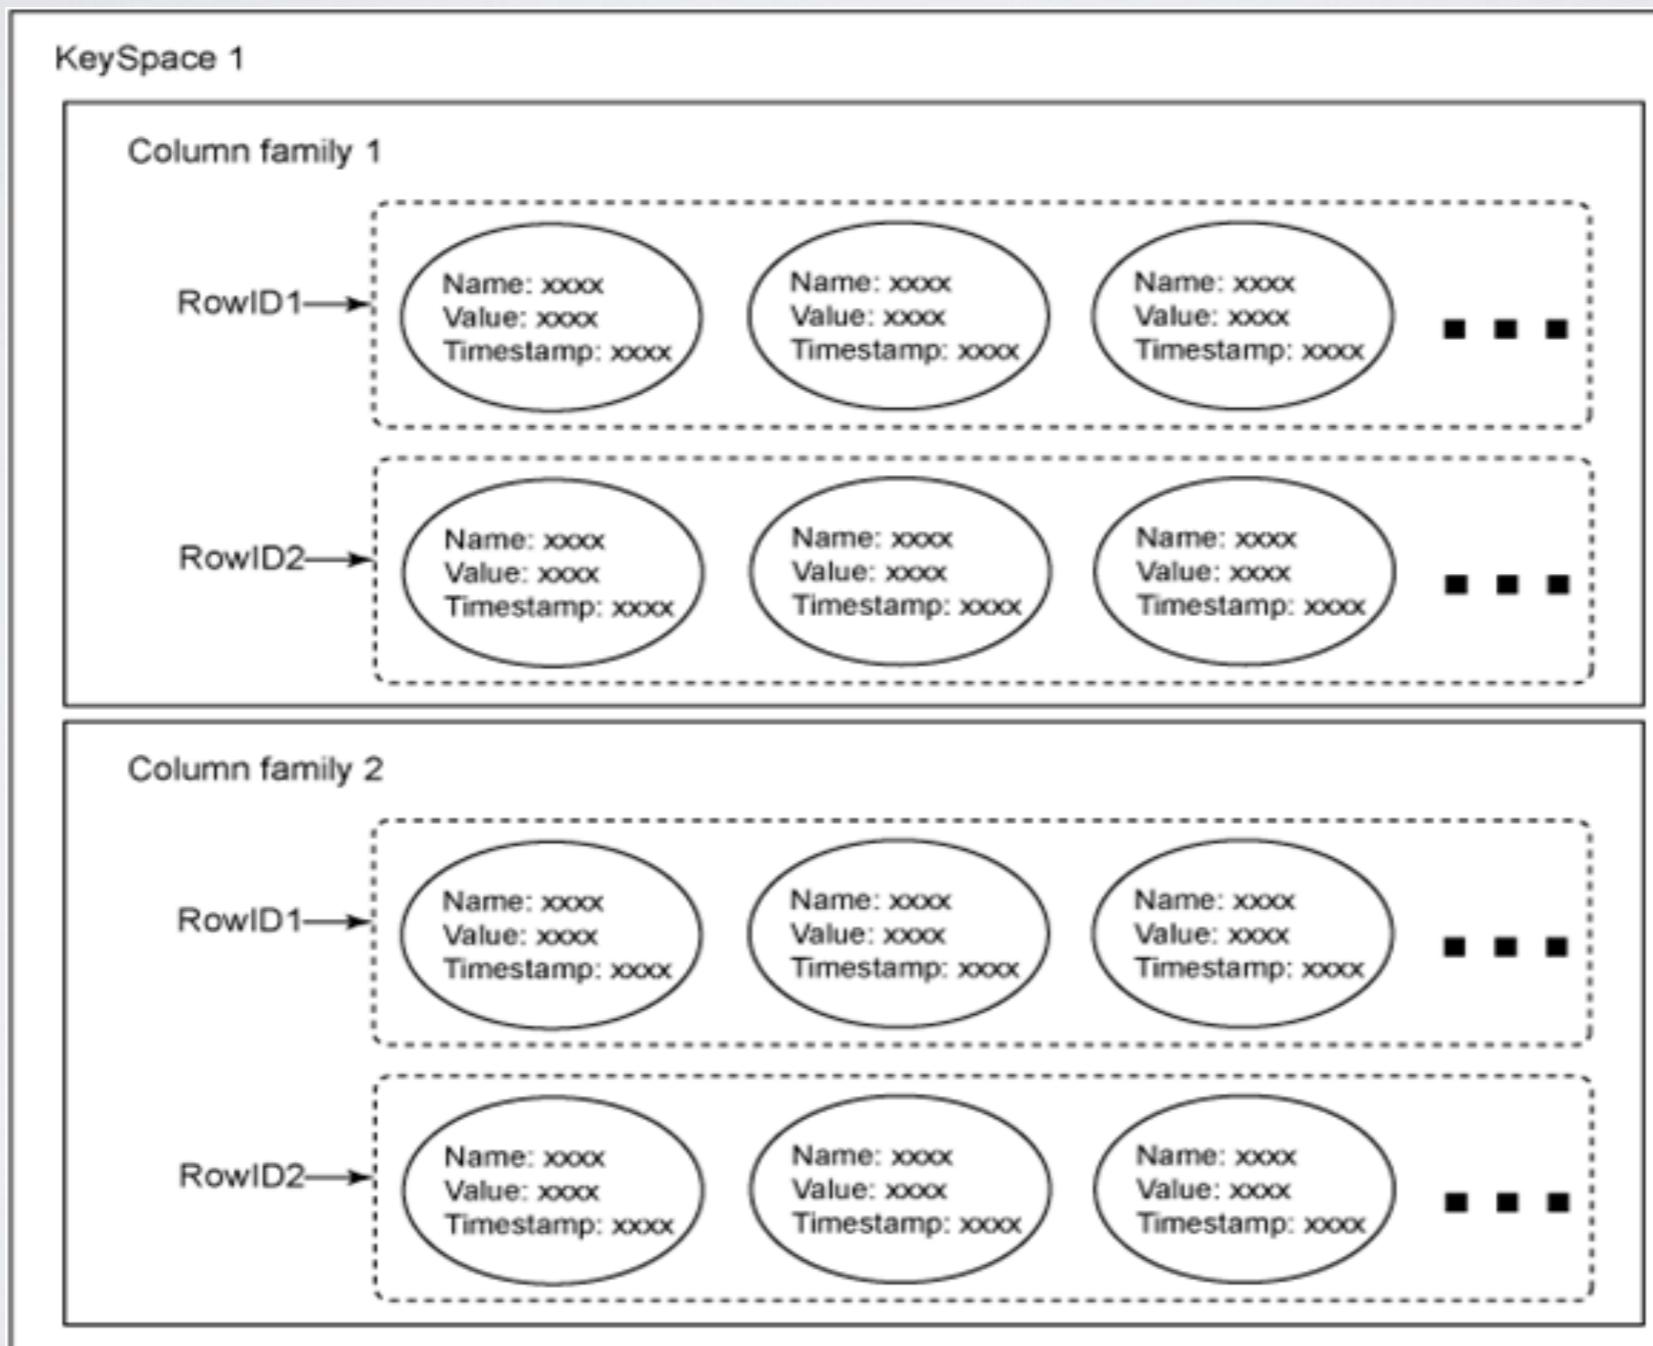
592	f0137940-8b8a-11df	22615e20-8b82-11df	...
	-	-	...

• • •

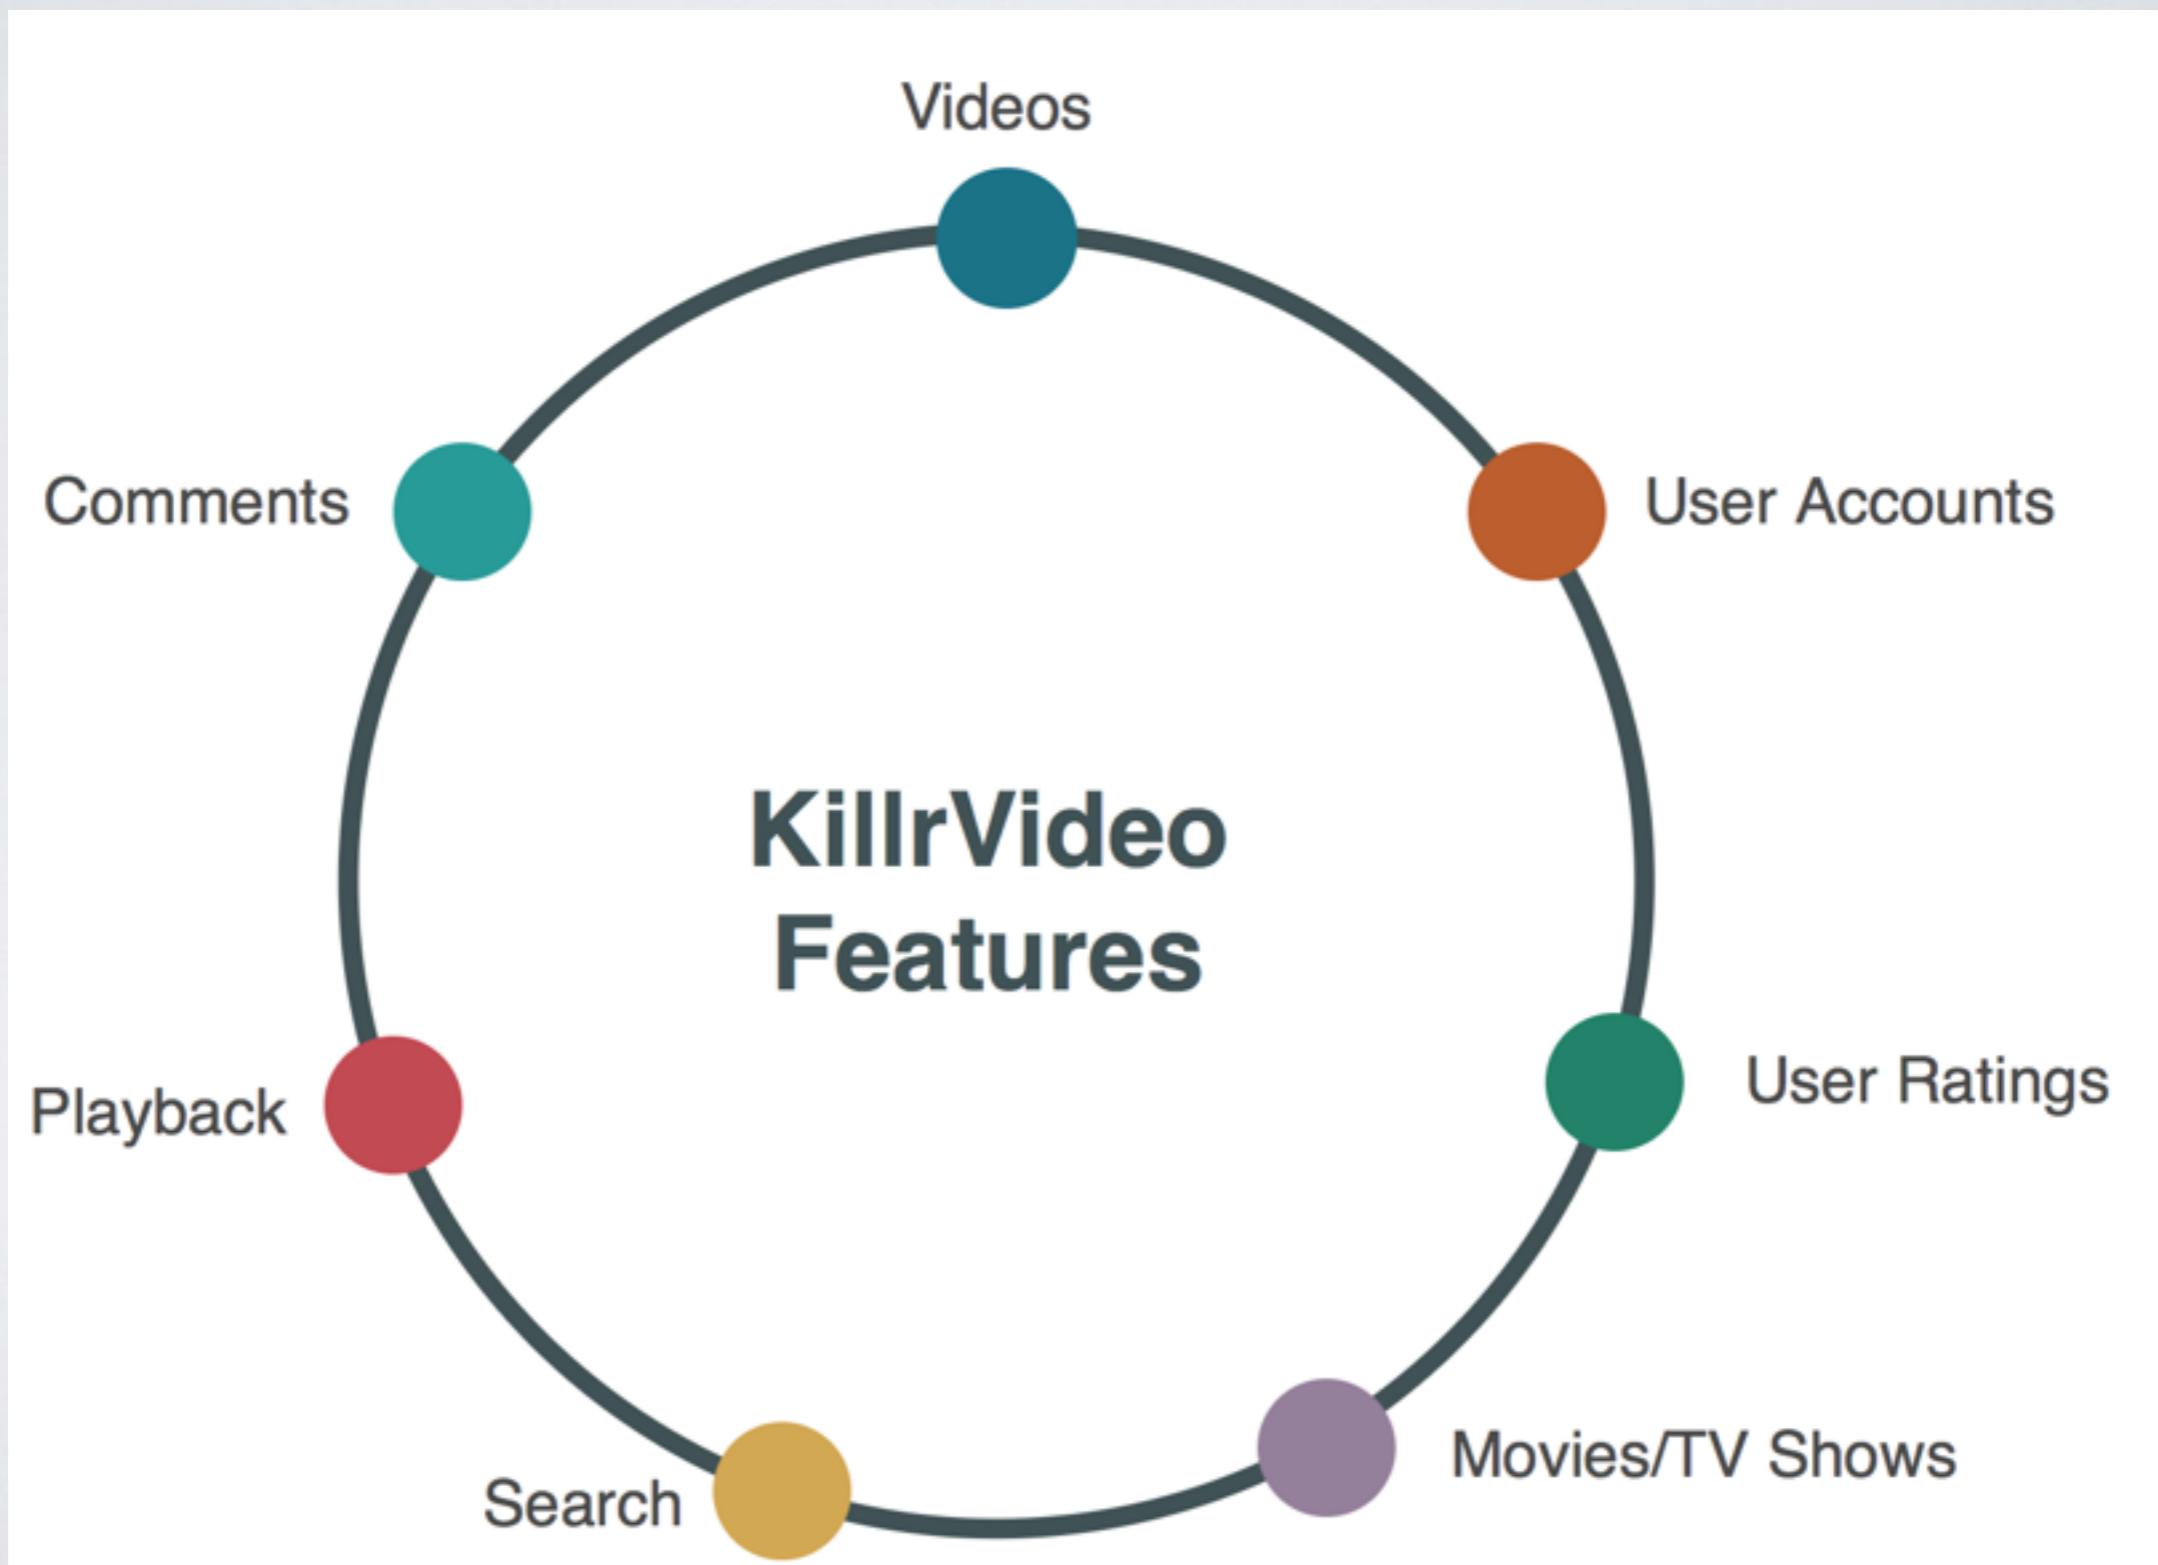
UUID AND TIMEUUID

- Universally Unique Identifier
 - Ex: 52b11d6d-16e2-4ee2-b2a9-5ef1e9589328
 - Generate via `uuid()`
- TIMEUUID embeds a TIMESTAMP value
 - Ex: 1be43390-9fe4-11e3-8d05-425861b86ab6
 - Sortable
 - Generate via `now()`

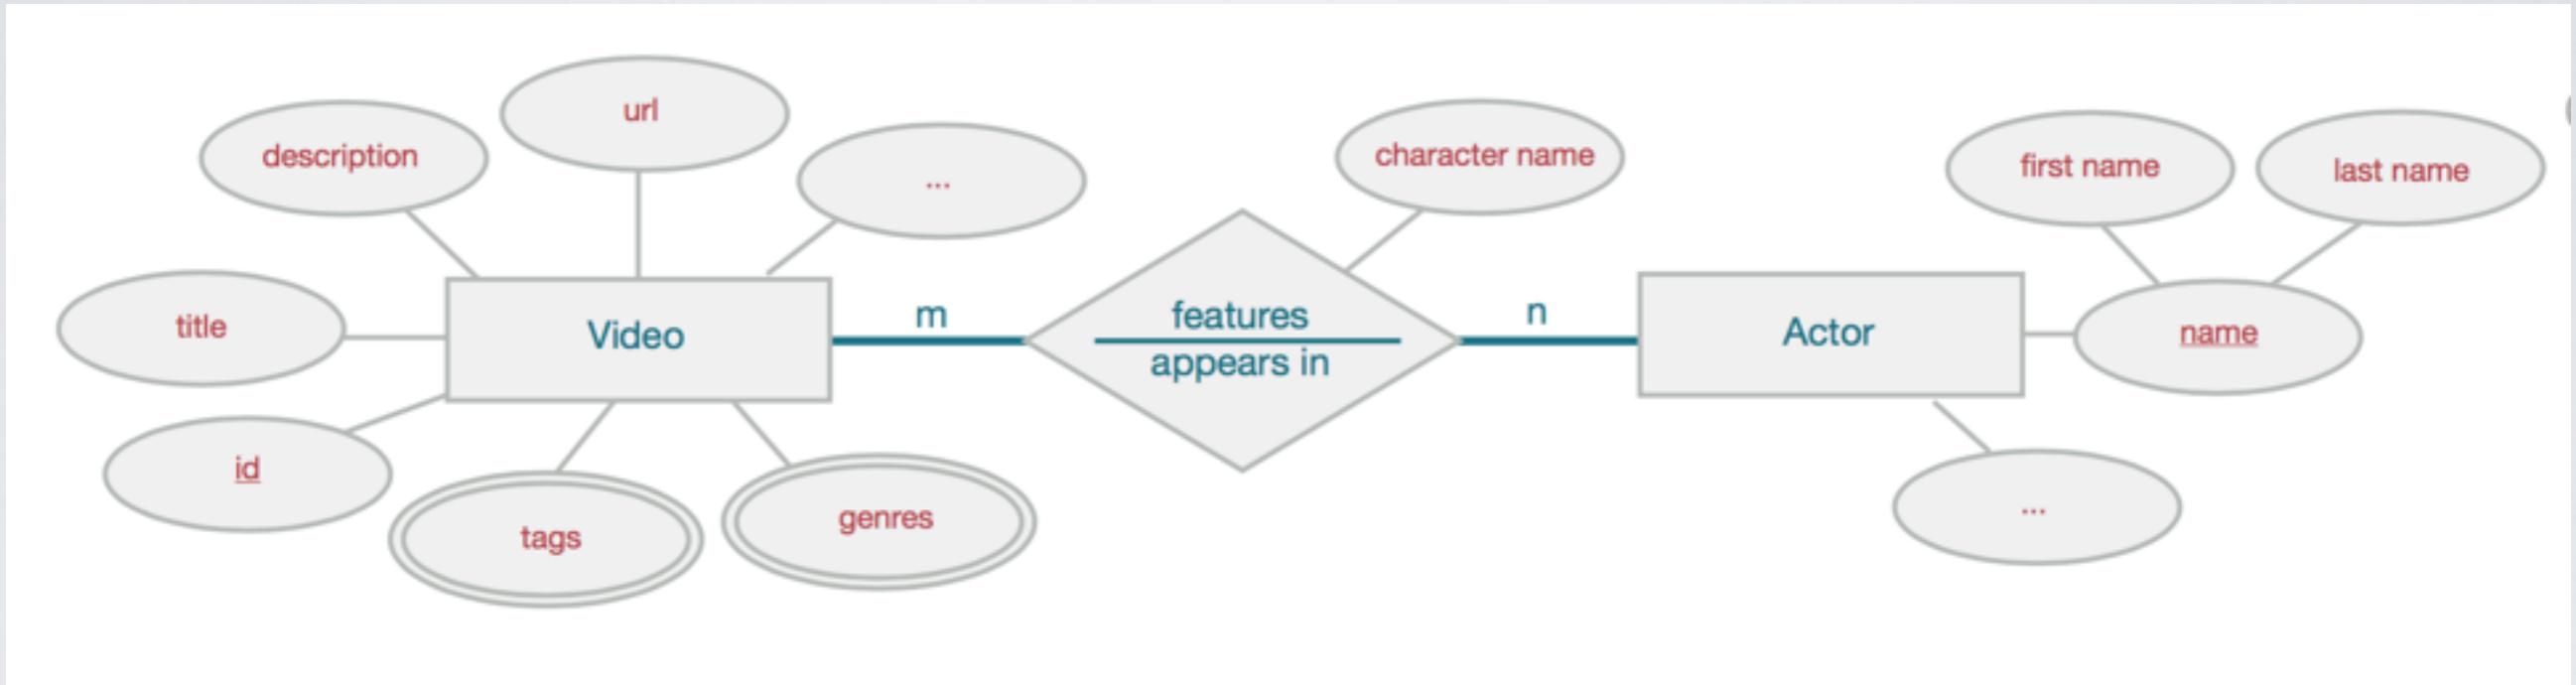
KEYSPACE



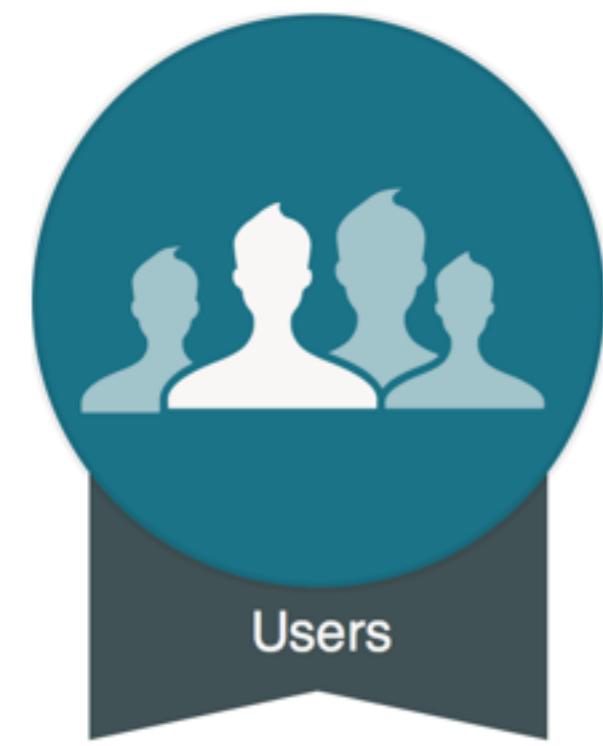
KILLRVIDEO.COM



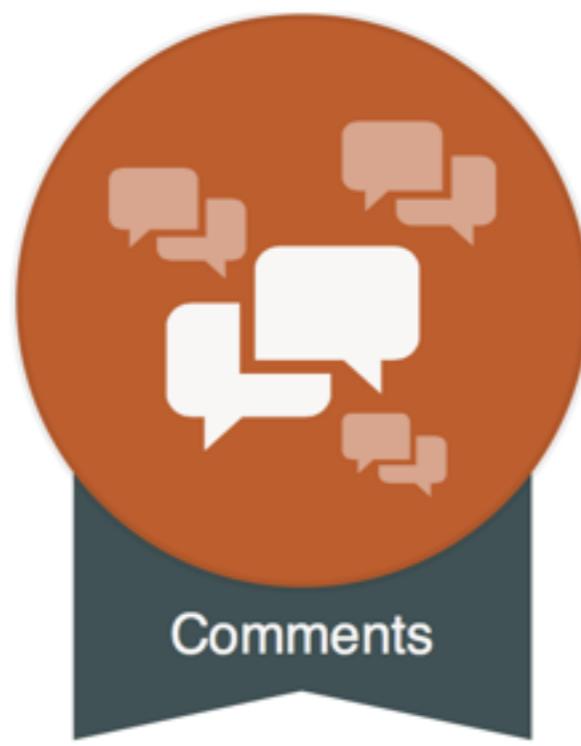
CONCEPTUAL DATA MODELING



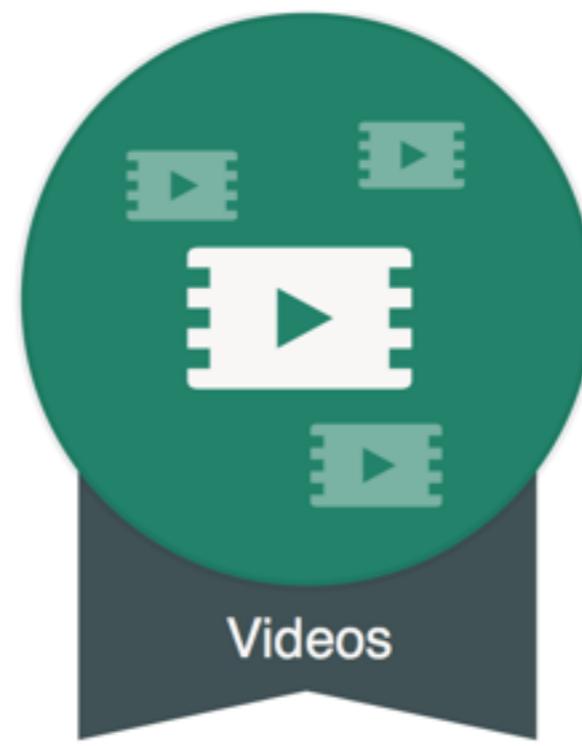
PURPOSE



Users



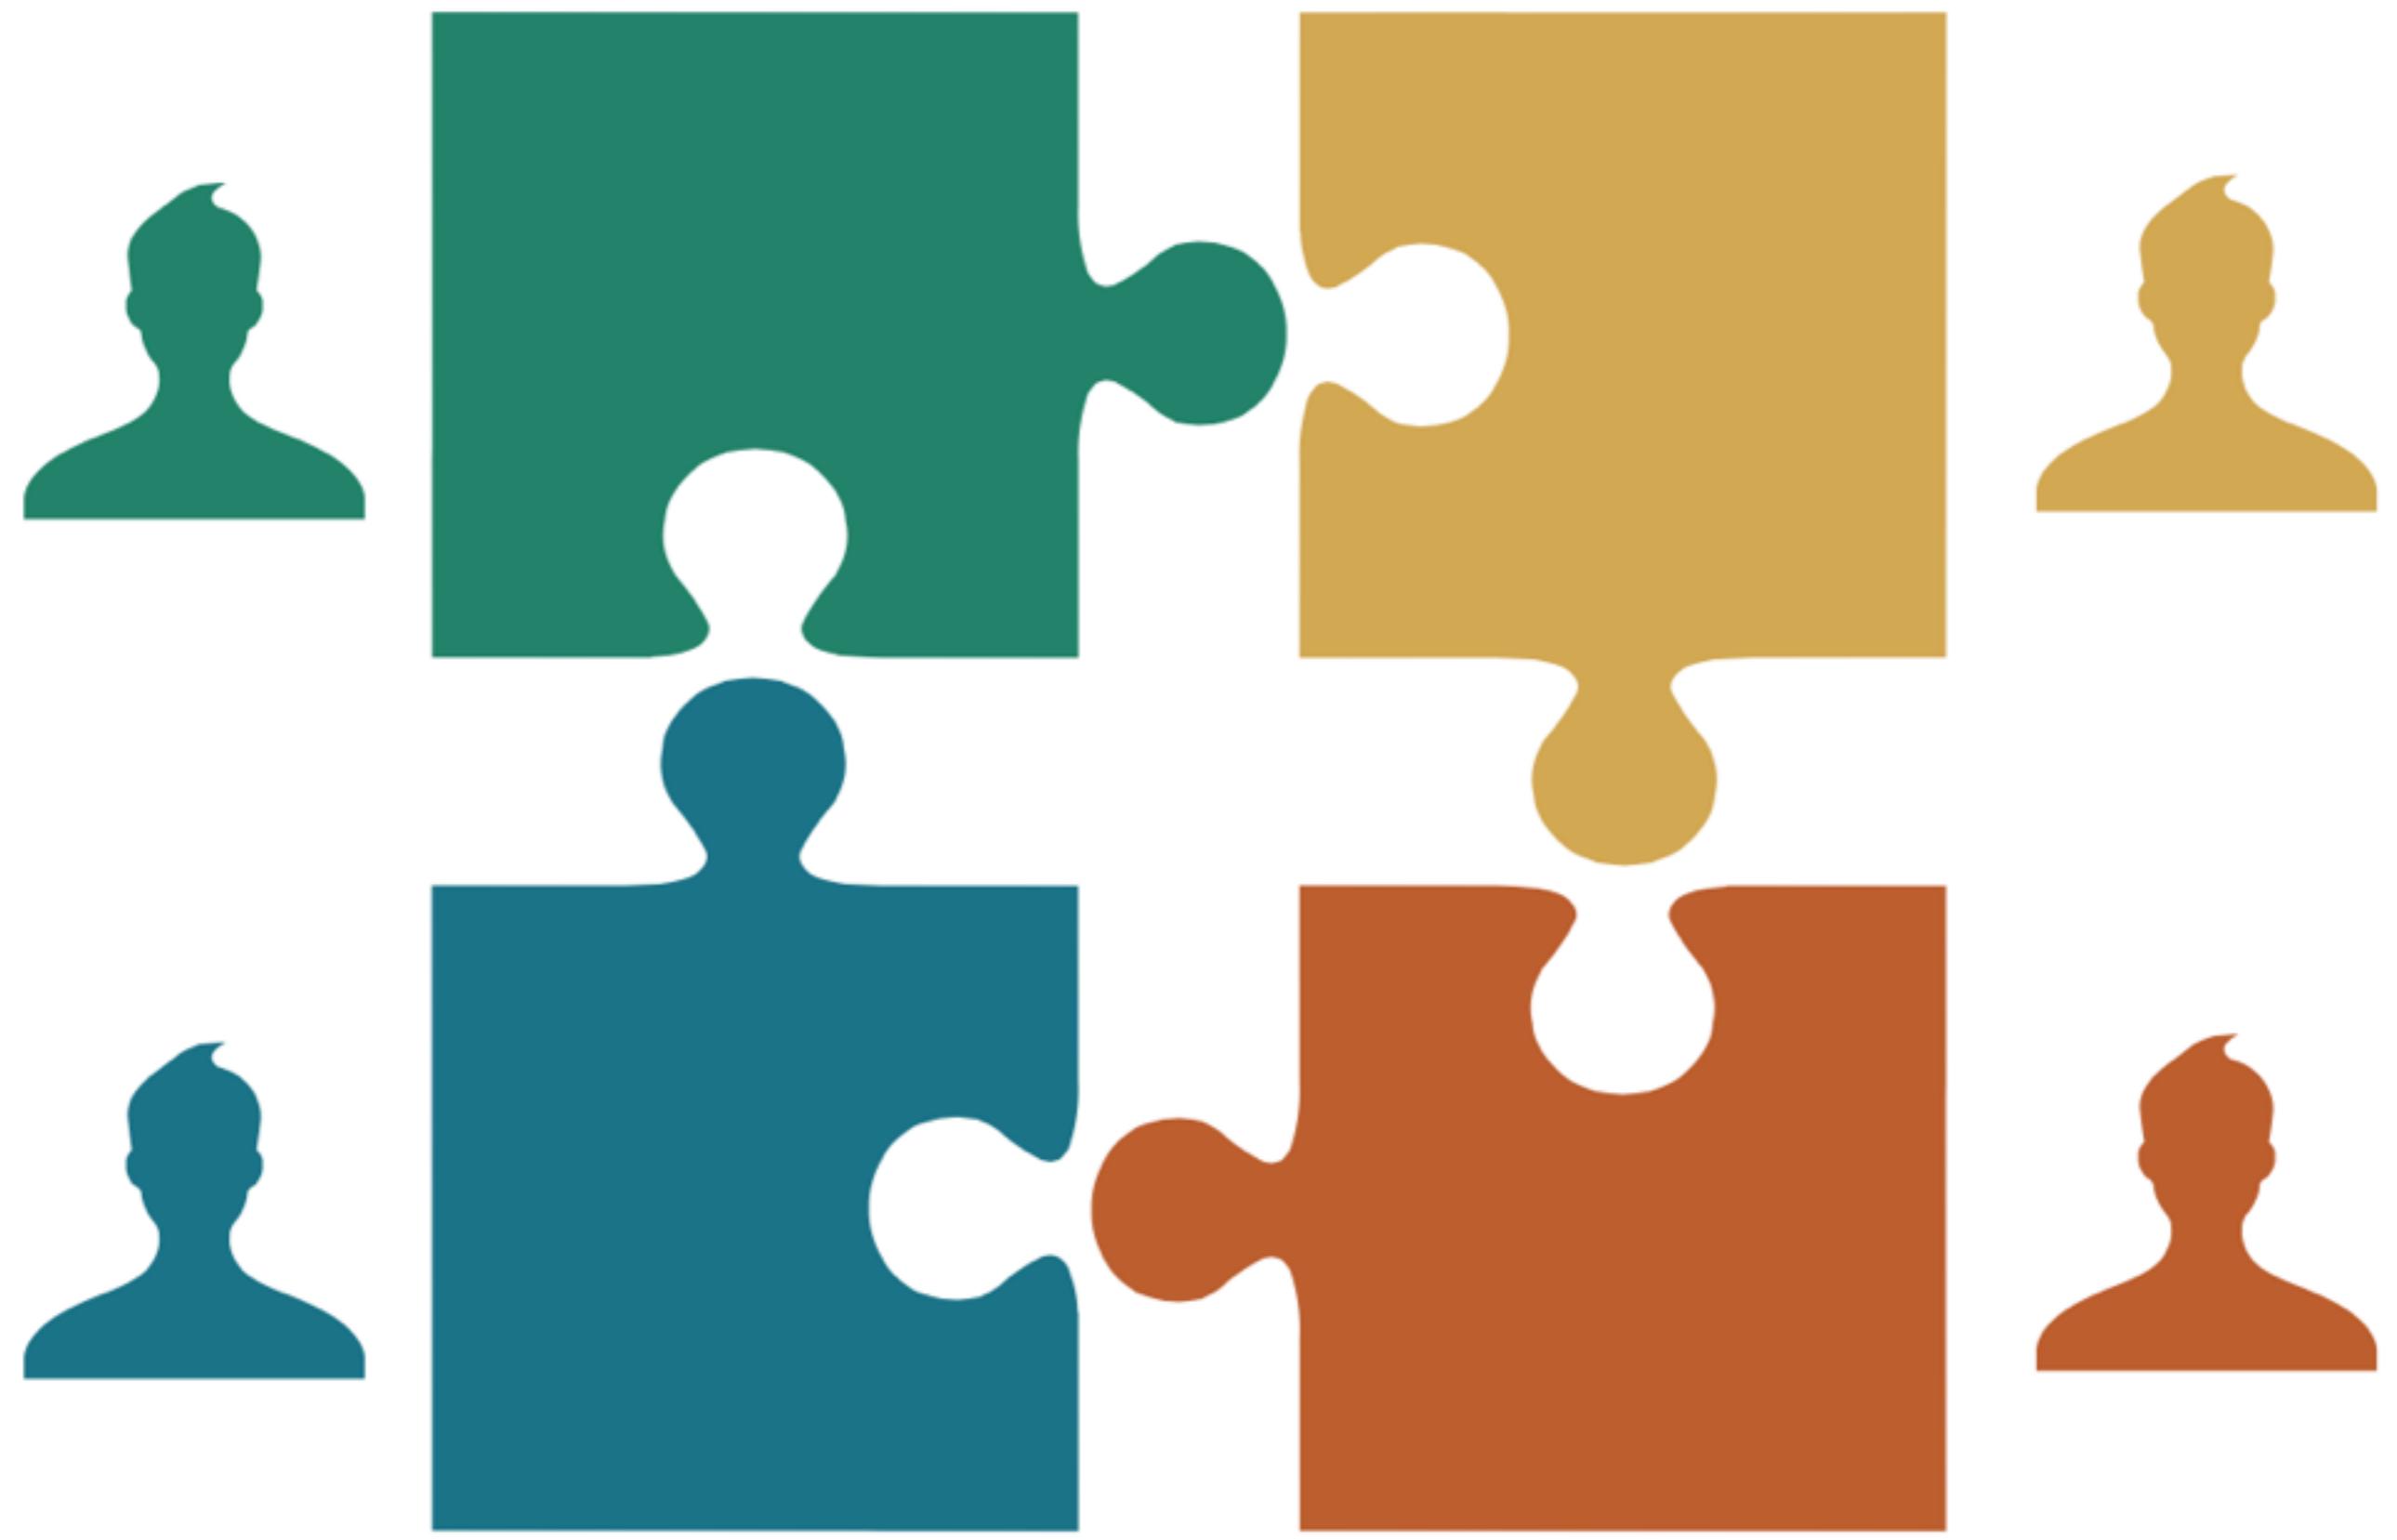
Comments



Videos



Ratings



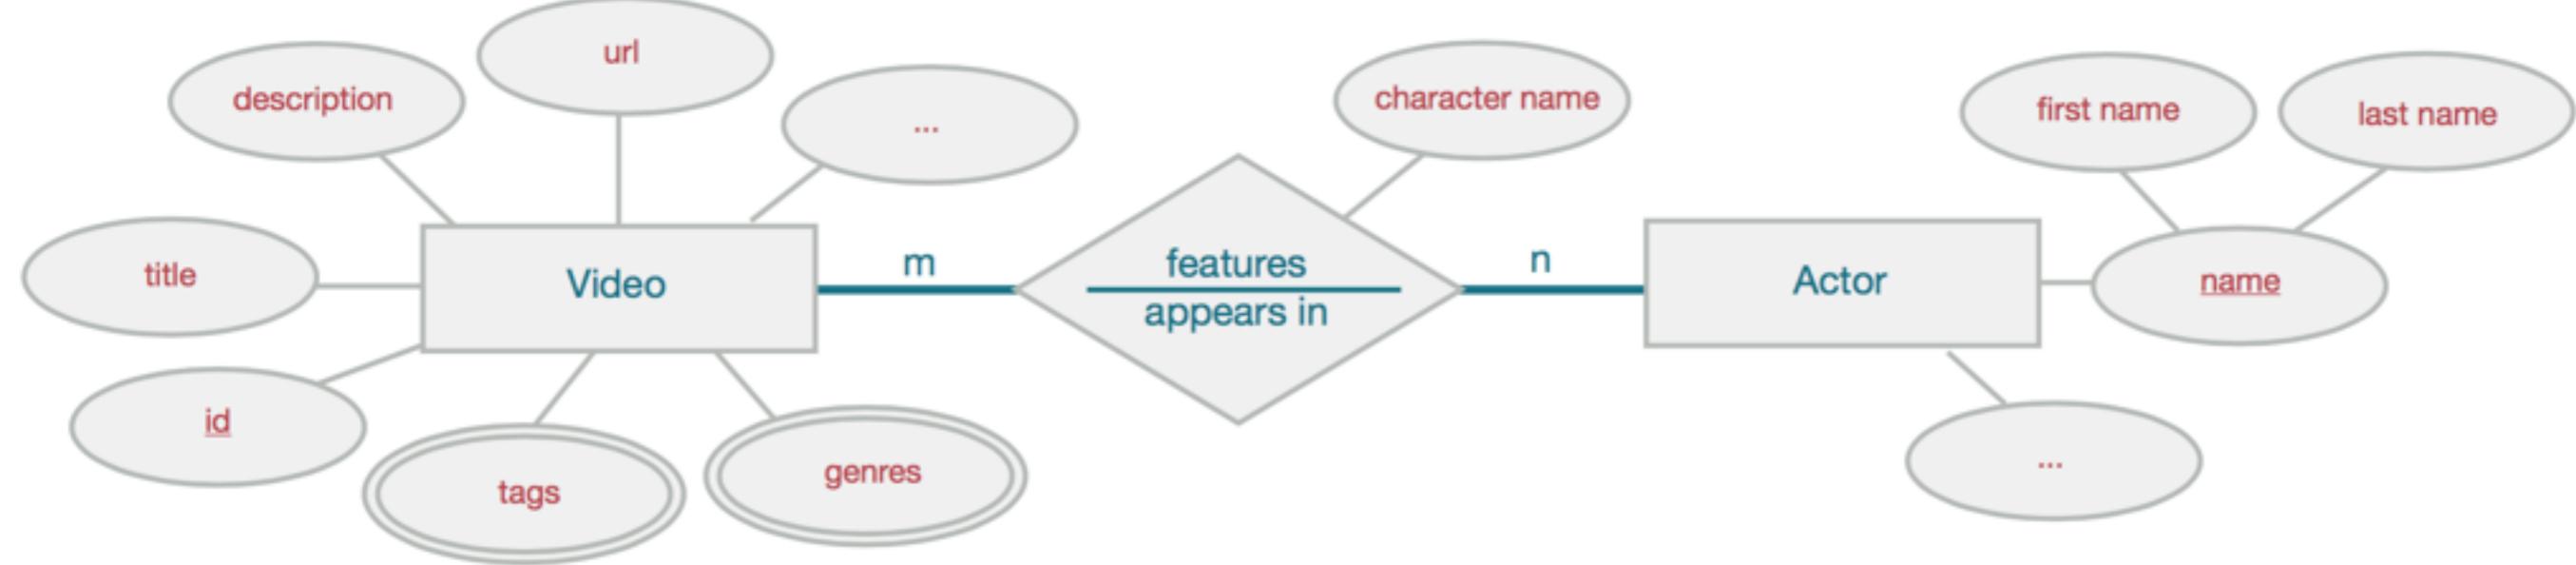
ABSTRACTION



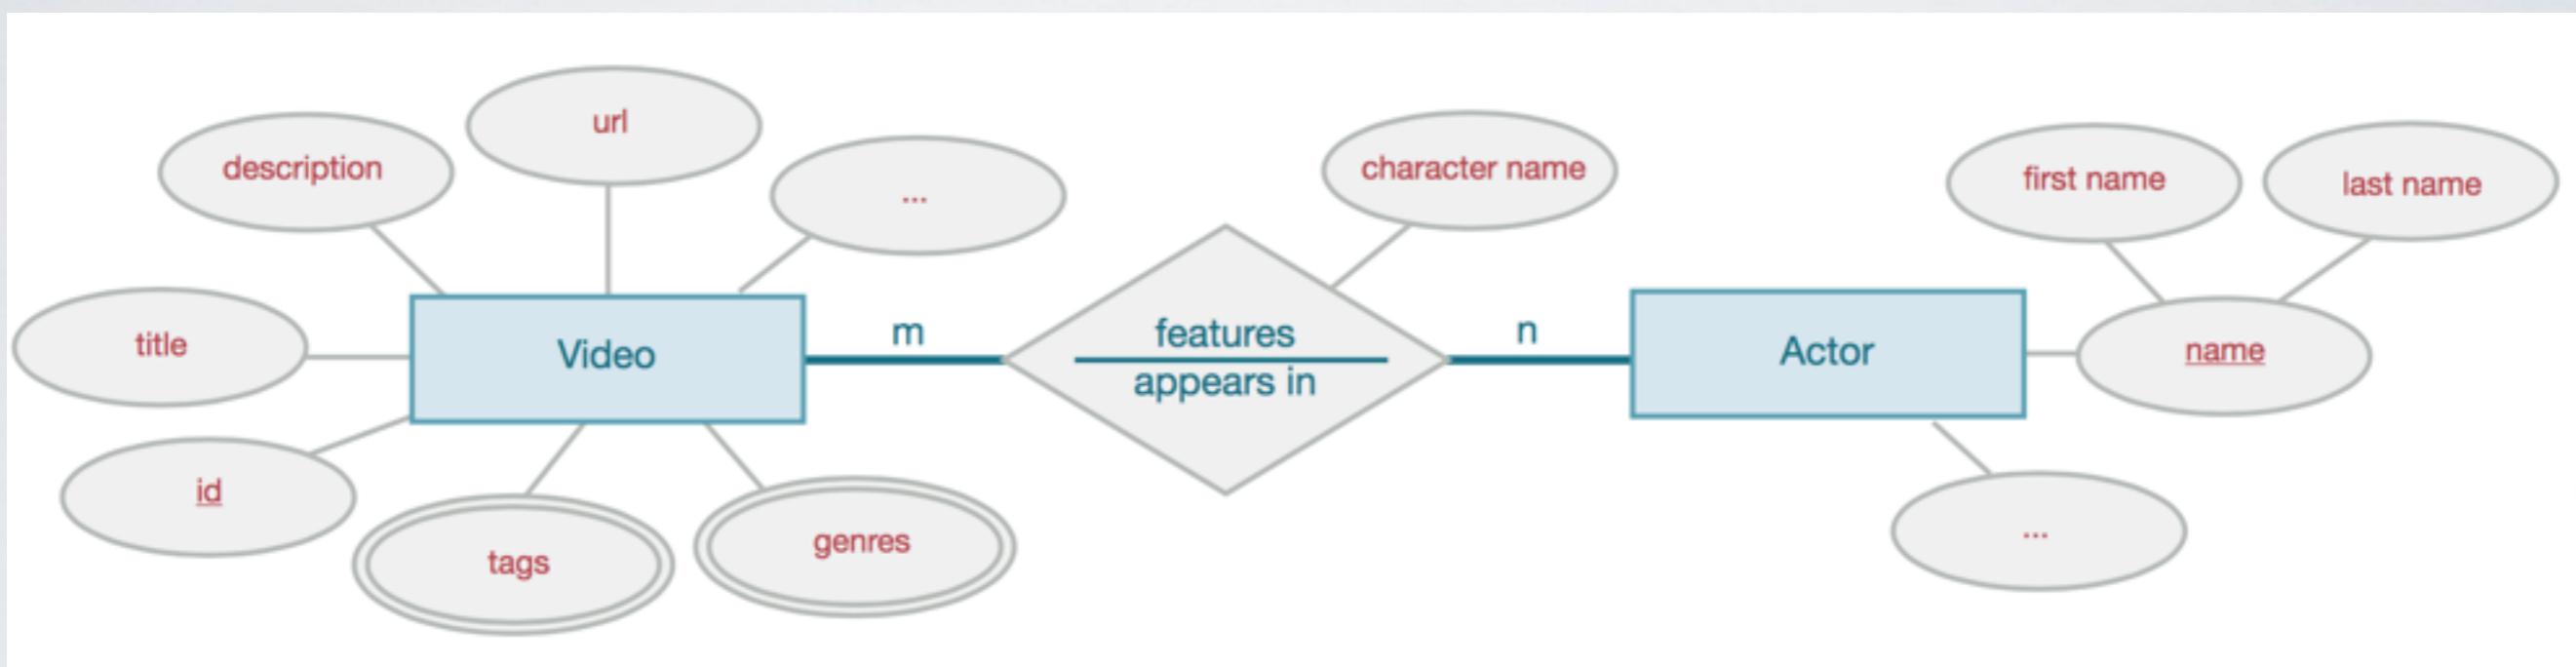
ABSTRACTION



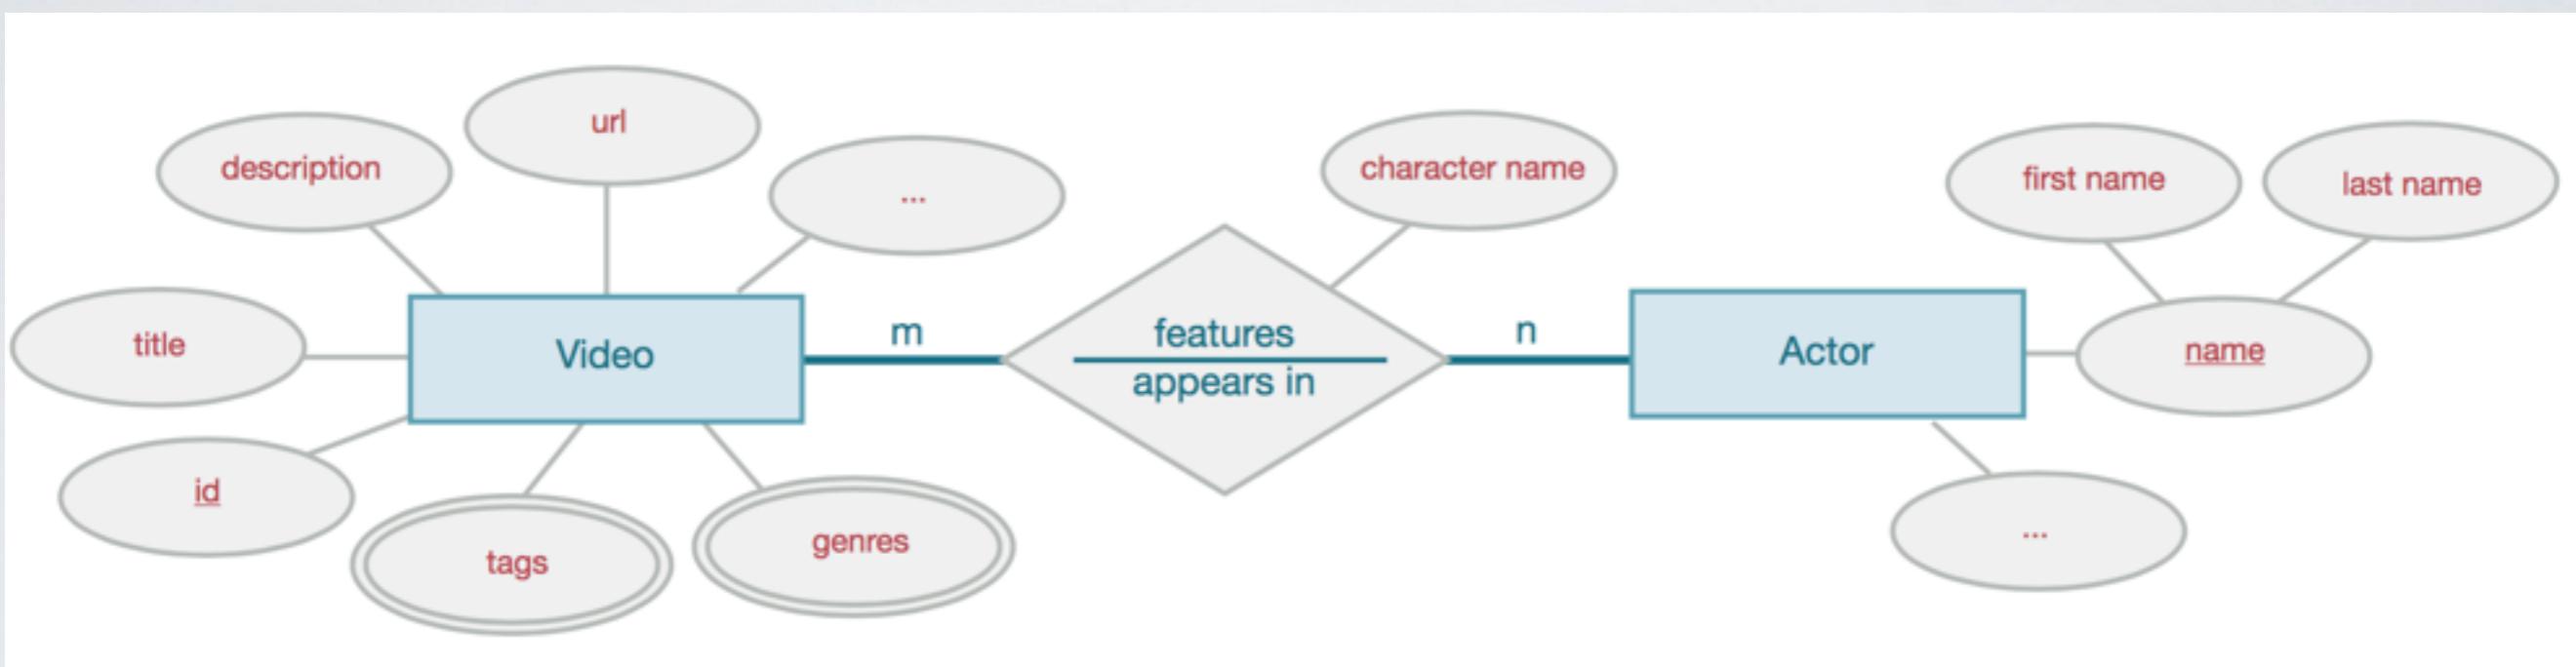
ENTITY-RELATIONSHIP (ER) MODEL



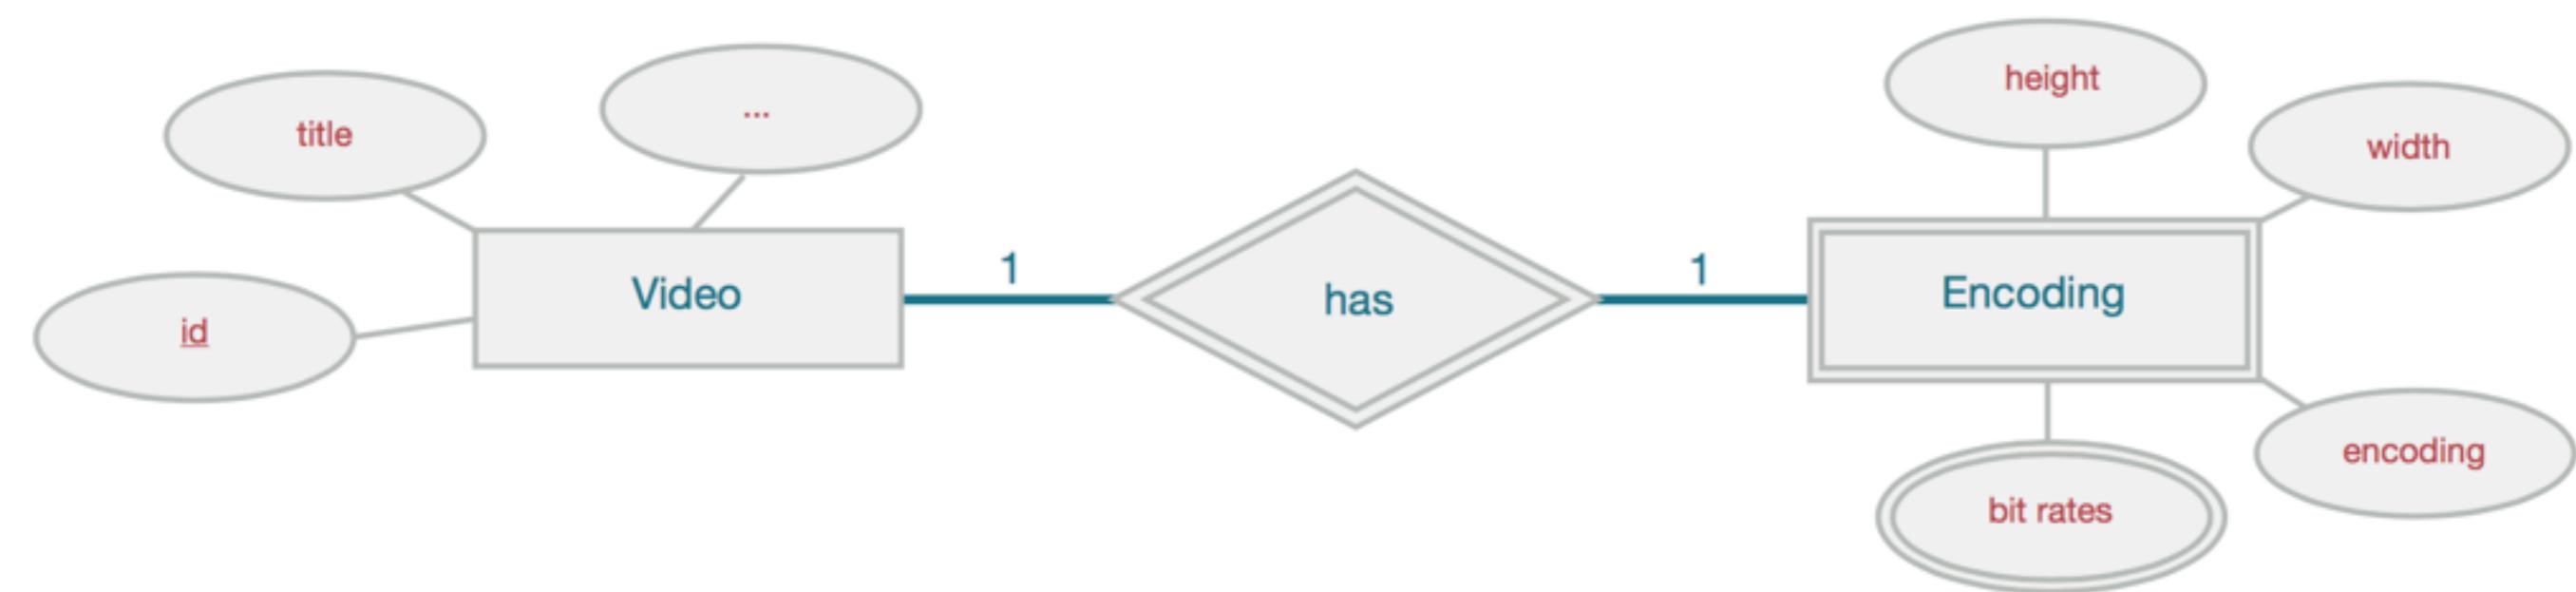
ENTITY TYPE



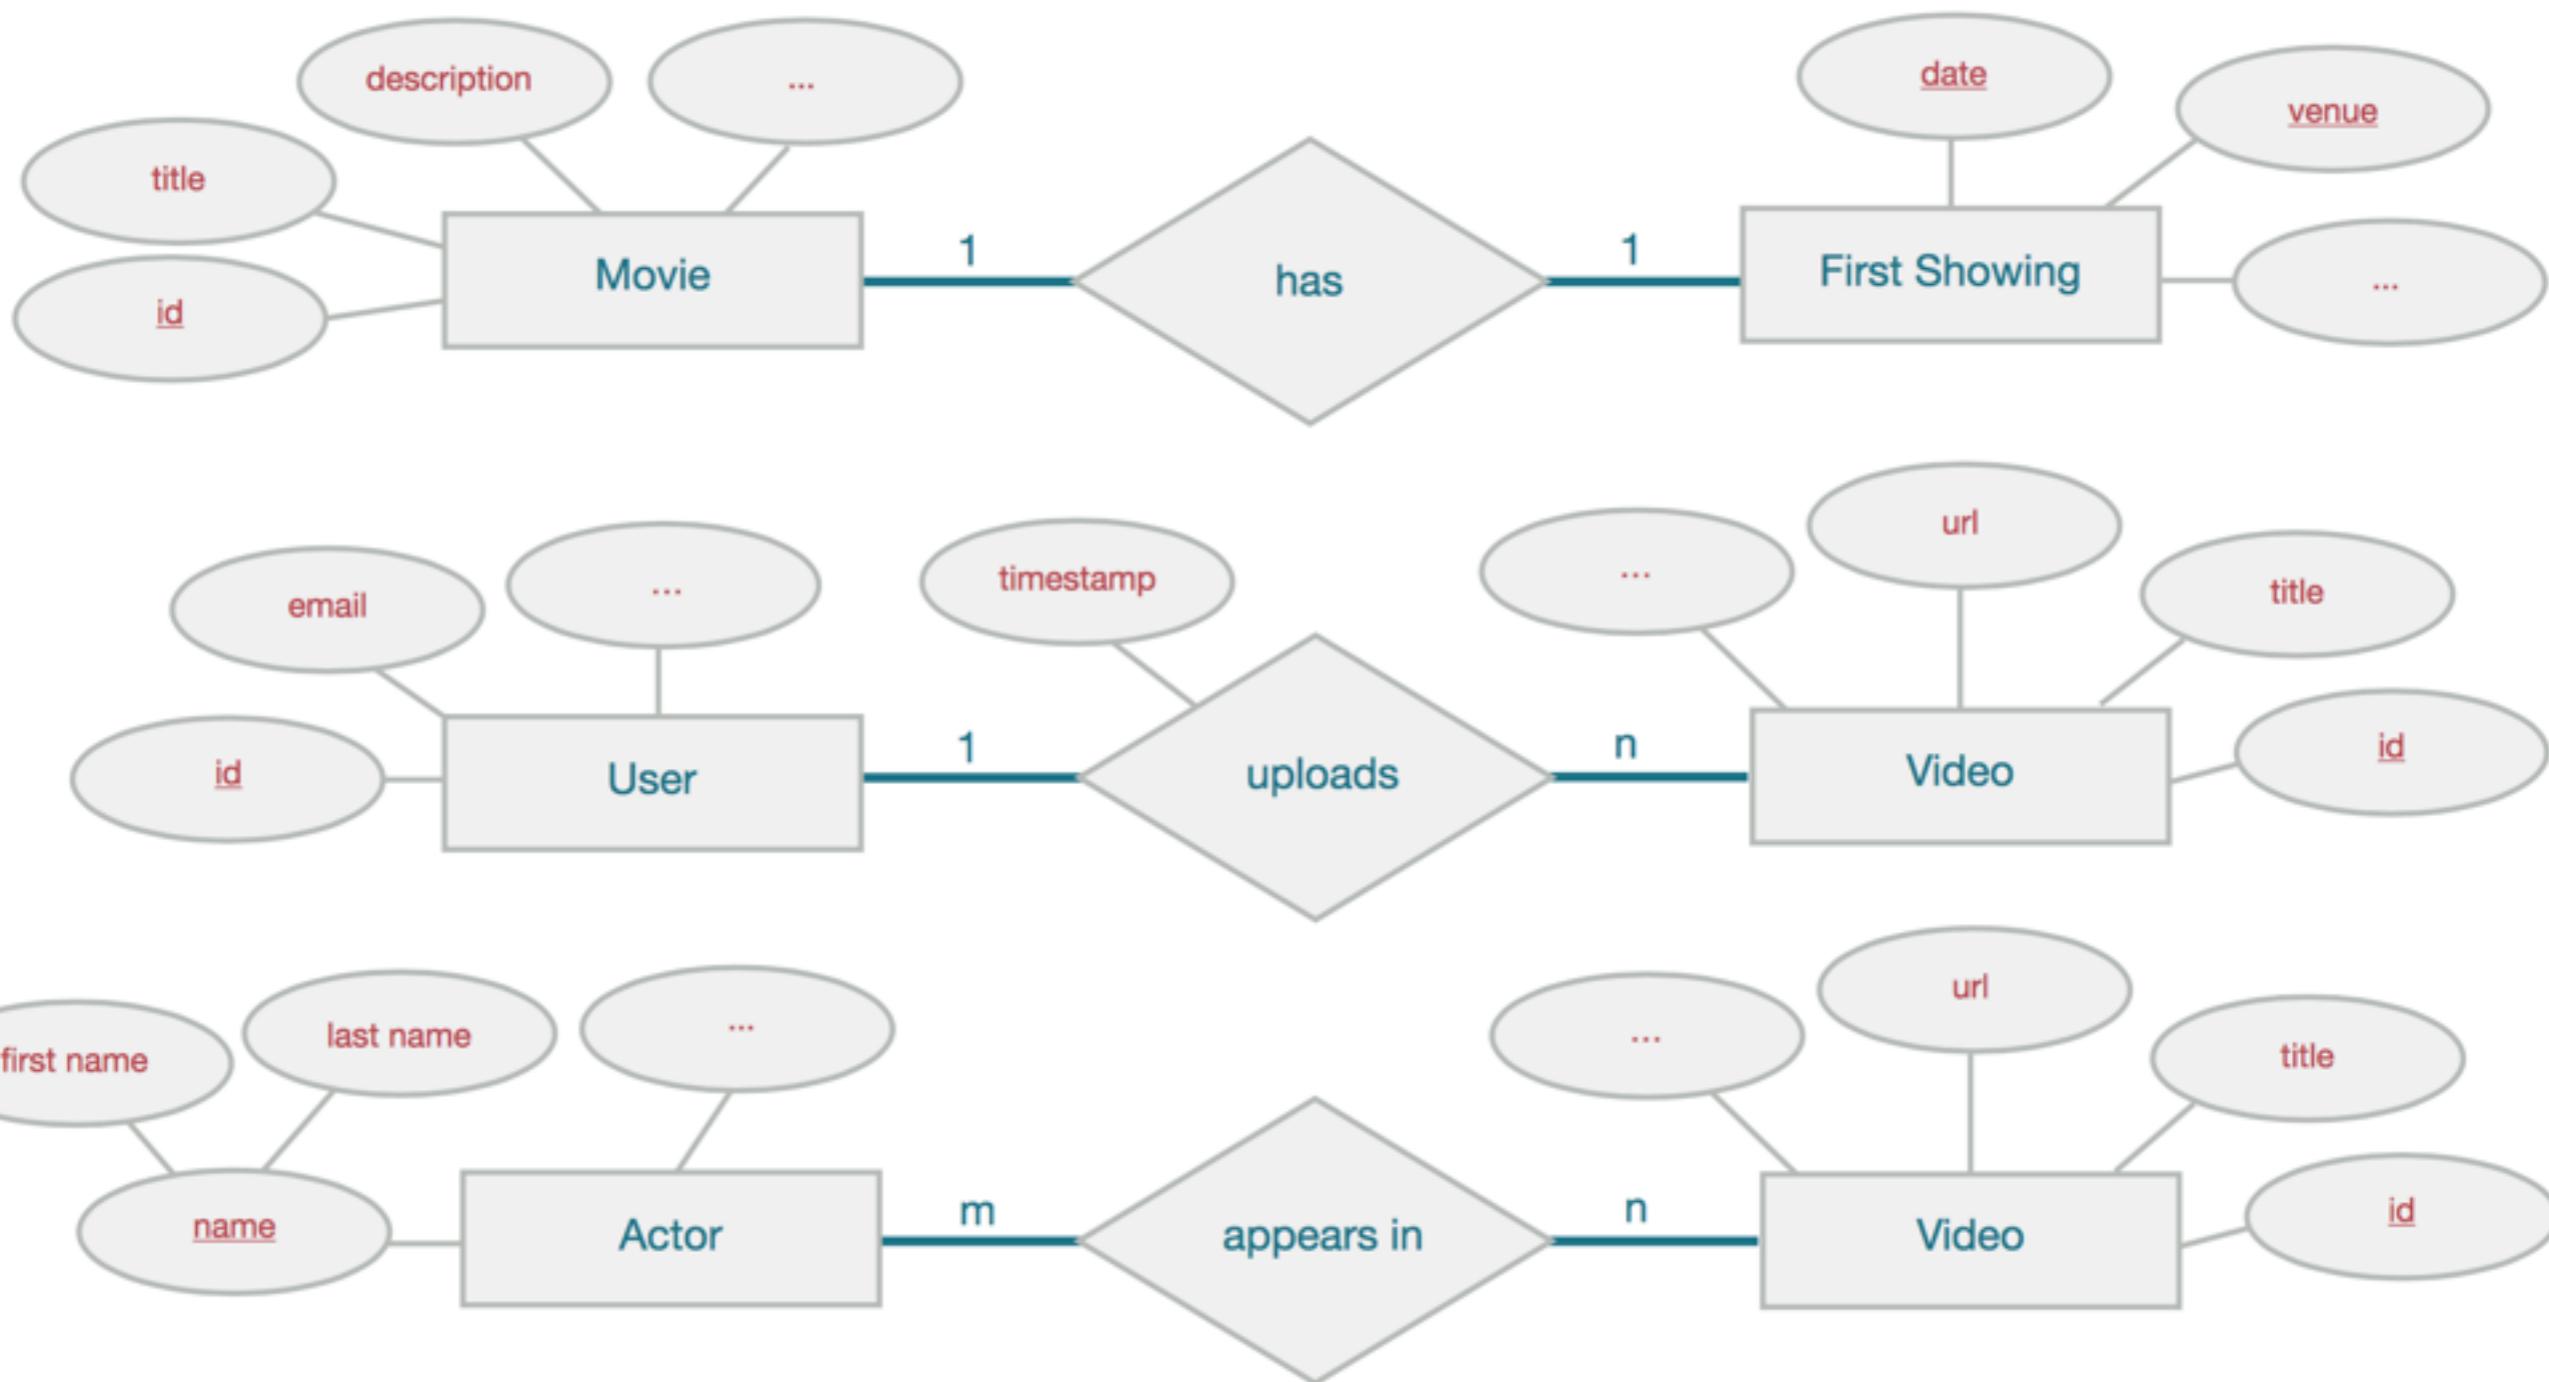
RELATIONSHIP TYPE



WEAK ENTITY TYPES



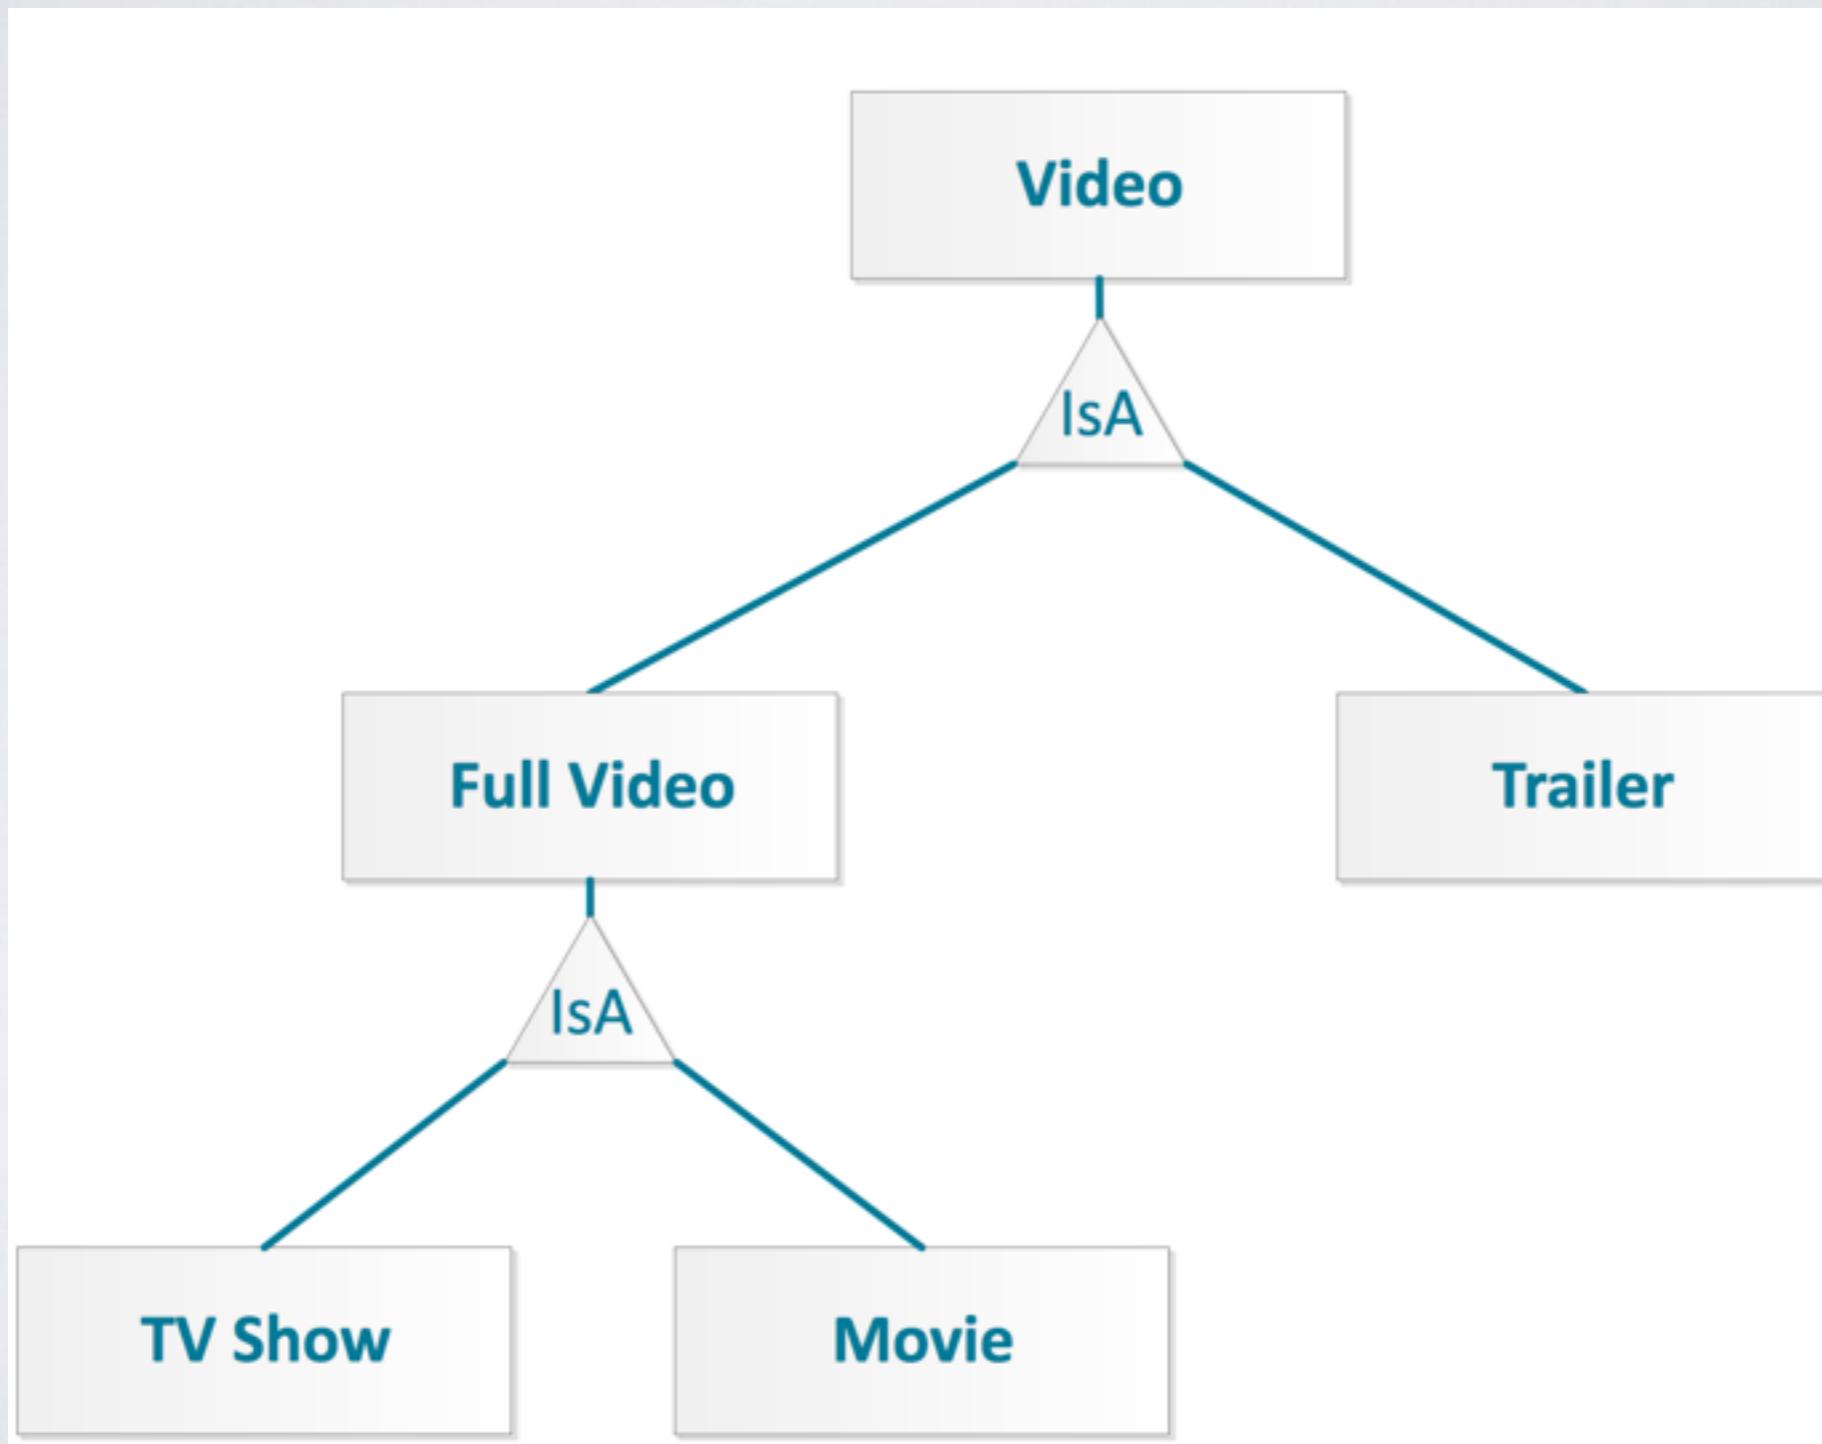
RELATIONSHIP KEYS



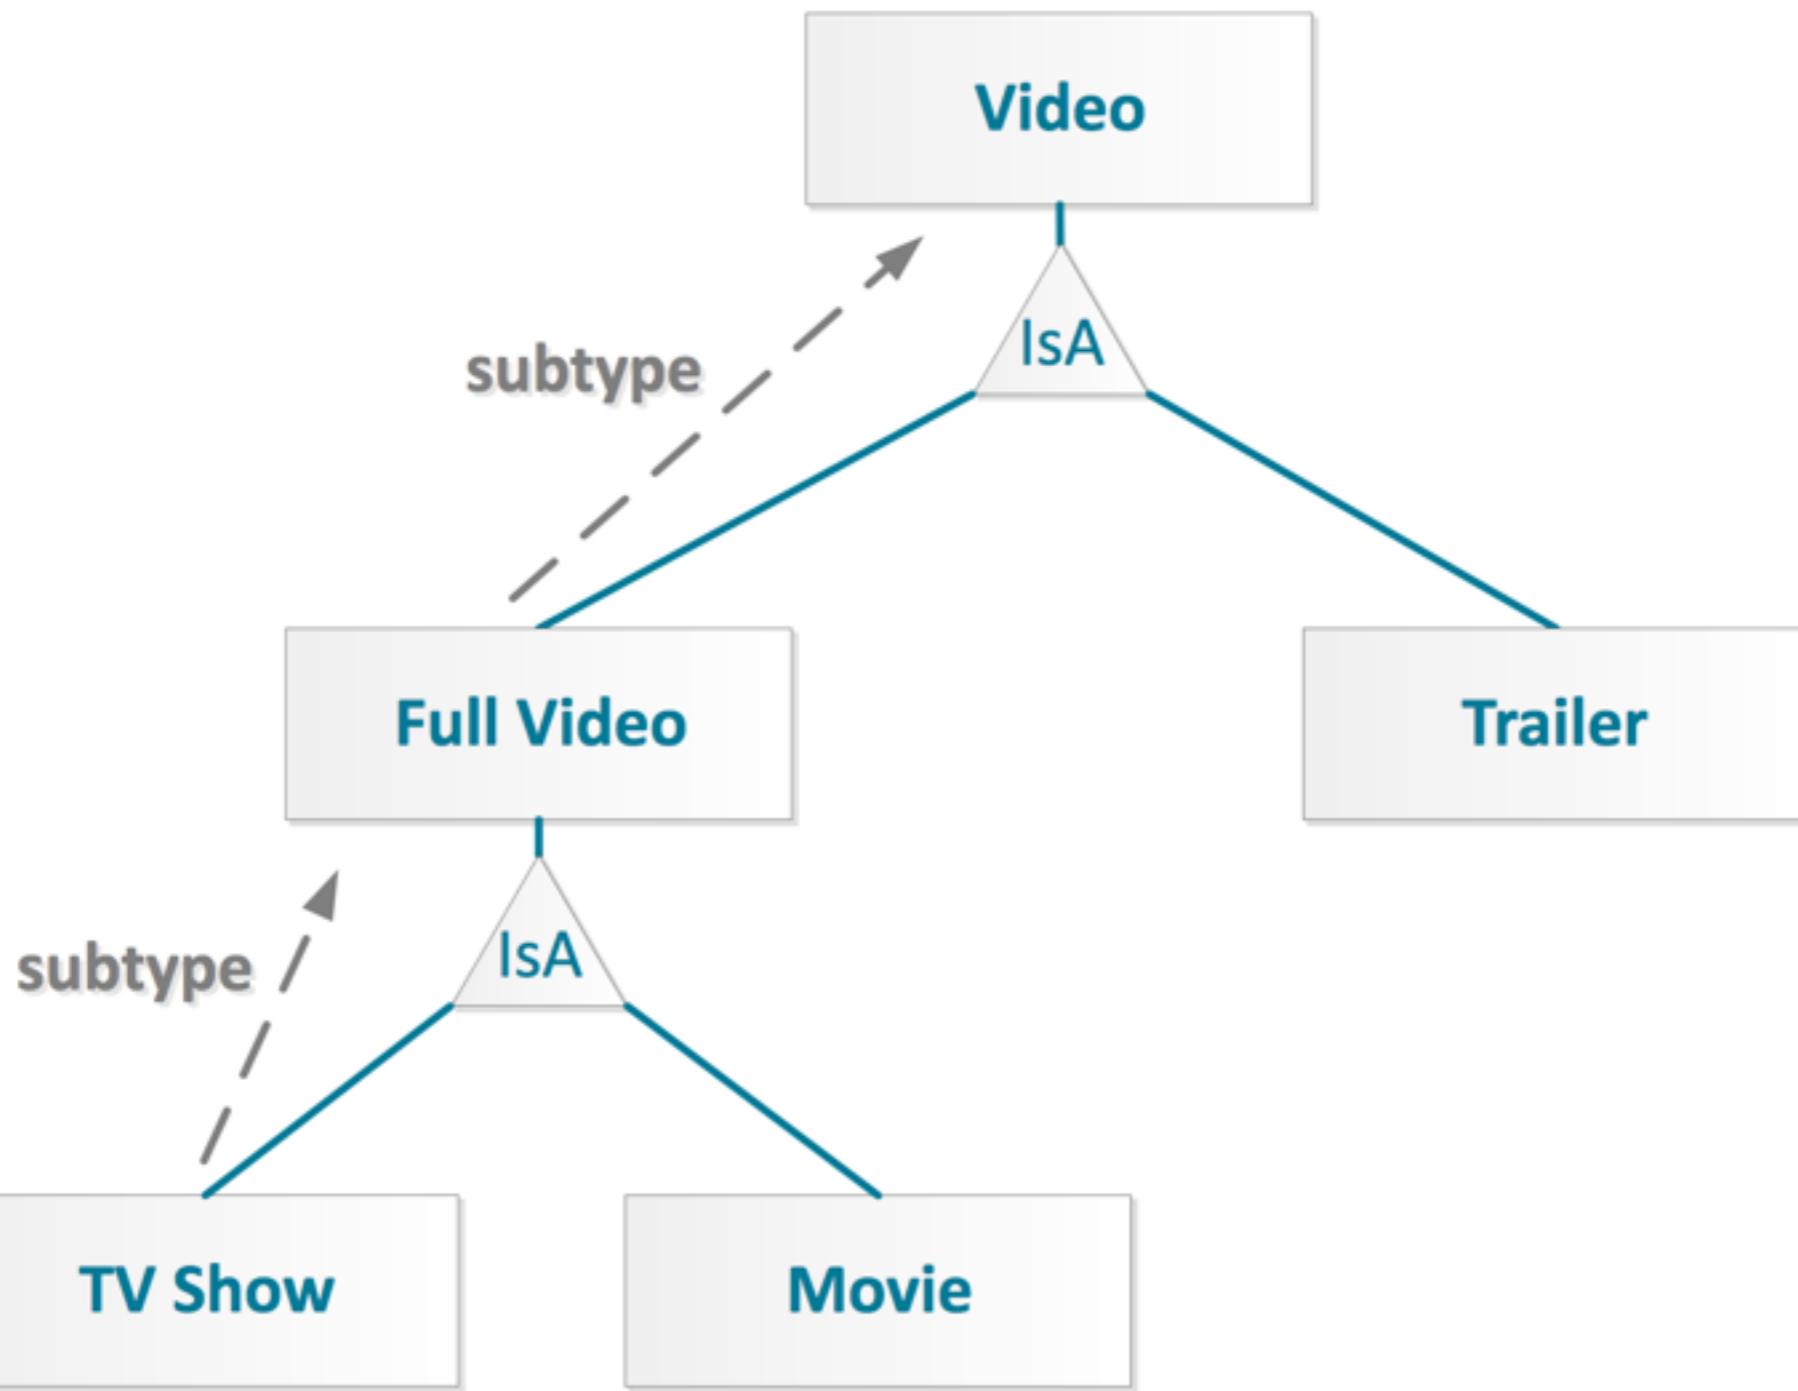
VIDEOS

- Video
- Full Video
- Trailer
- TV Show
- Movie

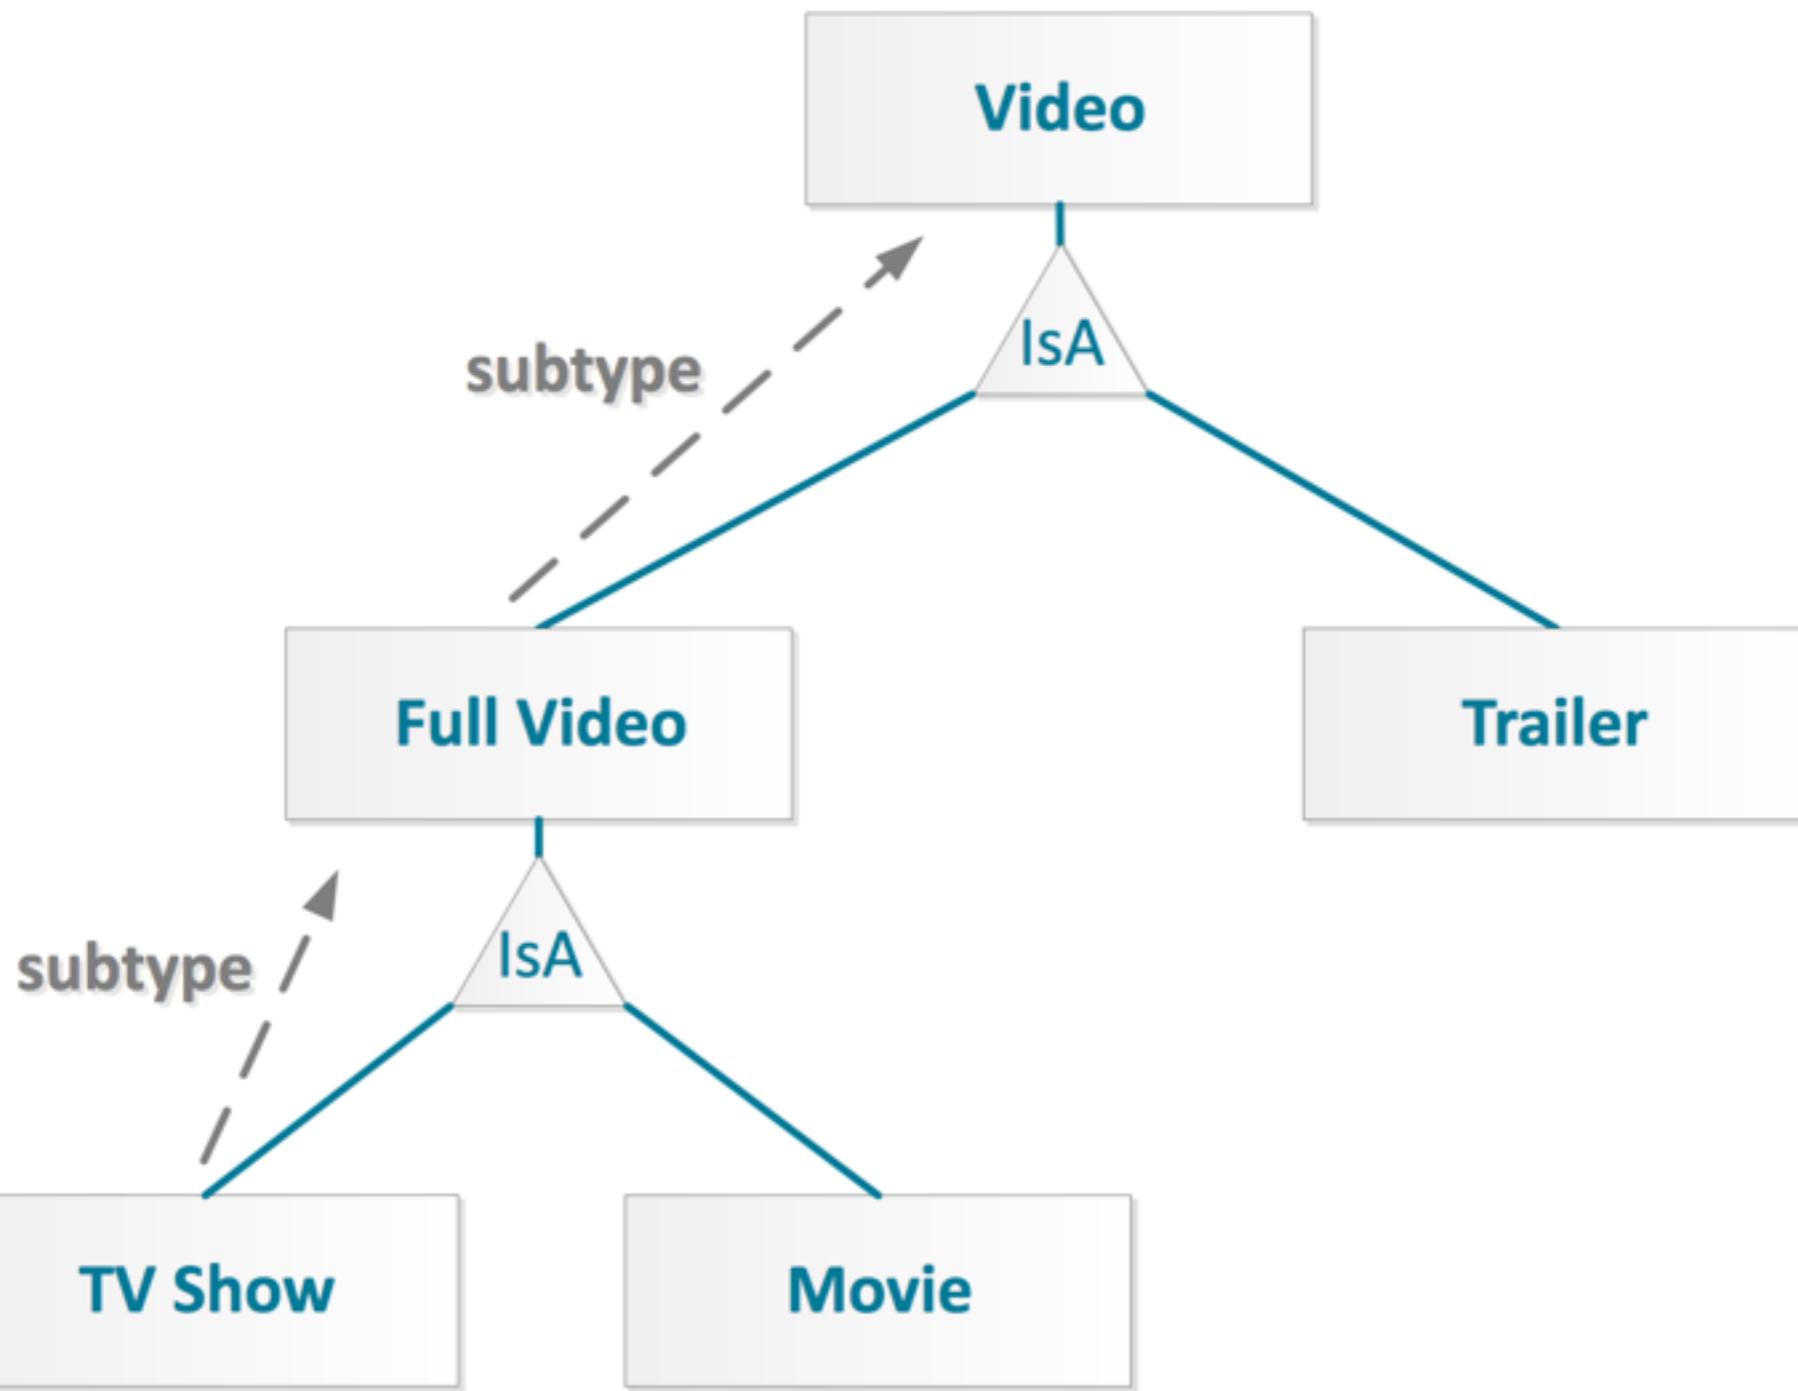
ENTITY TYPE HIERARCHY



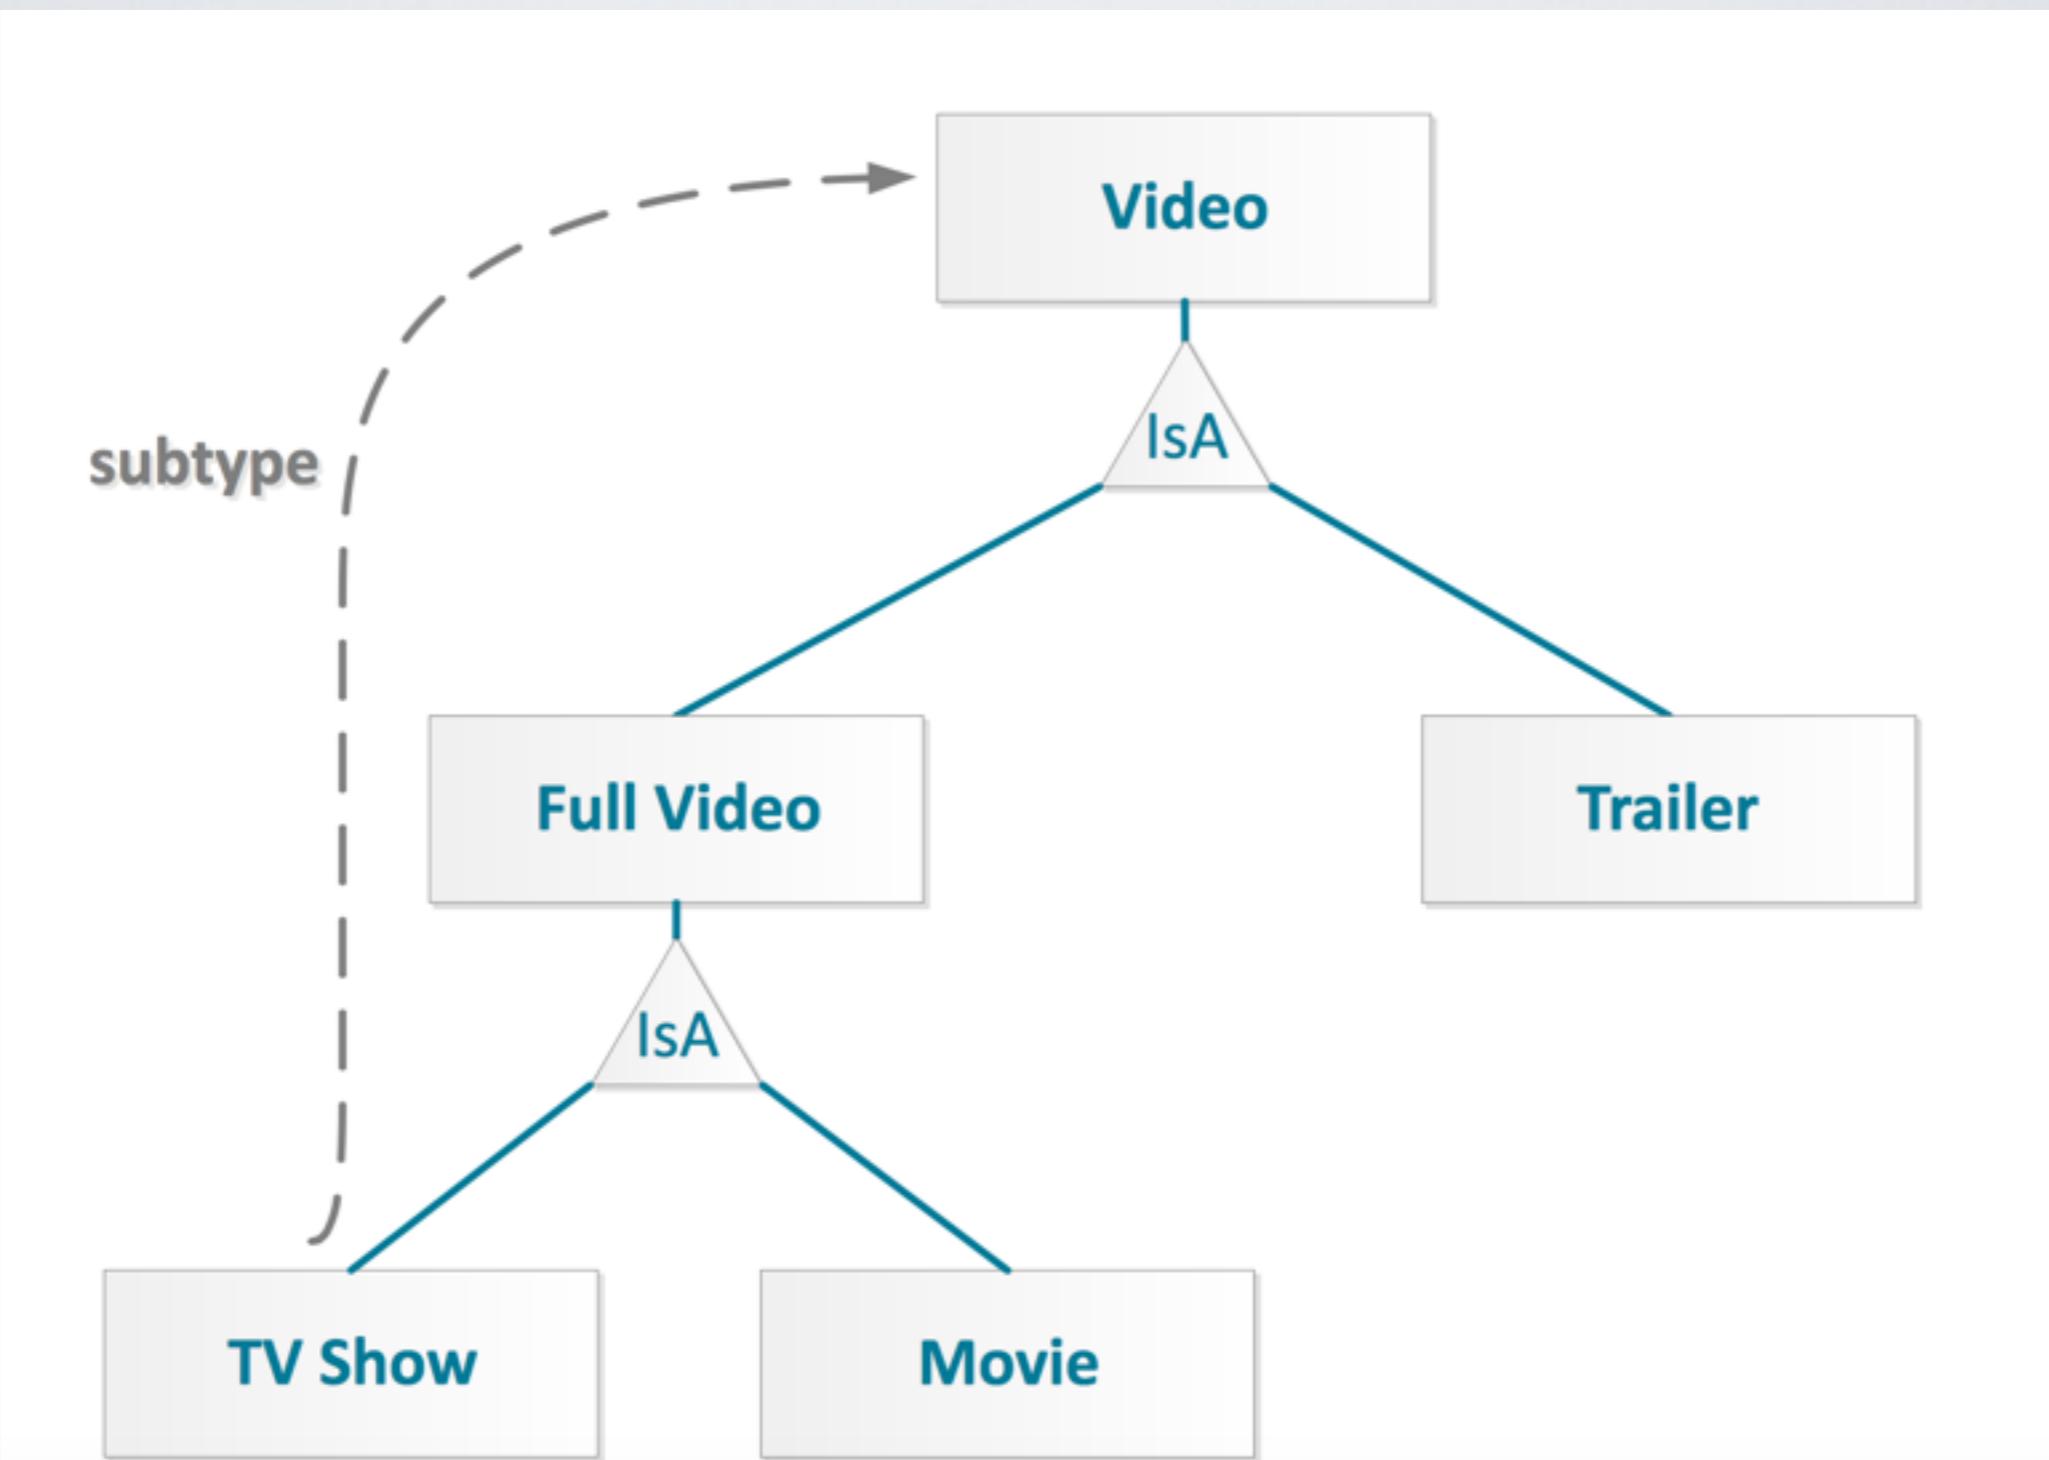
TRANSITIVITY

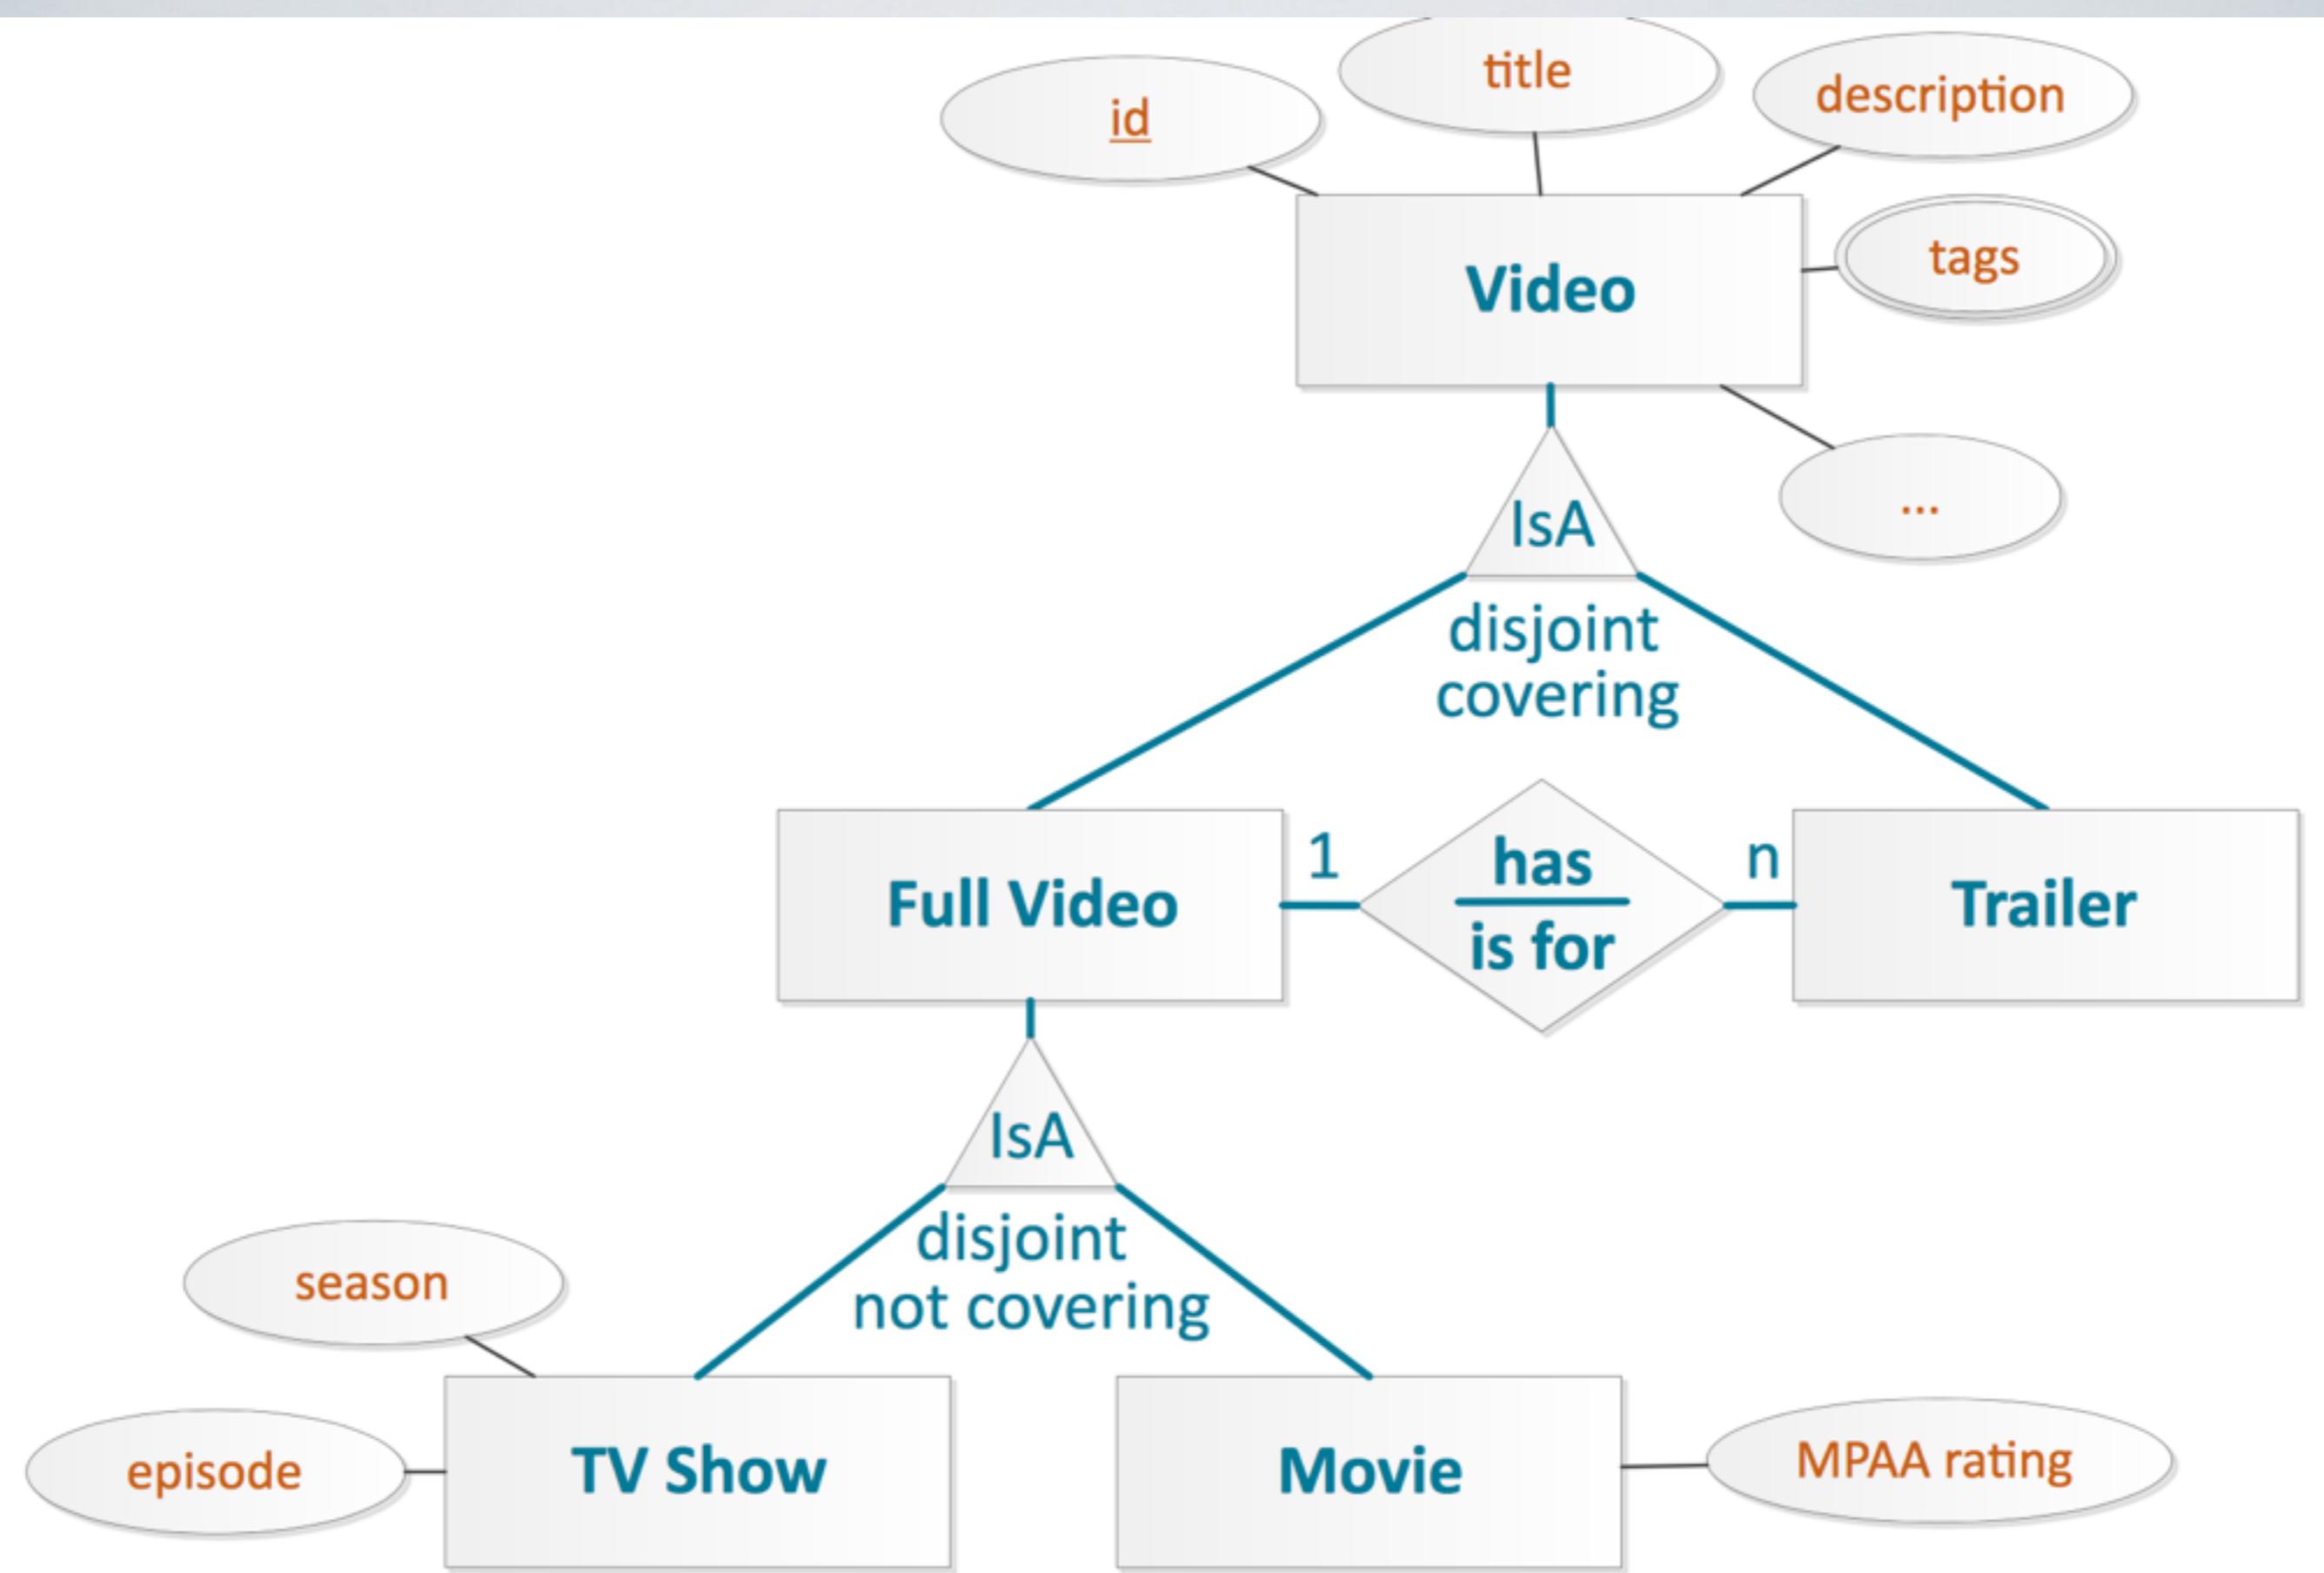


TRANSITIVITY

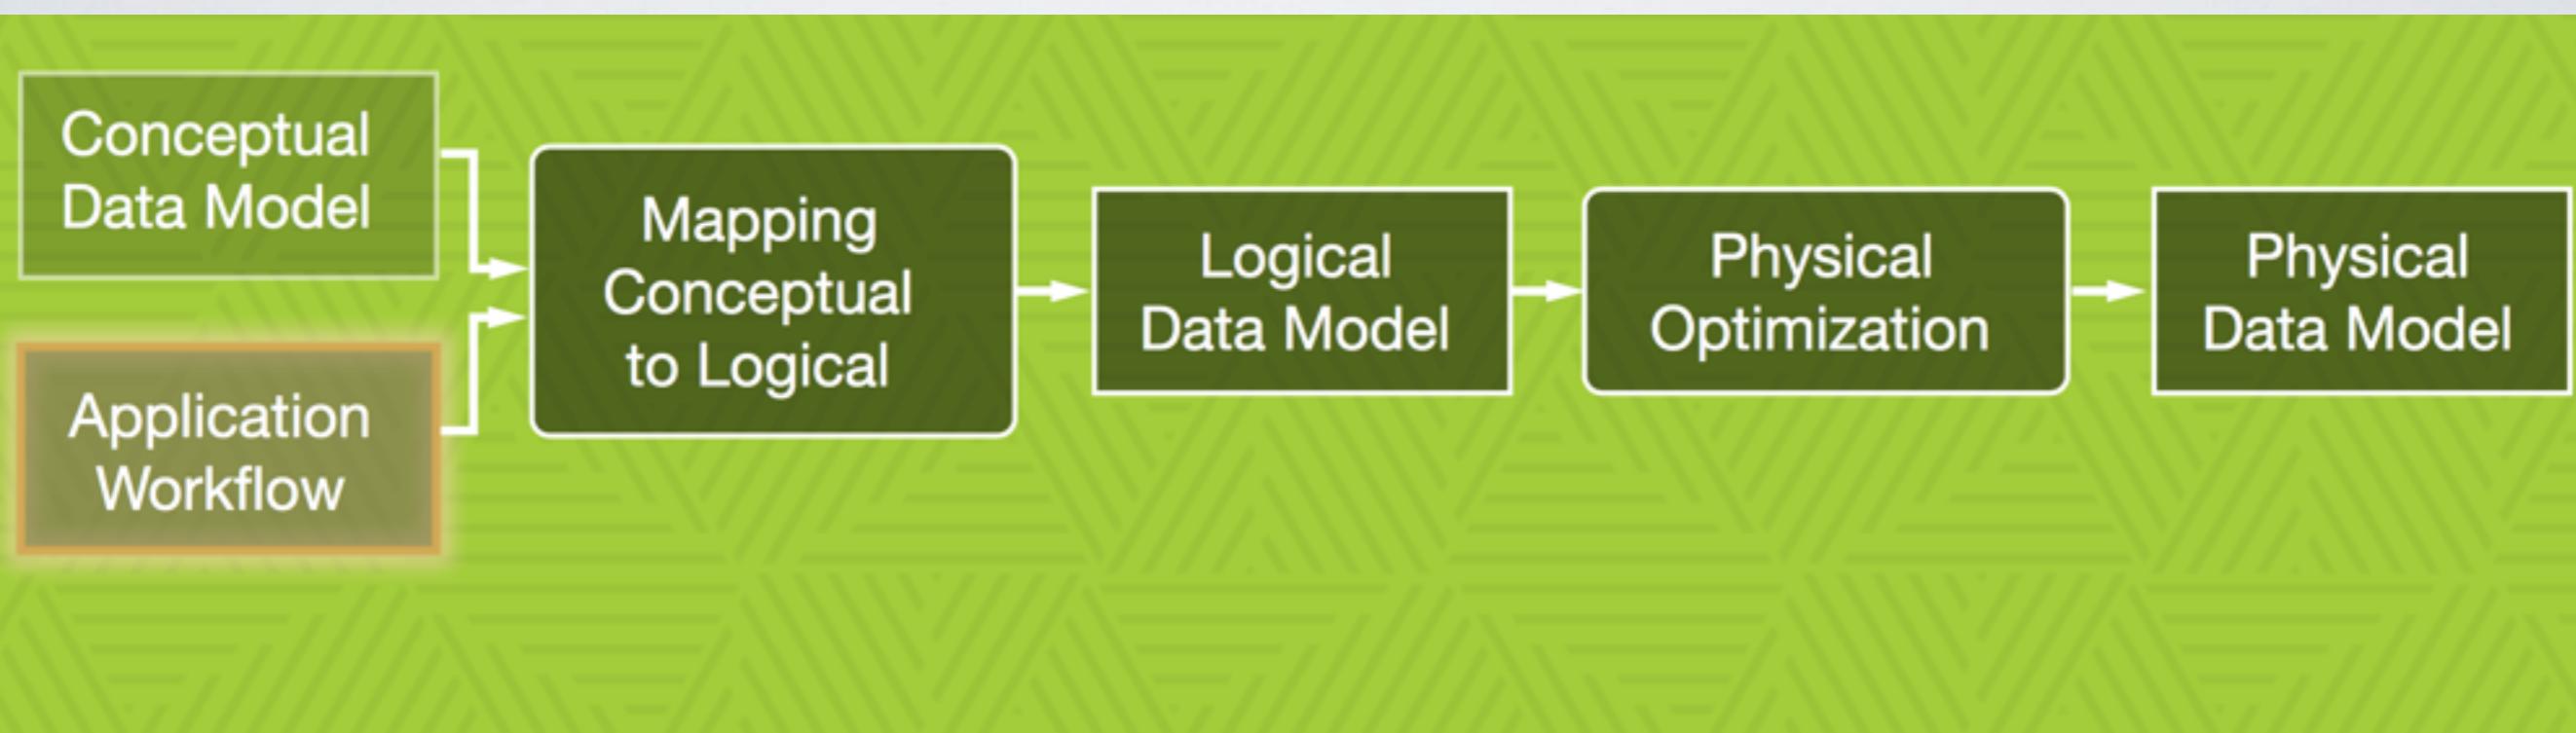


TRANSITIVITY



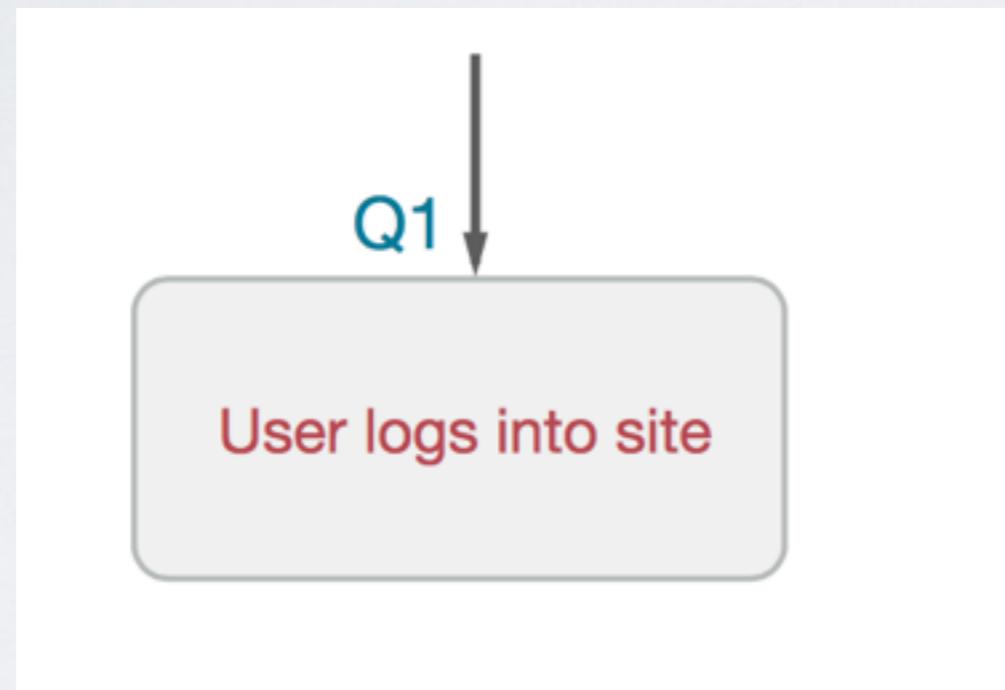


APPLICATION WORKFLOW MODEL



APPLICATION WORKFLOW MODEL

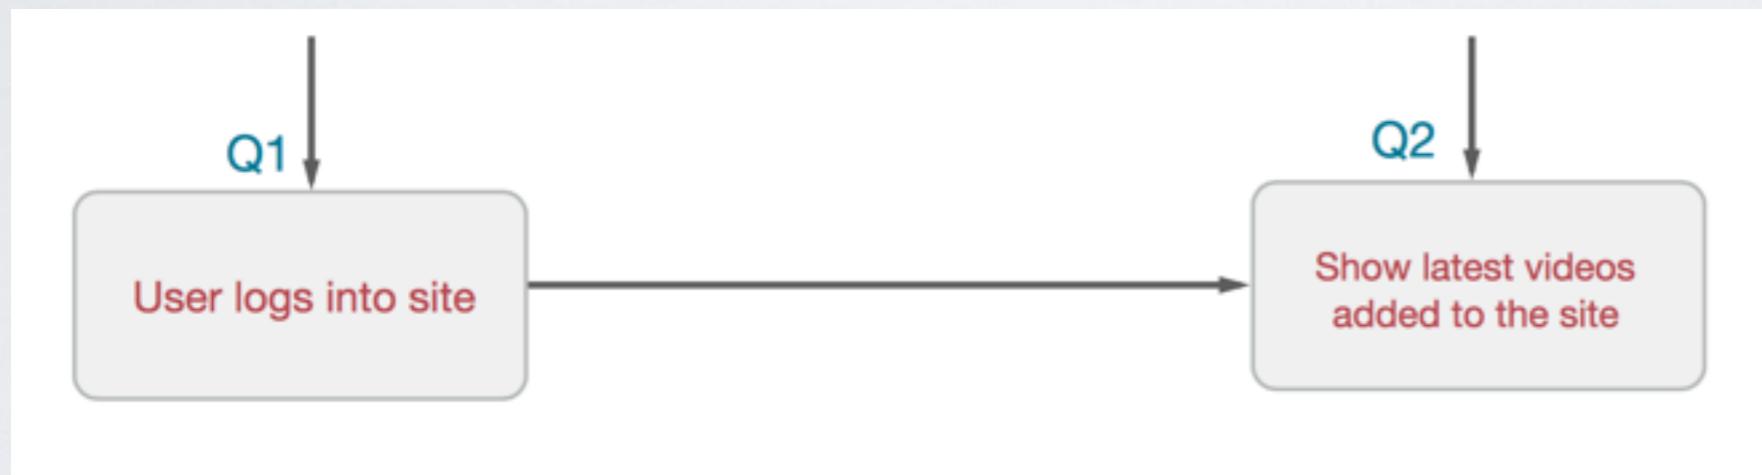
Task: Have a user login to a site



Access Patterns

Q1: Find a user with a specified email

Task: Show videos that were most recently uploaded

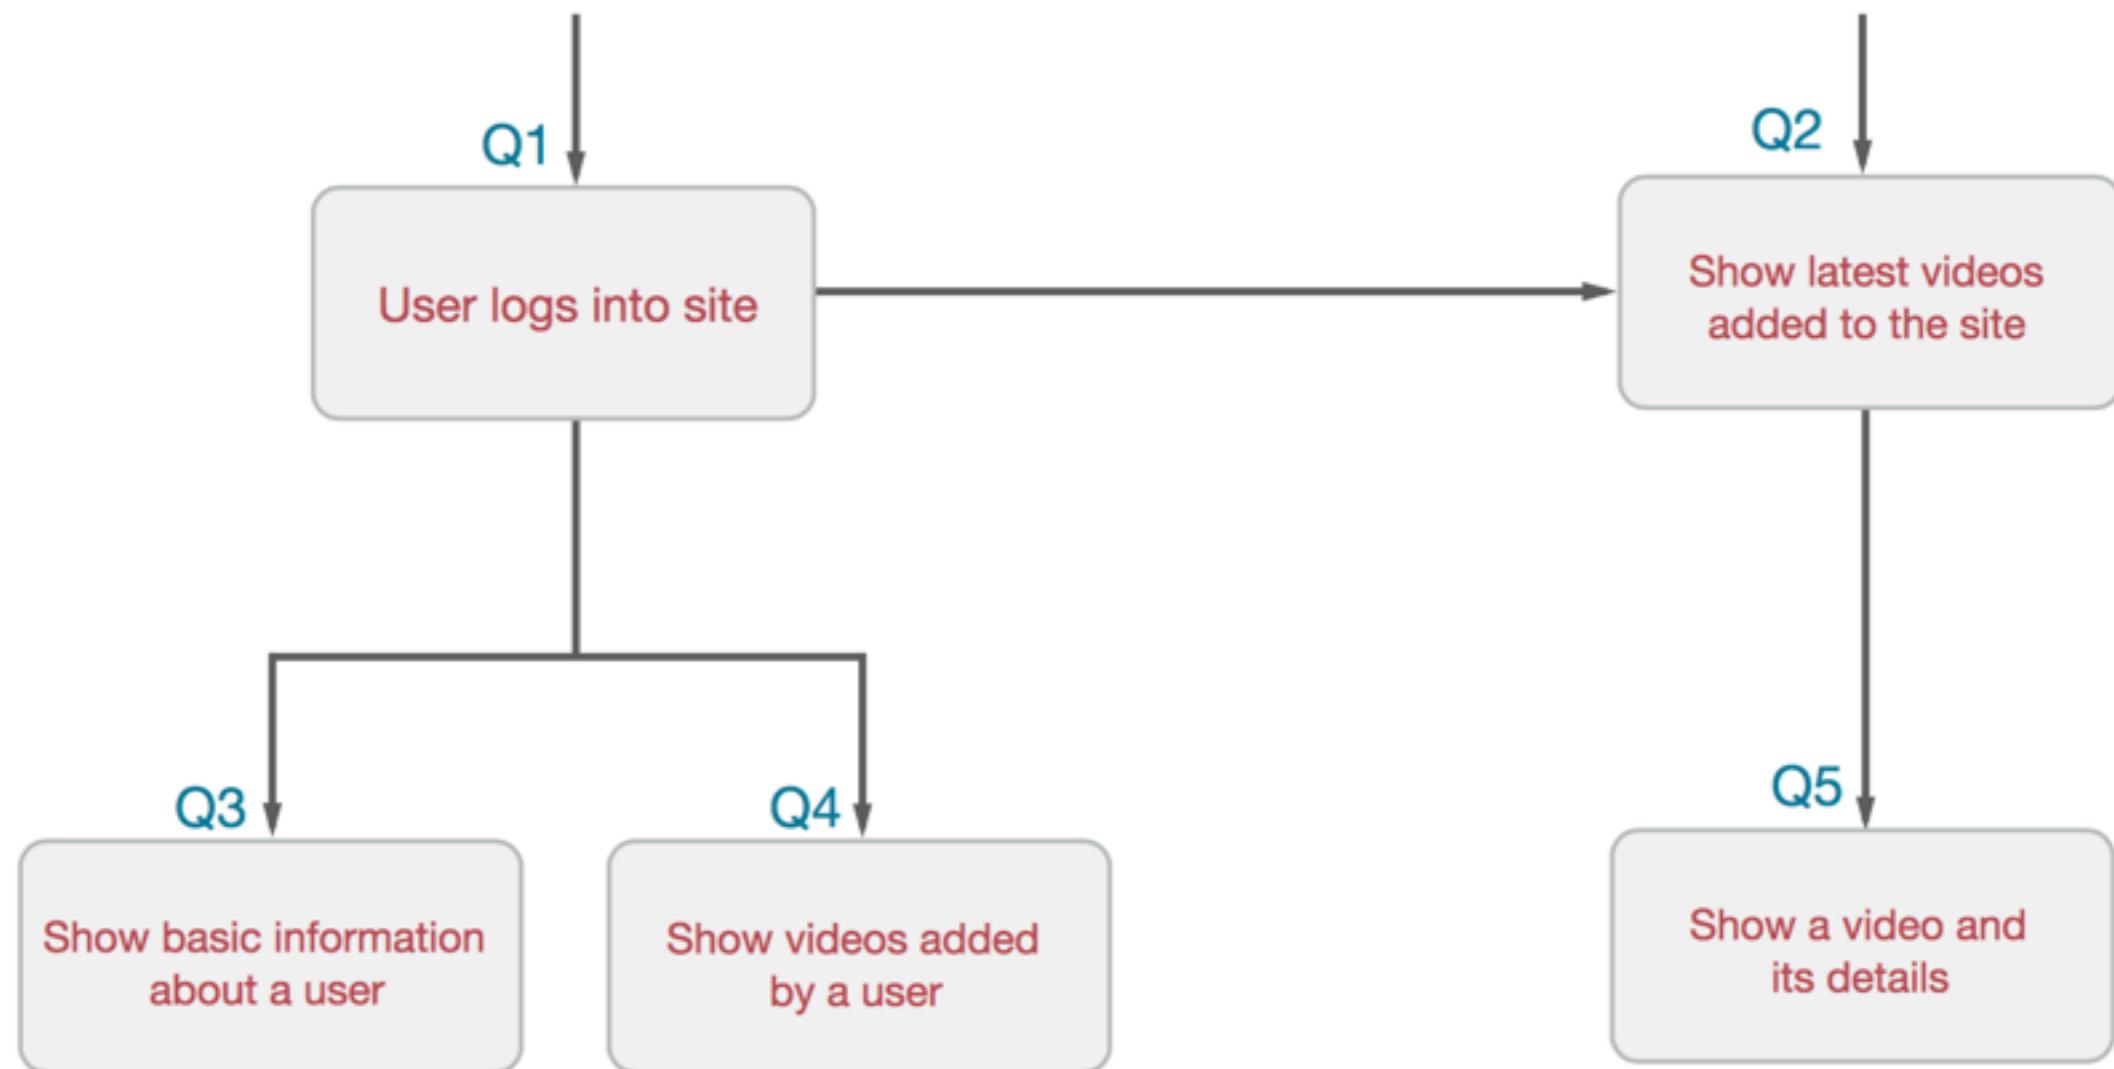


Access Patterns

Q1: Find a user with a specified email

Q2: Find most recently uploaded videos

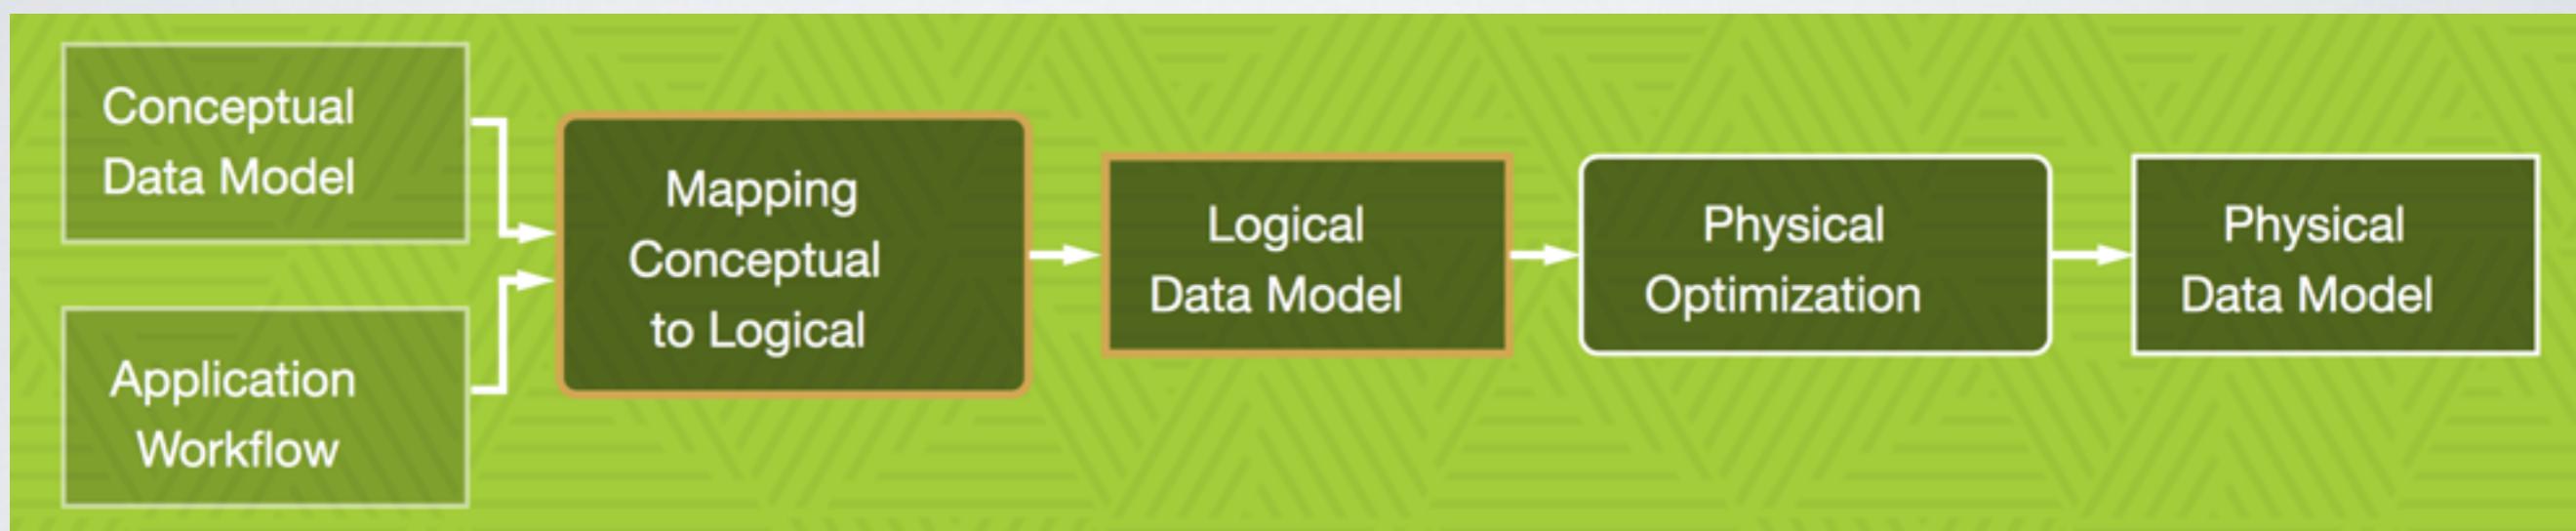
Task: Find a video that has a particular video id



ACCESS PATTERNS

- Q1: Find a user with a **specified email**
- Q2: Find most recently uploaded **videos**
- Q3: Find a **user with a specified id**
- Q4: Find videos uploaded by a **user with a known id**(show most recently uploaded videos first)
- Q5: Find a video with a **specified video id**

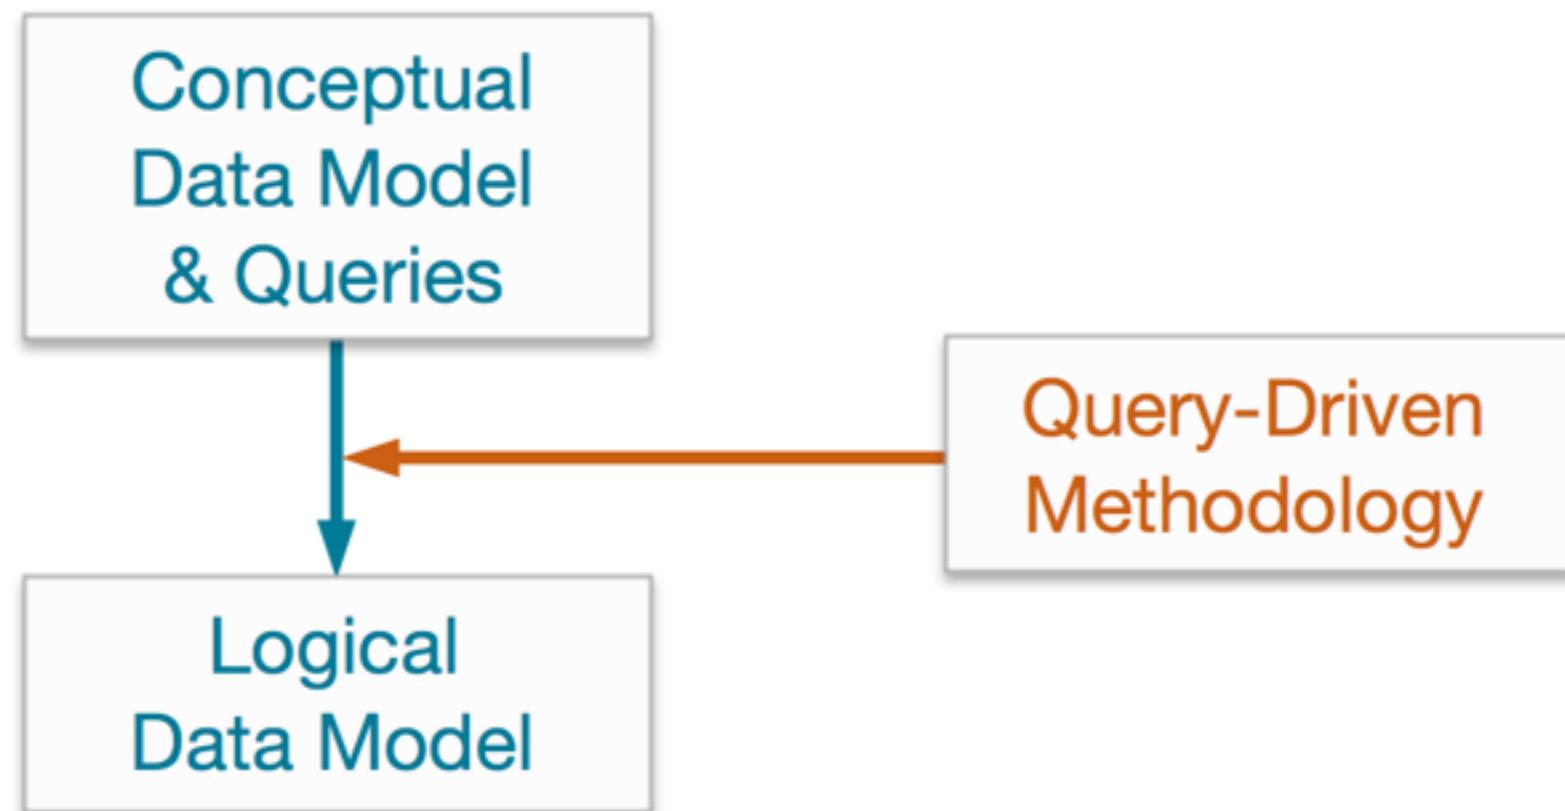
MAPPING CONCEPTUAL TO LOGICAL



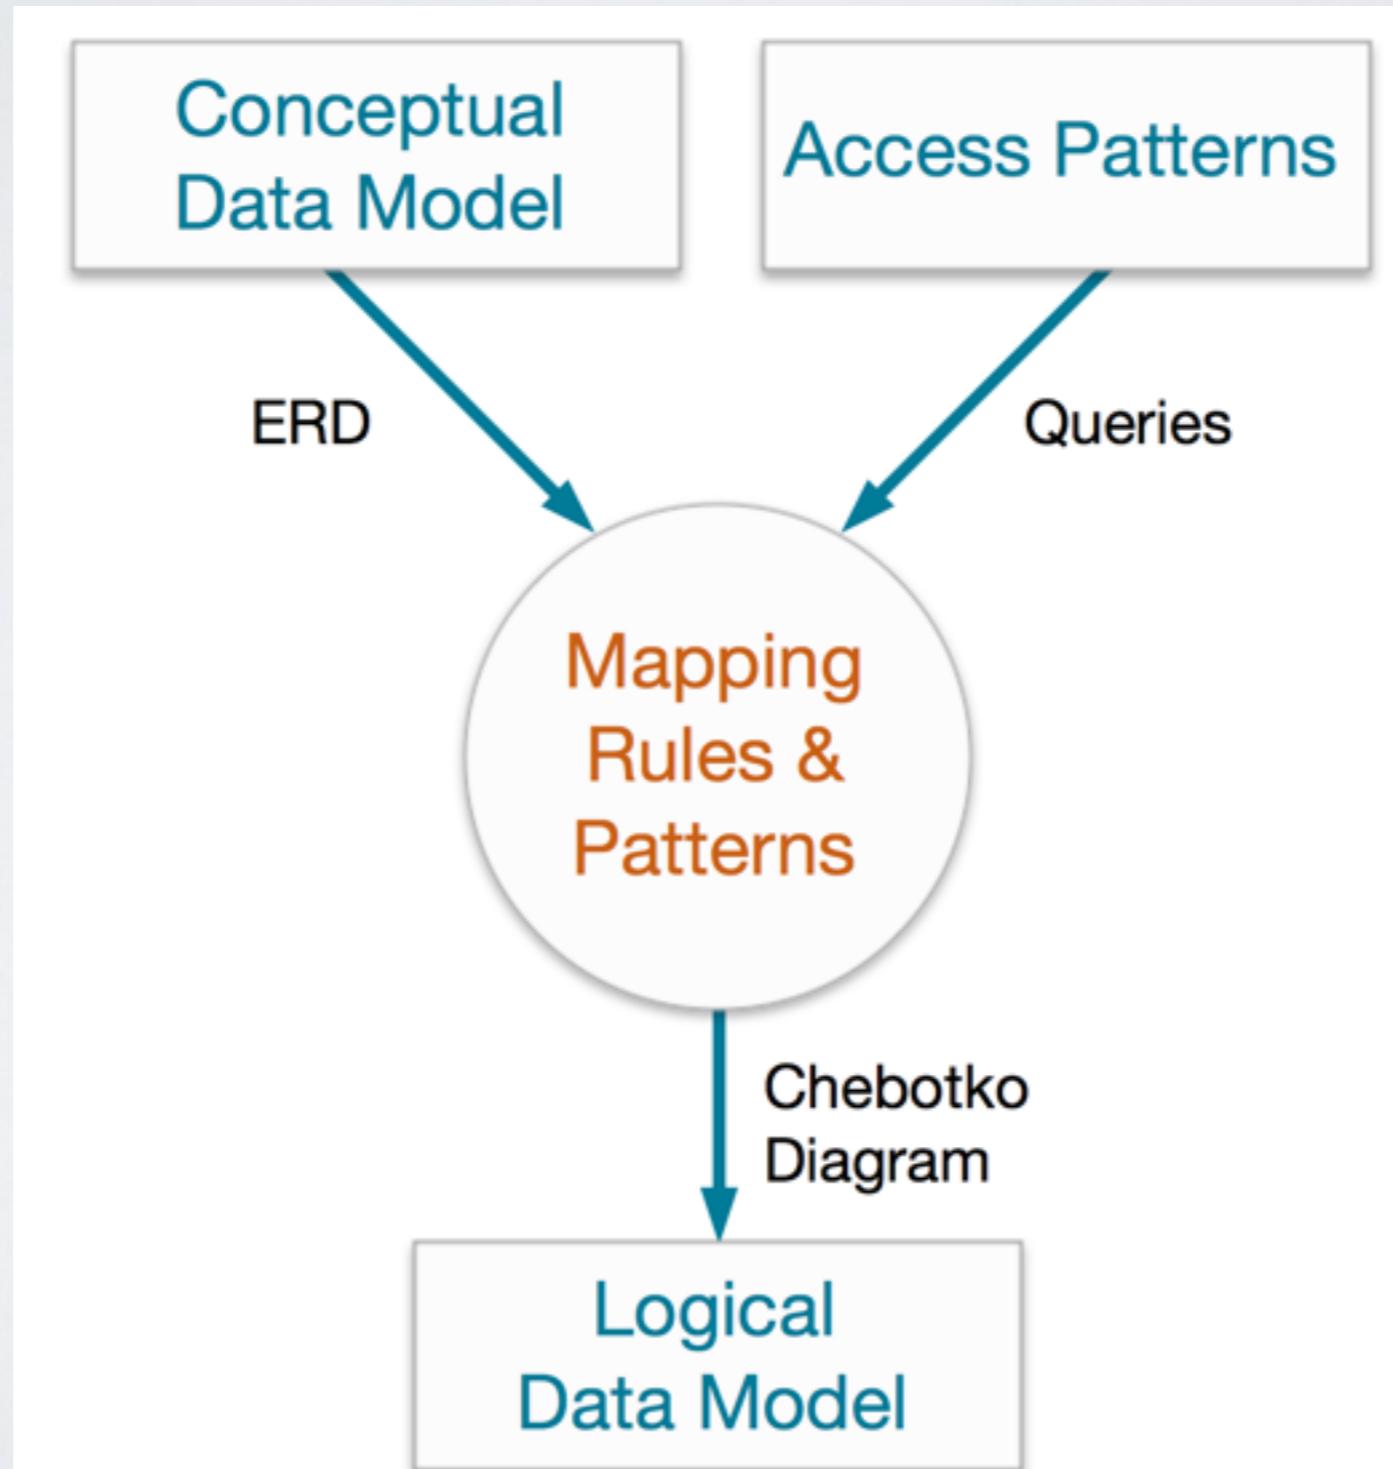
DATA MODELING METHODOLOGY

Process to design a logical model

- Uses a top-down approach
- Can be algorithmically defined
- Effective in the long run—vs. dark art

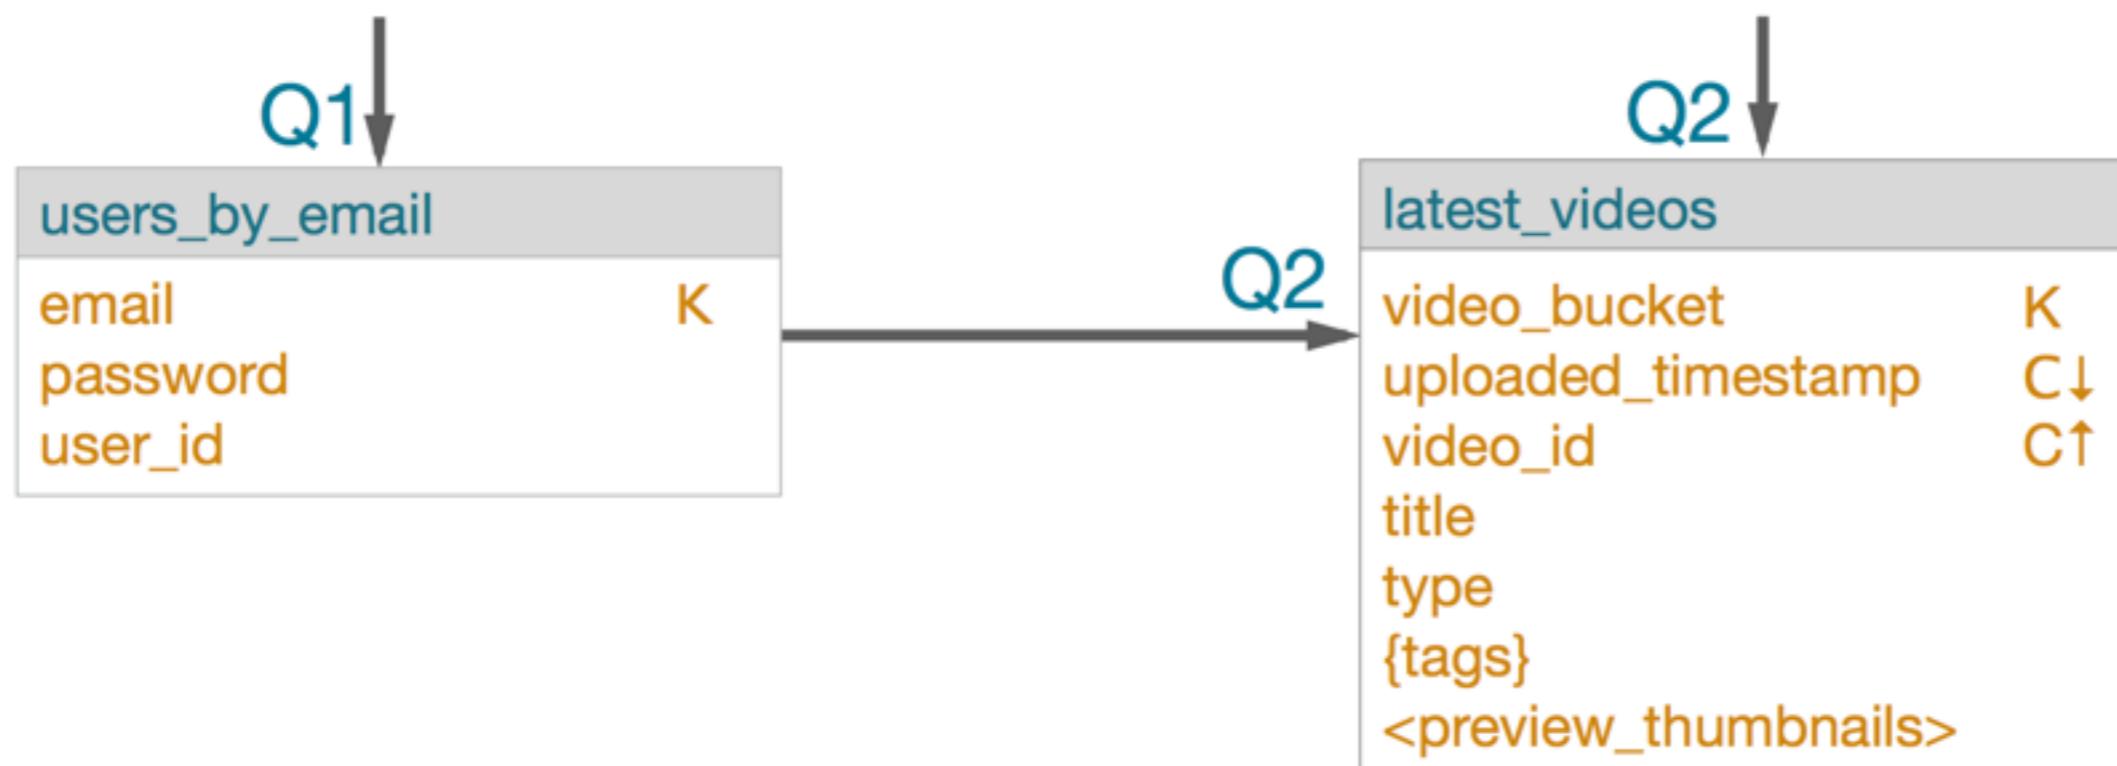


QUERY-DRIVEN DATA MODELING



CHEBOTKO DIAGRAM

- Visual diagram for Cassandra tables and access patterns



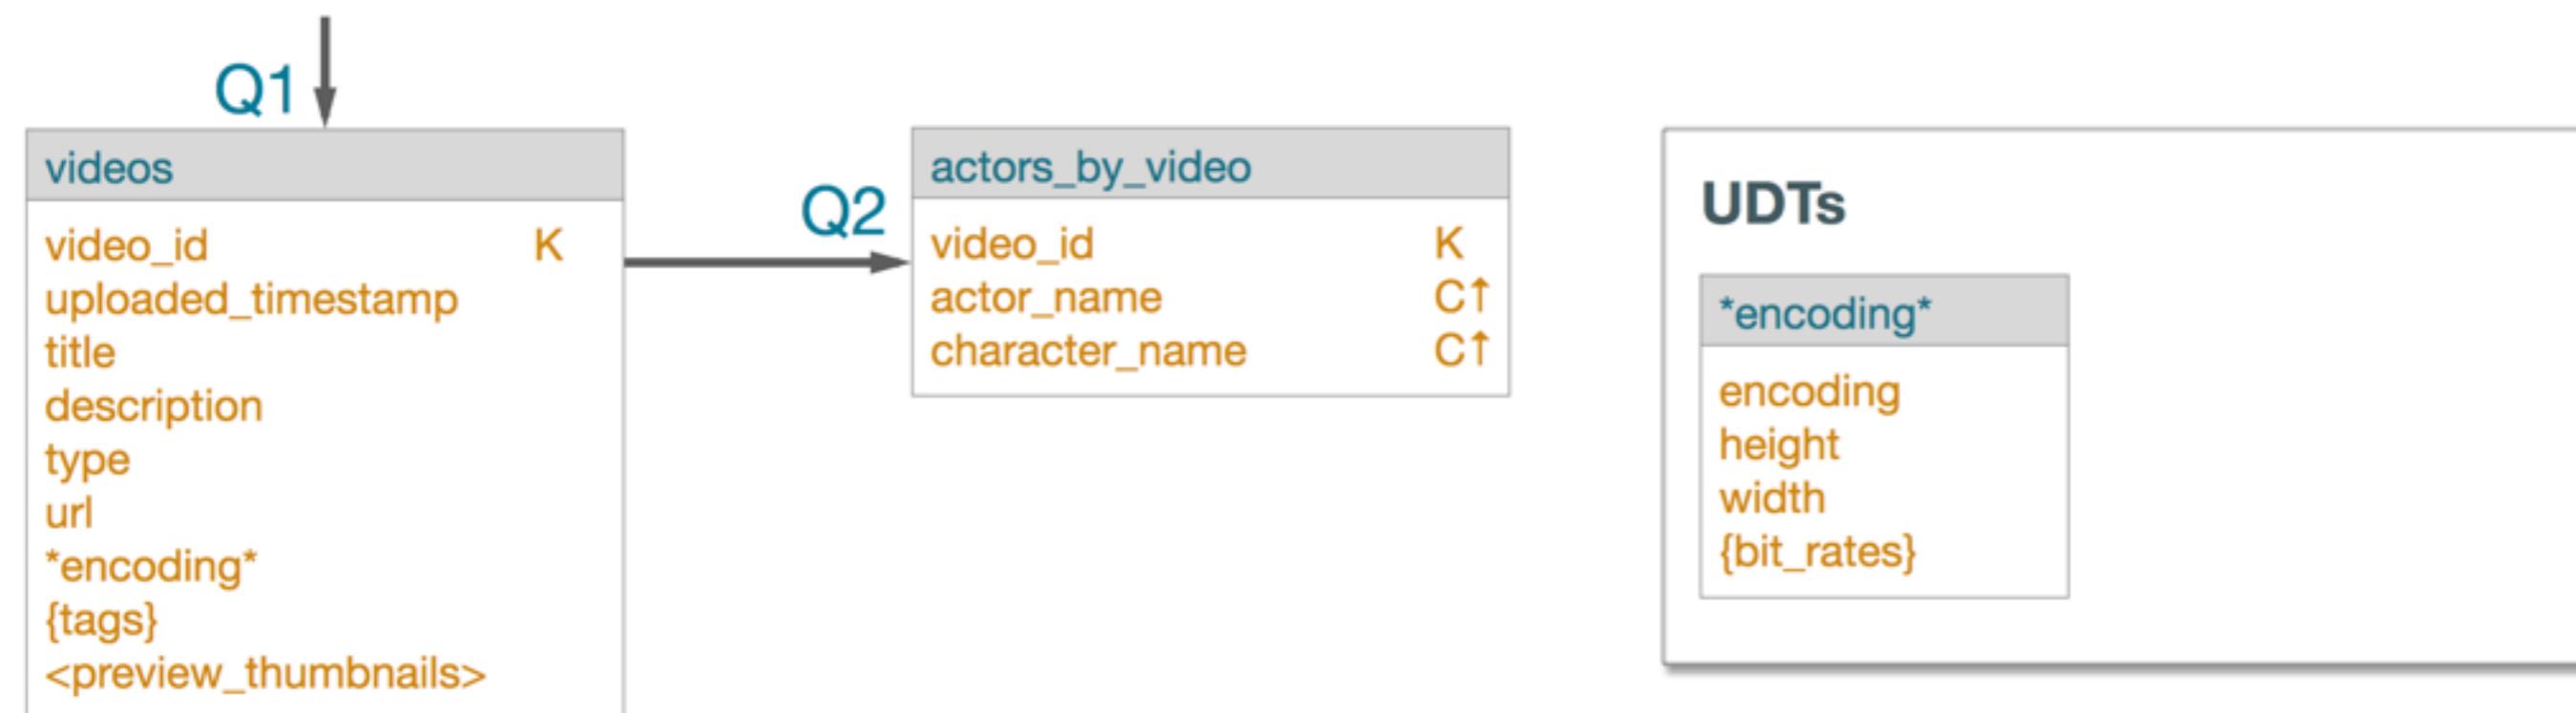
ACCESS PATTERNS

Q1: Find a user with a **specified email**

Q2: Find most recently uploaded videos

CHEBOTKO DIAGRAMS

- Graphical representation of Cassandra database schema design
- Documents the logical and physical data model



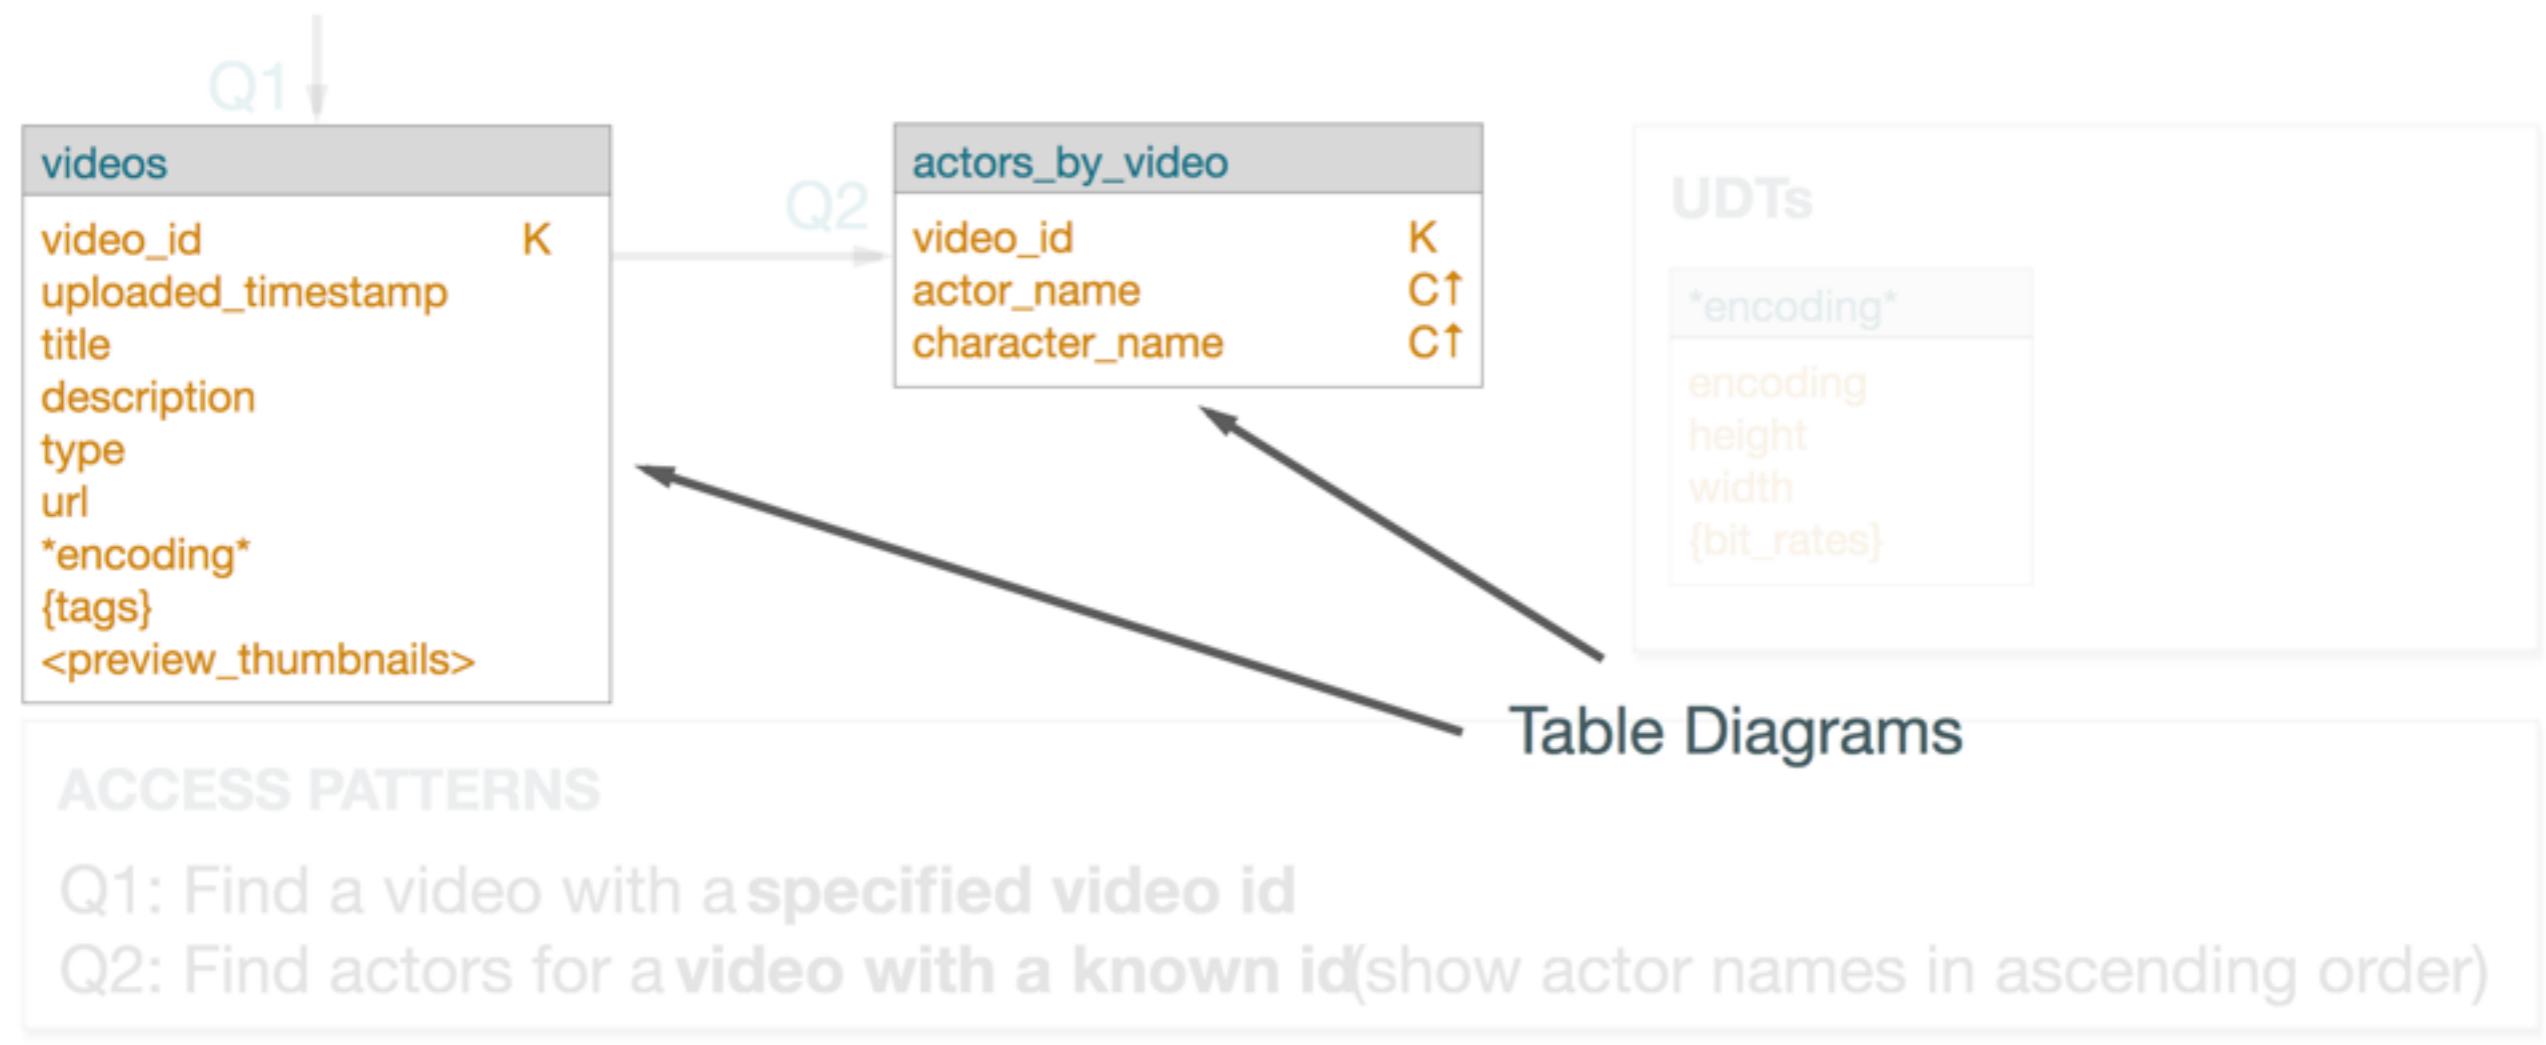
ACCESS PATTERNS

Q1: Find a video with a **specified video id**

Q2: Find actors for a **video with a known id** (show actor names in ascending order)

CHEBOTKO DIAGRAMS

- Graphical representation of Cassandra database schema design
- Documents the logical and physical data model



CHEBOTKO DIAGRAMS

- Graphical representation of Cassandra database schema design
- Documents the logical and physical data model



ACCESS PATTERNS

Q1: Find a video with a specified video id

Q2: Find actors for a video with a known id(show actor names in ascending order)

CHEBOTKO DIAGRAMS

- Graphical representation of Cassandra database schema design
- Documents the logical and physical data model



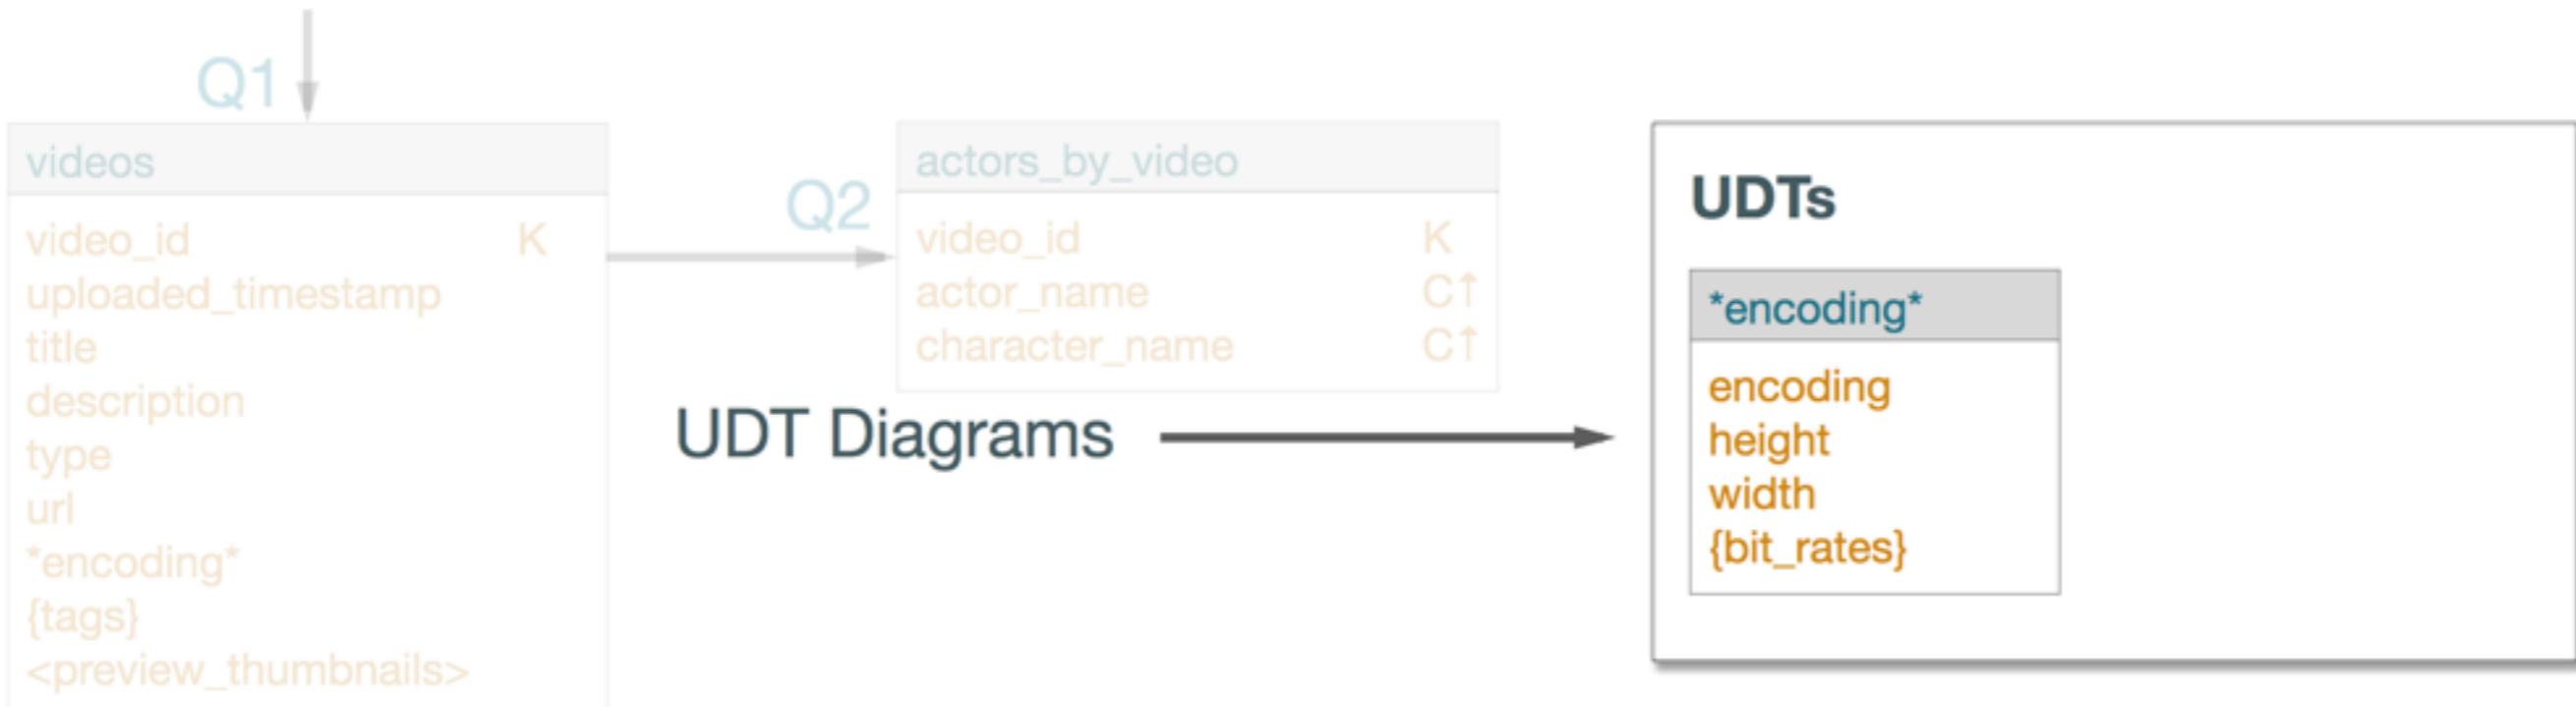
ACCESS PATTERNS

Q1: Find a video with a **specified video id**

Q2: Find actors for a **video with a known id** (show actor names in ascending order)

CHEBOTKO DIAGRAMS

- Graphical representation of Cassandra database schema design
- Documents the logical and physical data model



ACCESS PATTERNS

Q1: Find a video with a specified video id

Q2: Find actors for a video with a known id(show actor names in ascending order)

CHEBOTKO DIAGRAM NOTATION

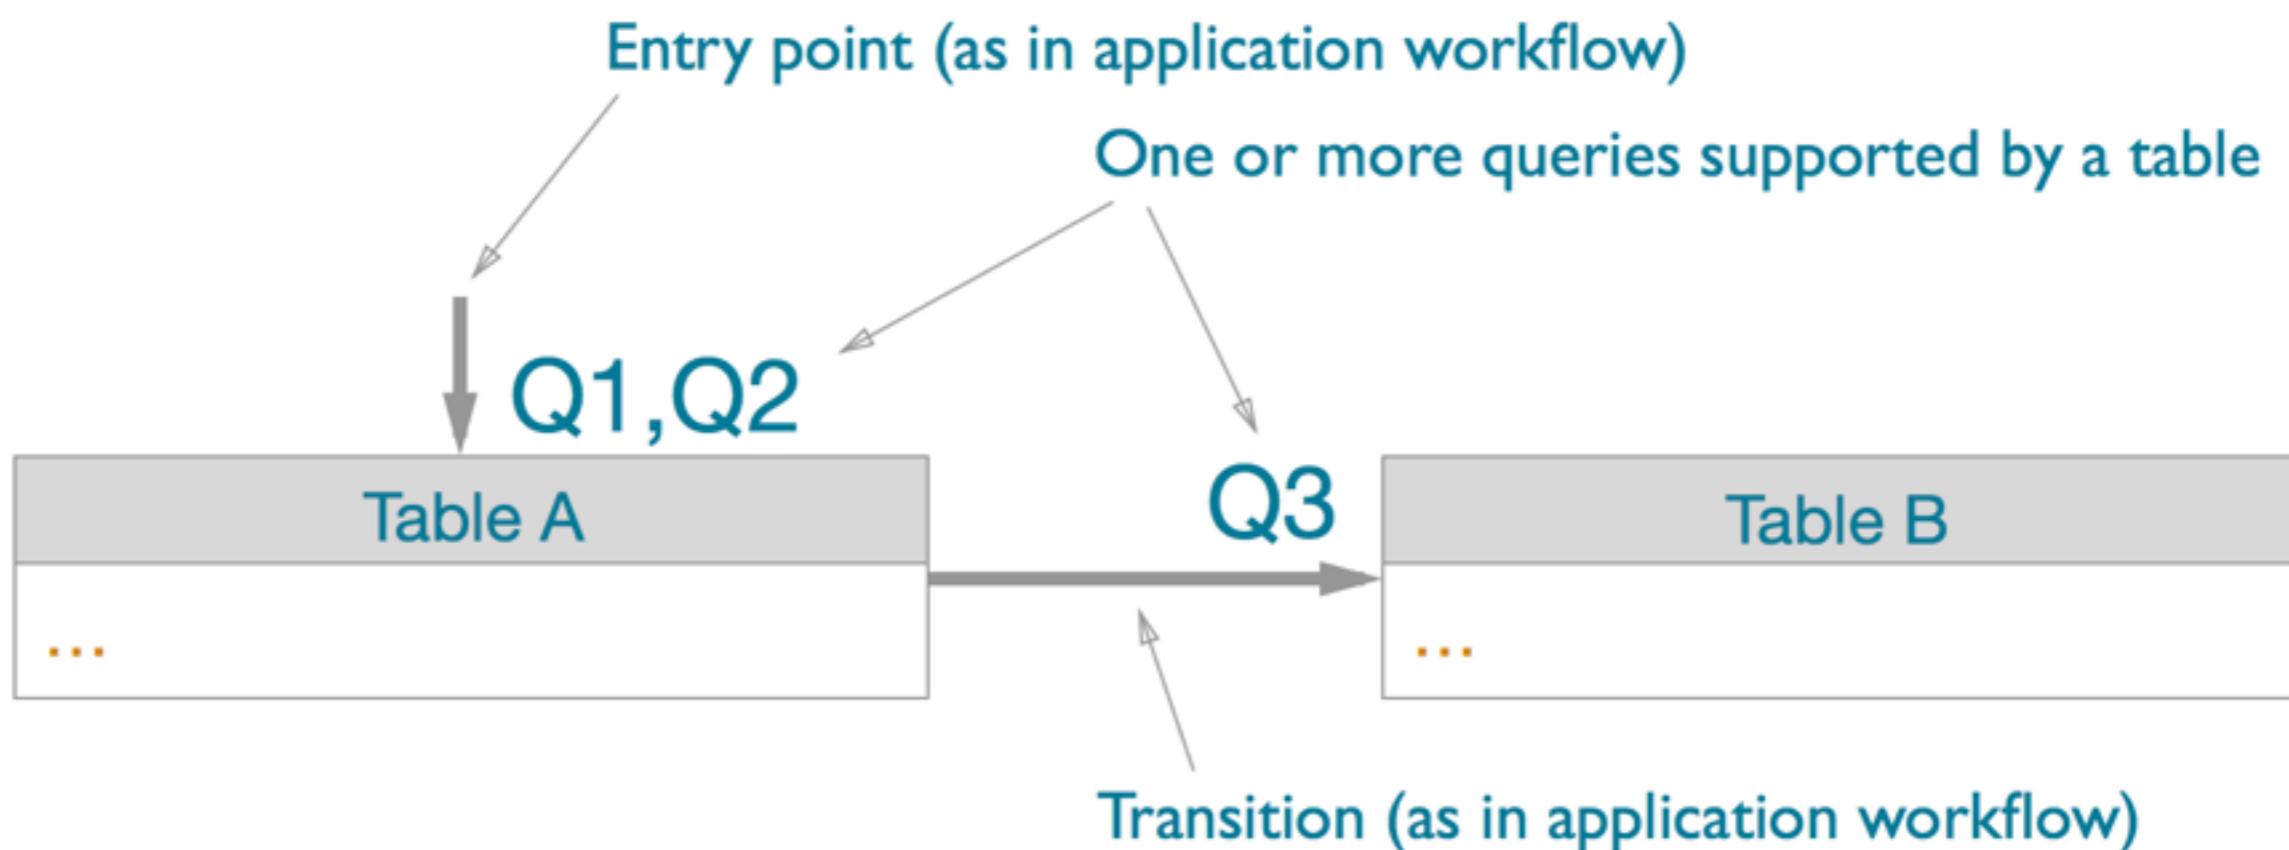
Table representation

- Logical-level shows column names and properties
- Physical-level also shows the column data type

table_name	CQL Type		
column_name_1	K	← Partition key column	
column_name_2	C↑	← Clustering key column (ASC)	
column_name_3	C↓	← Clustering key column (DESC)	
column_name_4	S	← Static column	
column_name_5	IDX	← Secondary index column	
column_name_6	++	← Counter column	
[column_name_7]		← Collection column (list)	
{column_name_8}		← Collection column (set)	
<column_name_9>		← Collection column (map)	
column_name_10	UDT Name	← UDT column	
(column_name_11)	CQL Type	← Tuple column	
column_name_12	CQL Type	← Regular column	

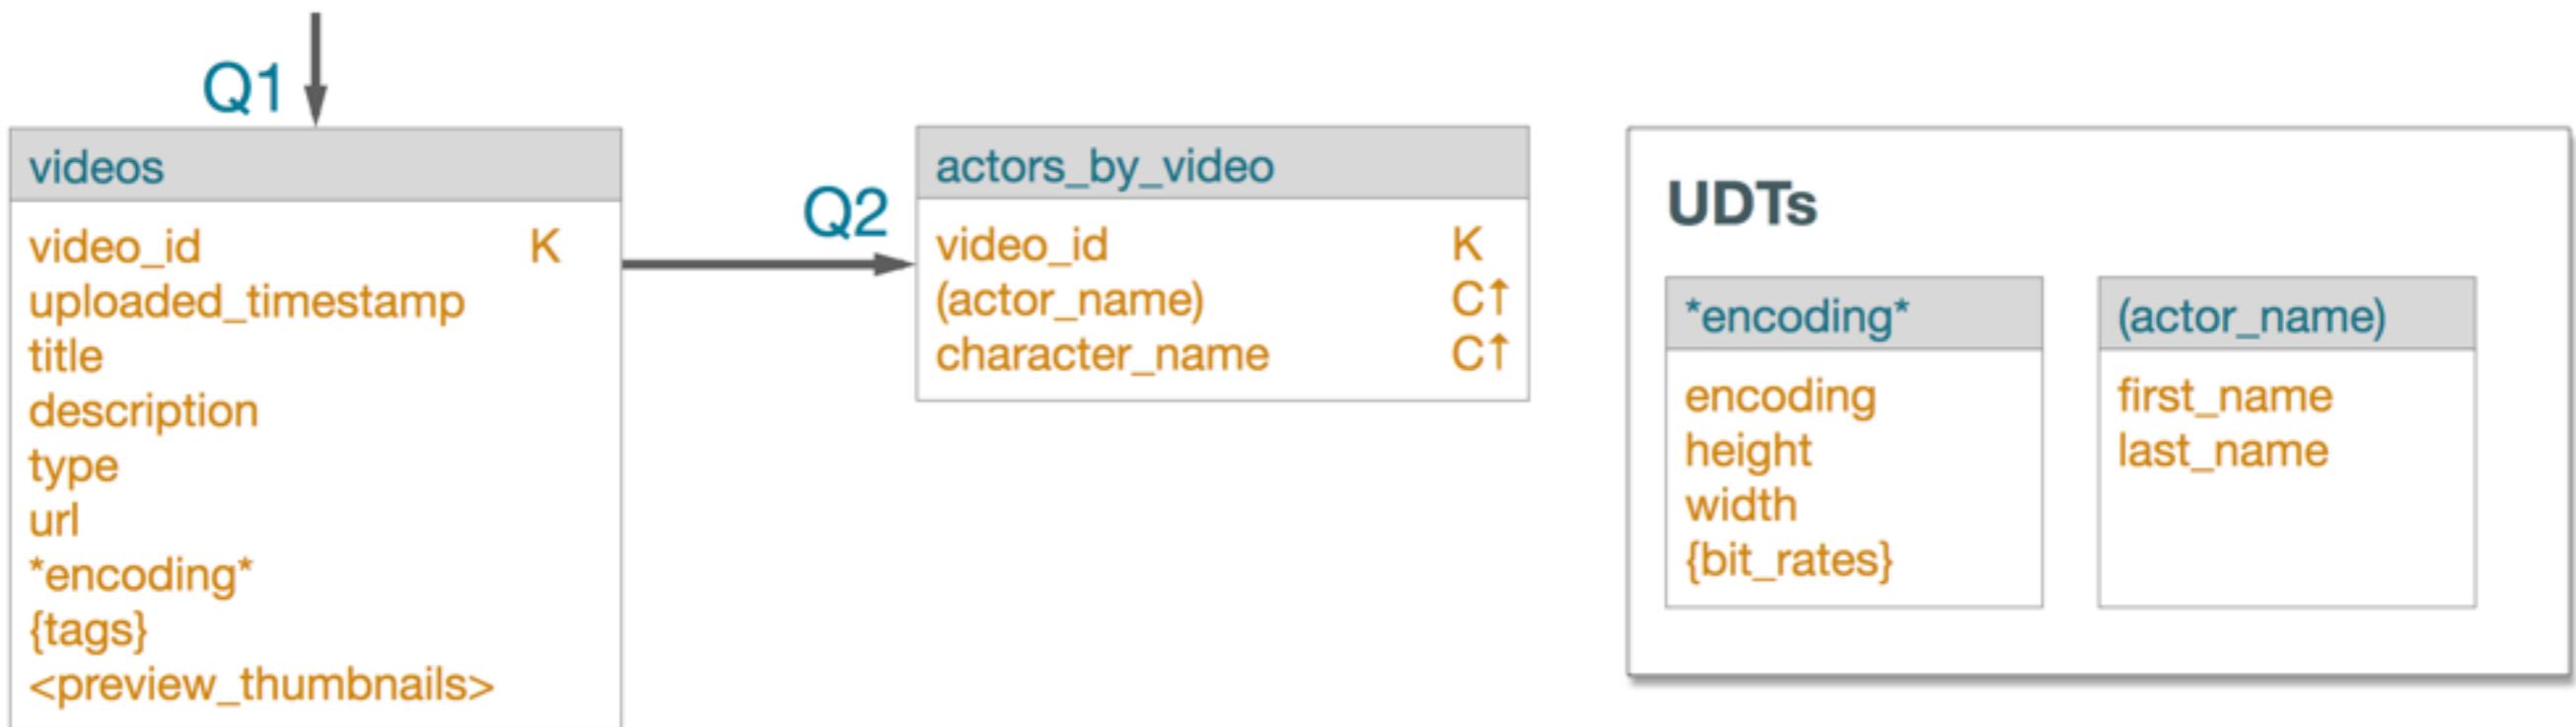
ACCESS PATTERNS

- Directed links and query label shows how tables are accessed
- Similar to the application workflow, but now with our logical information



LOGICAL UDT DIAGRAM

- Represents user defined types and tuples

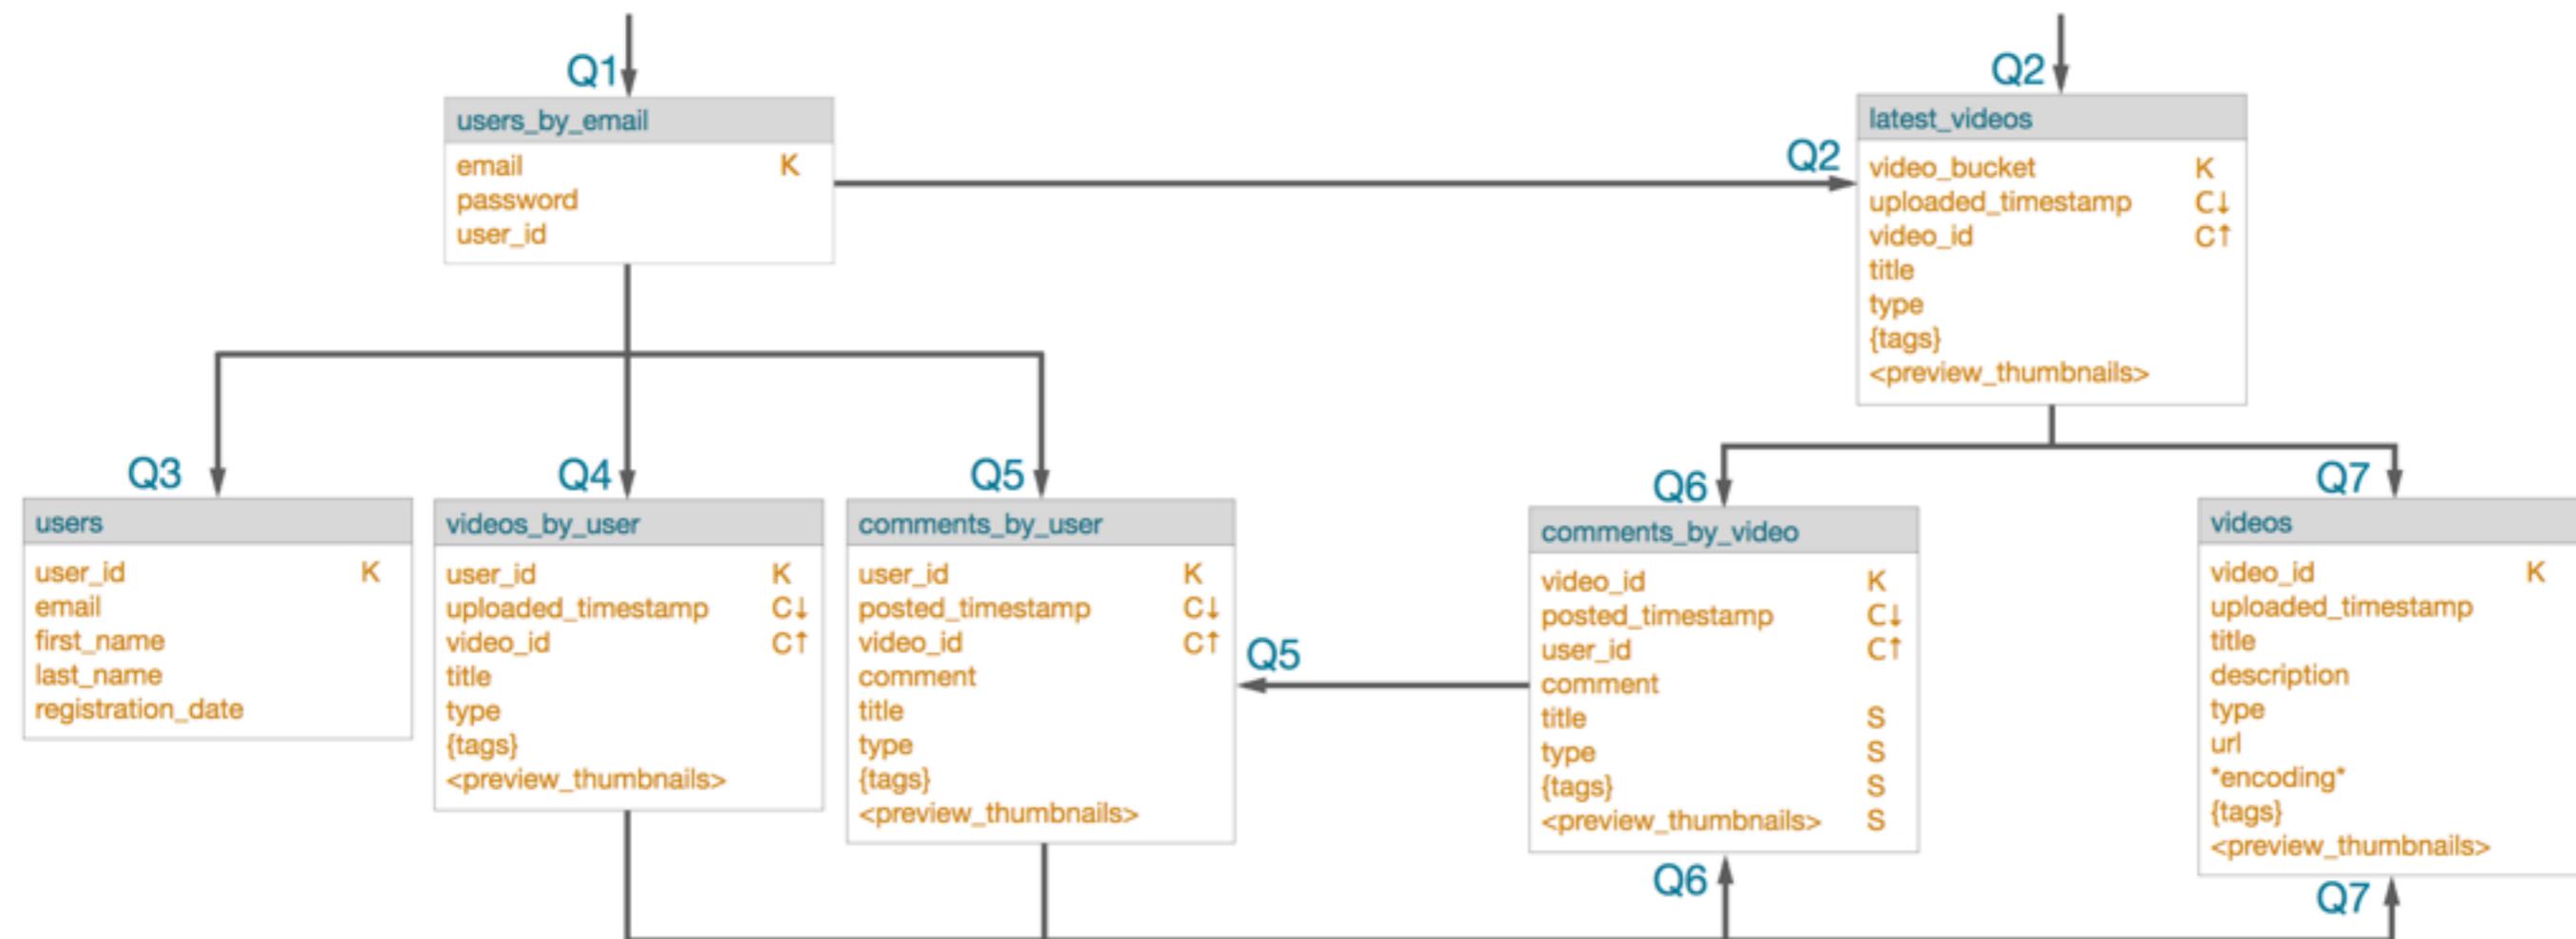


PHYSICAL UDT DIAGRAM

- Represents user defined types and tuples



EXAMPLE CHEBOTKO DIAGRAM



ACCESS PATTERNS

- Q1: Find a user with a **specified email**
- Q2: Find most recently uploaded videos
- Q3: Find a **user with a specified id**
- Q4: Find videos uploaded by a **user with a known id** (show most recently uploaded videos first)
- Q5: Find comments posted by a **user with a known id** (show most recently posted comments first)
- Q6: Find comments posted for a **video with a known id** (show most recent comments first)
- Q7: Find a video with a **specified video id**

UDTs

encoding
encoding
height
width
{bit_rates}

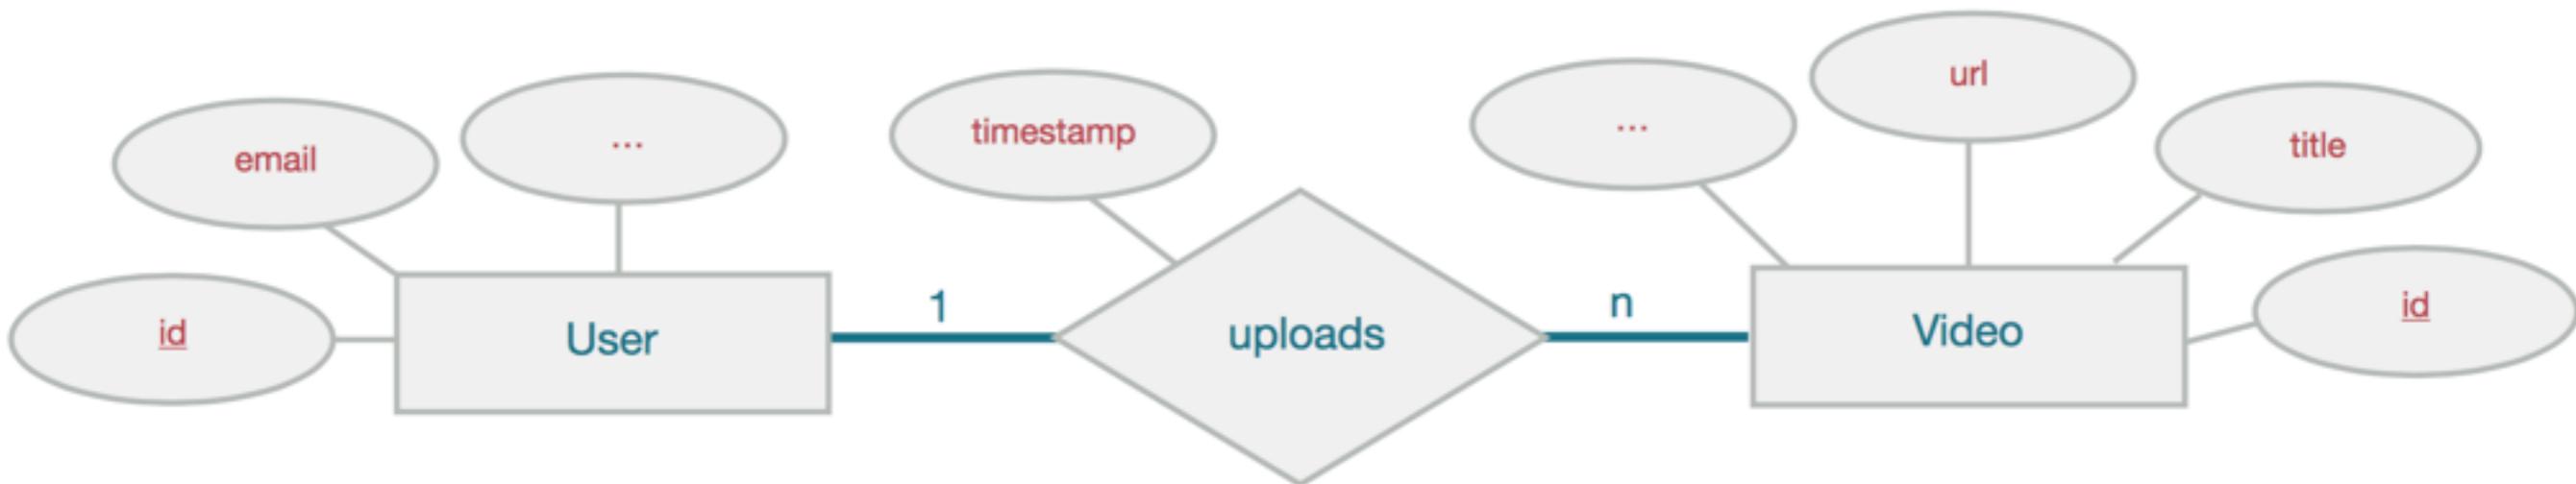
CASSANDRA DATA MODELING PRINCIPLES

- Know your data
- Know your queries
- Nest data
- Duplicate data

KNOW YOUR DATA

Understanding the data is the key to successful design

- Data captured by conceptual data model
- Define what is stored in database
- Preserve properties so that data is organized correctly



KNOW YOUR DATA

Key constraints affect schema design

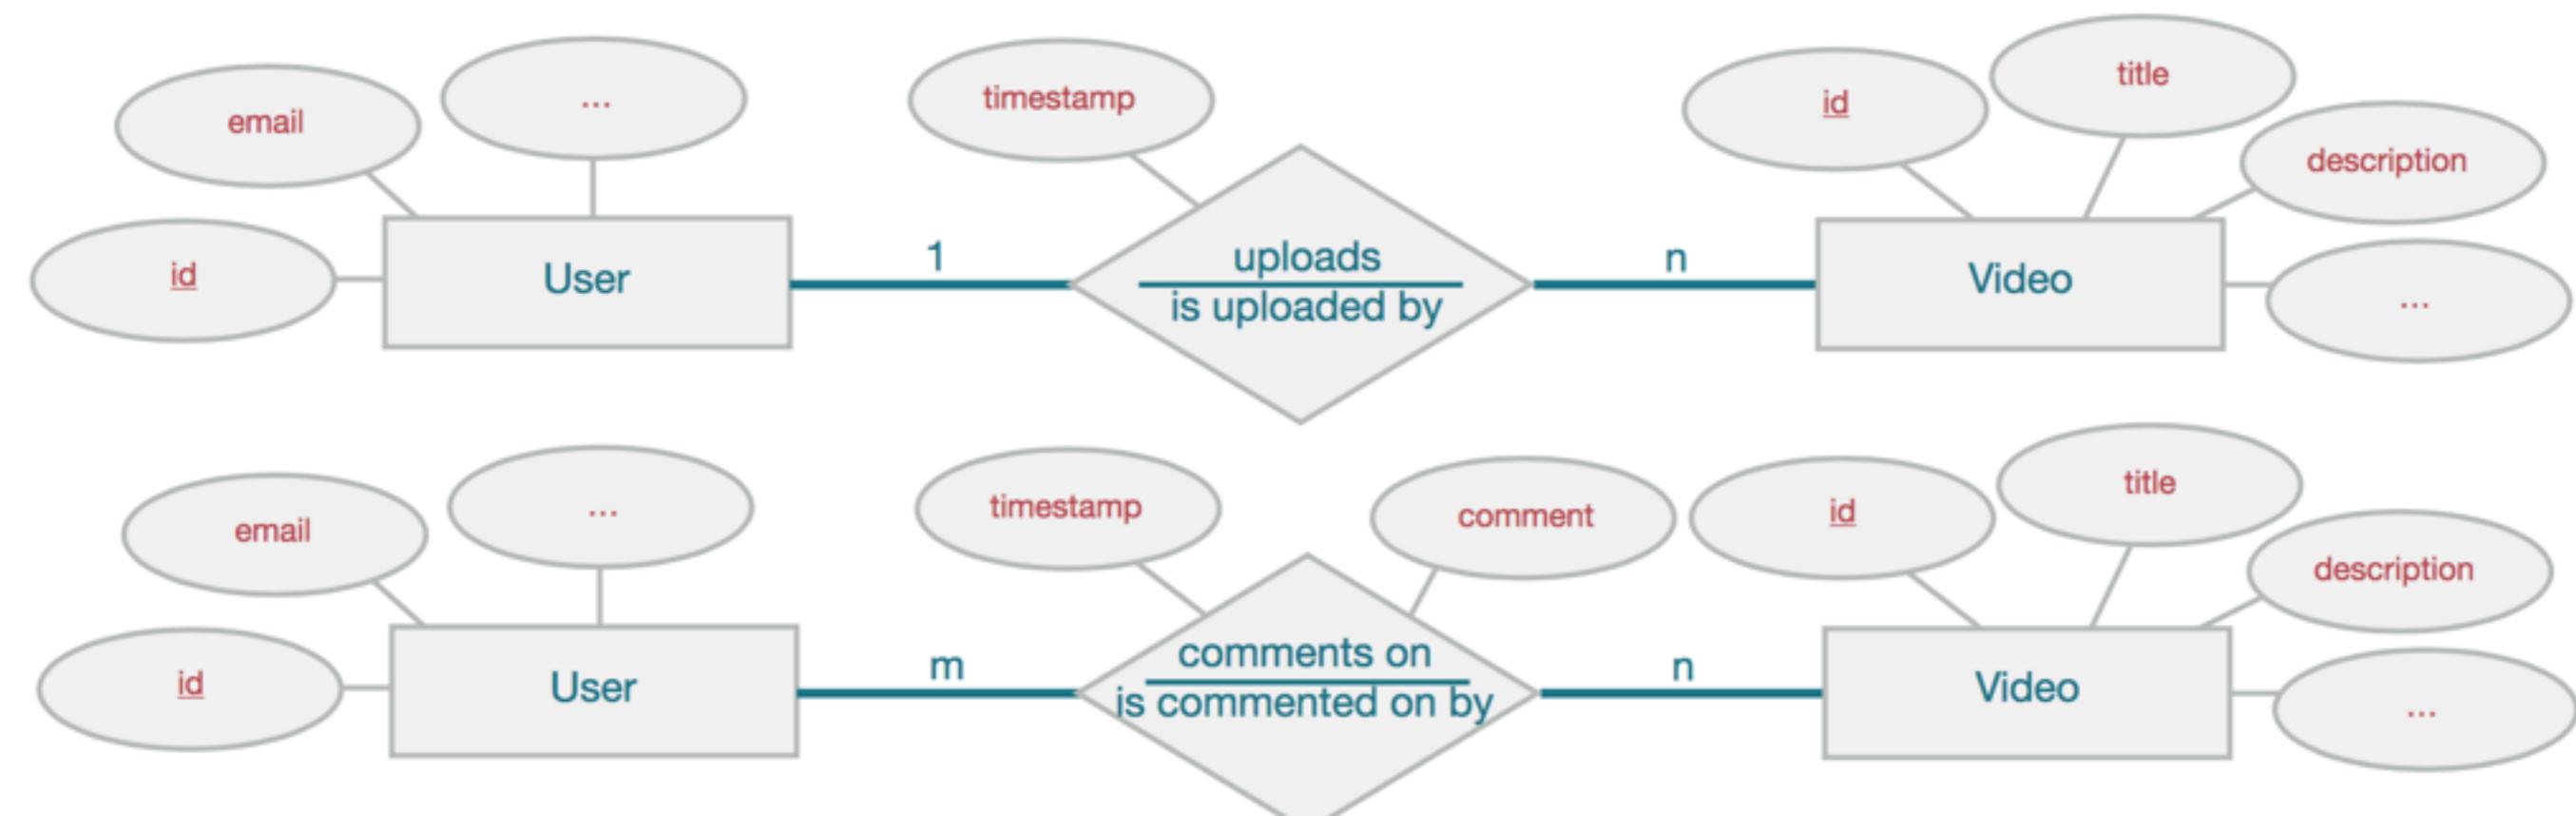
- Entity and relationship keys affect the table primary keys
- Primary key uniquely identifies a row / entity / relationship
- Composed of a key and possibly additional columns

videos	
video_id	K
uploaded_timestamp	
user_id	
title	
description	
type	
encoding	
{tags}	
<preview_thumbnails>	
{genres}	

videos_by_user	
user_id	K
uploaded_timestamp	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

KNOW YOUR DATA

Cardinality constraints affect the key for relationships



KNOW YOUR QUERIES

Queries directly affect schema design

- Queries captured by application workflow model
- Table schema design changes if queries change



ACCESS PATTERNS

Q1: Find a user with a **specified email**

Q2: Find most recently uploaded**videos**

KNOW YOUR QUERIES

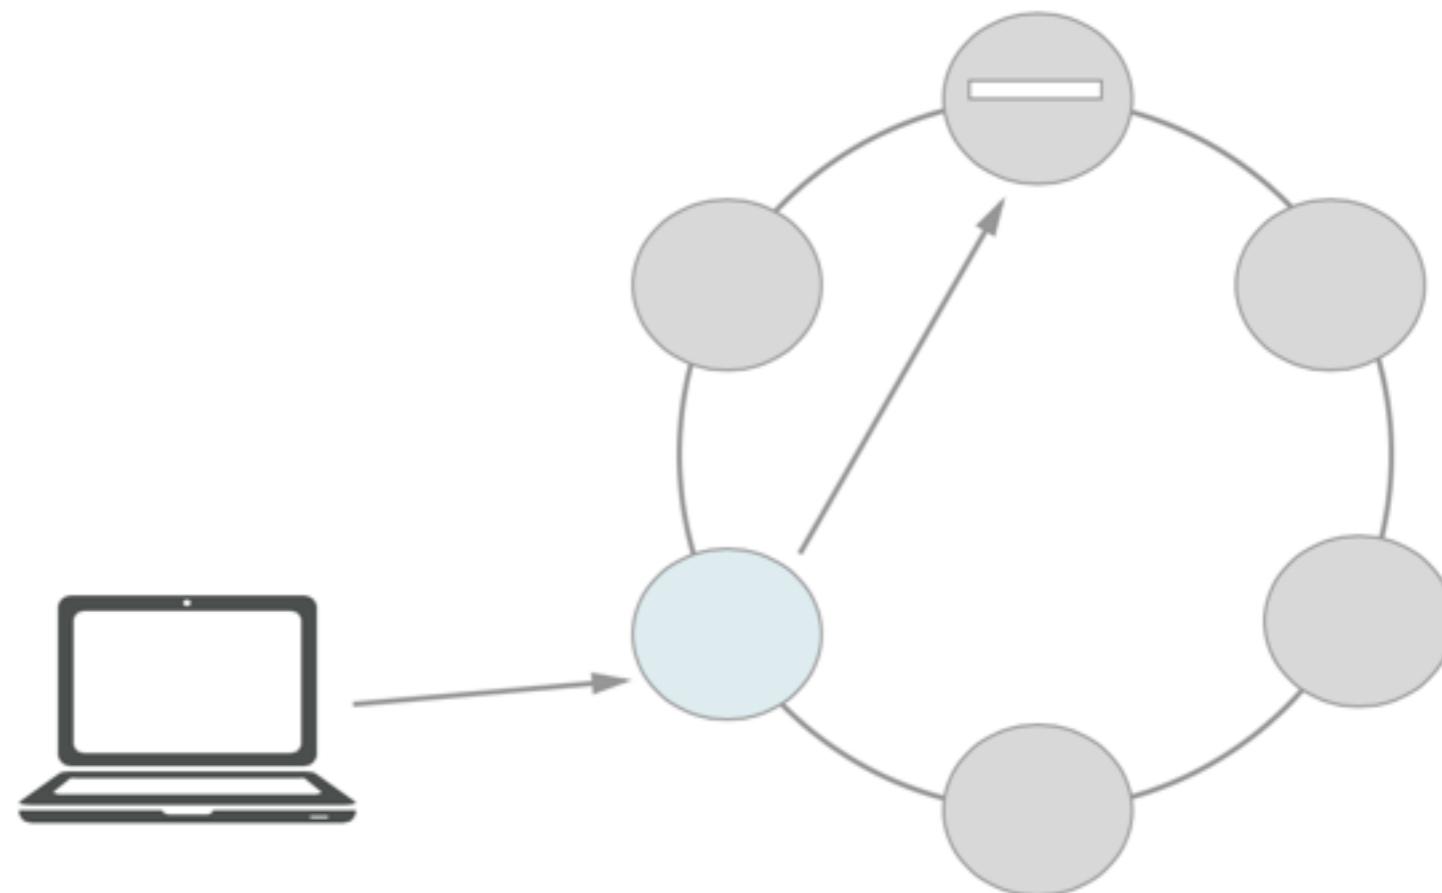
Schema design organizes data to efficiently run queries

- Partition per query – **ideal**
- Partition+ per query – **acceptable**
- Table scan – **anti-pattern**
- Multi-table – **anti-pattern**

KNOW YOUR QUERIES

Partition per query

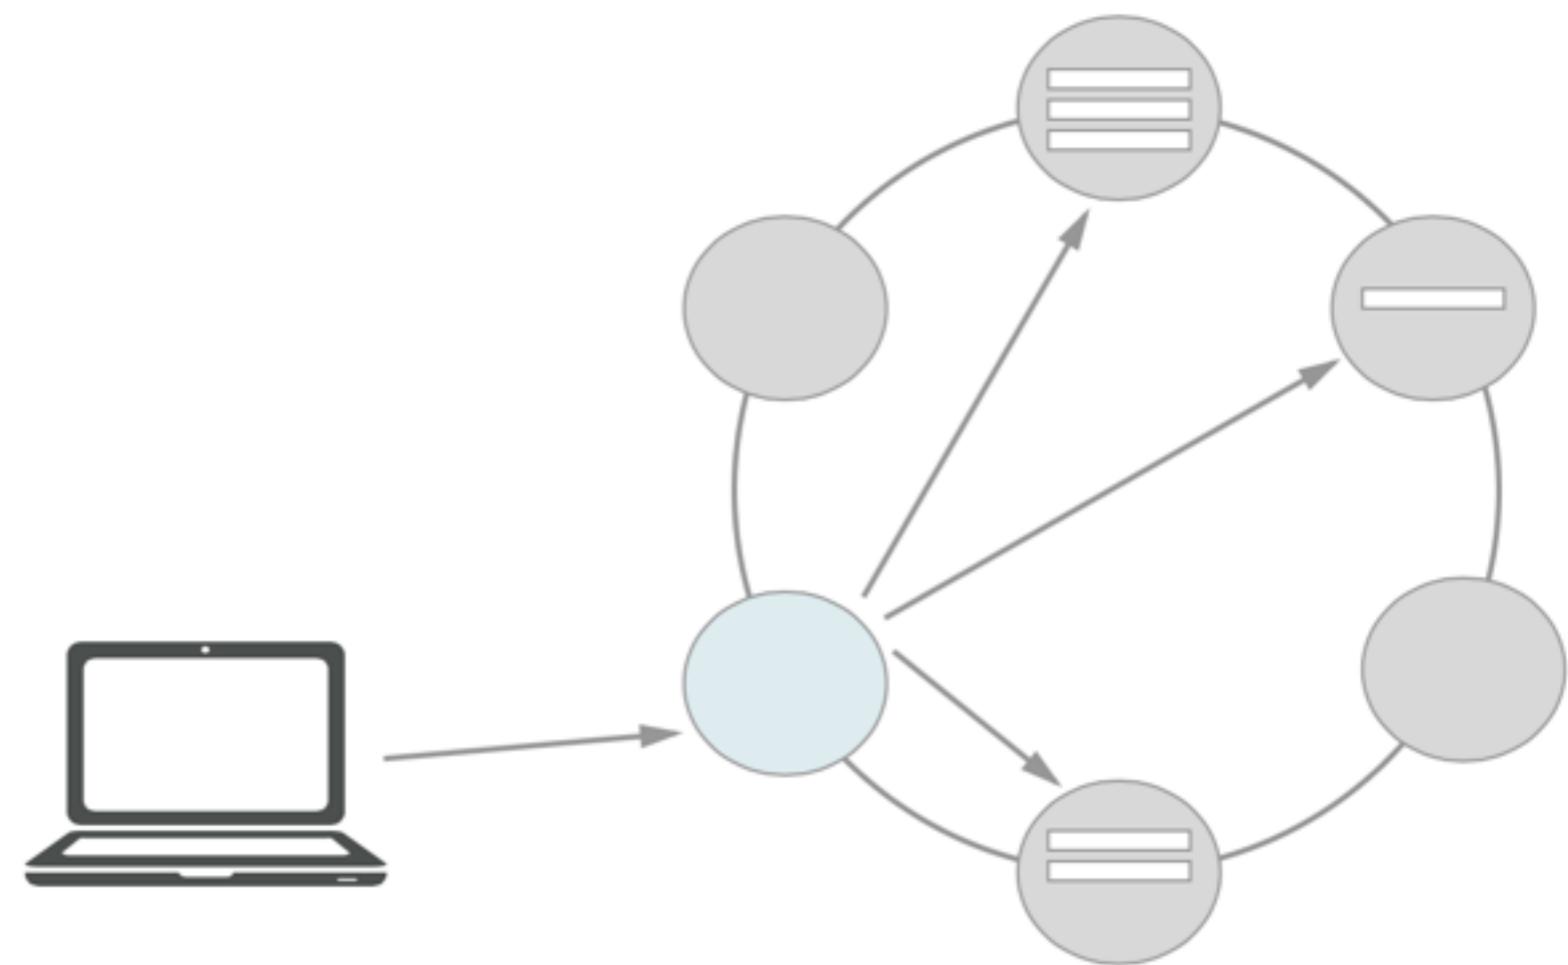
- The most efficient access pattern
- Query accesses only one partition to retrieve results
- Partition can be single-row or multi-row



KNOW YOUR QUERIES

Partition+ per query

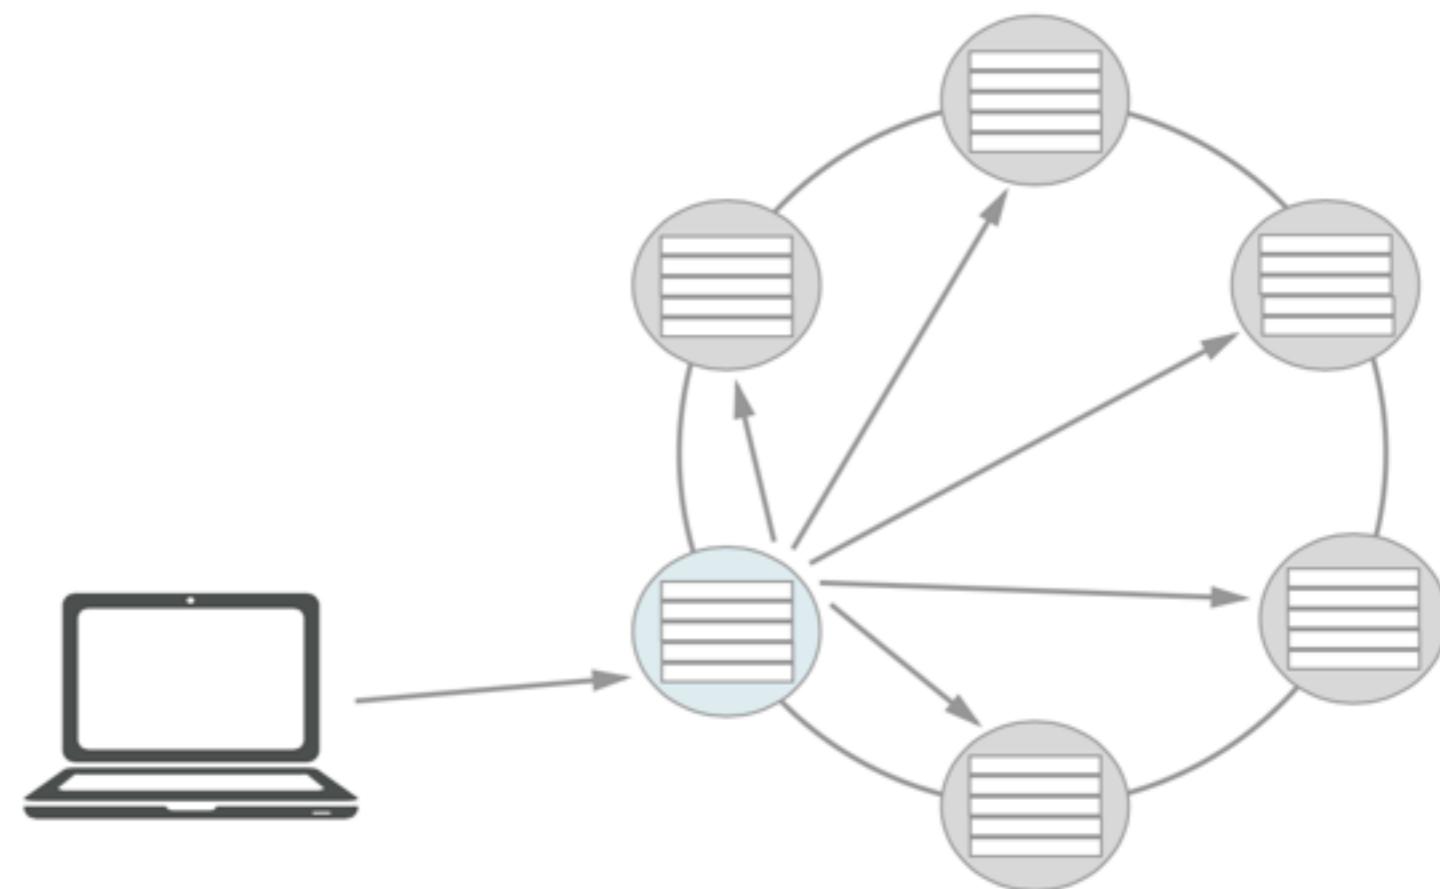
- Less efficient access pattern but not necessarily bad
- Query needs to access multiple partitions to retrieve results



KNOW YOUR QUERIES

Table scan and Multi-table

- Least efficient type of query but may be needed in some cases
- Query needs to access all partitions in a table(s) to retrieve results



NEST DATA

Data nesting is the main data modeling technique

- Nesting organizes multiple entities into a single partition
- Supports partition per query data access
- Three data nesting mechanisms
 - Clustering columns—multi-row partitions
 - Collection columns
 - User-defined type columns

NEST DATA

Clustering columns - primary data nesting mechanism

- Partition key identifies an entity that other entities will nest into
- Values in a clustering column identify the nested entities
- Multiple clustering columns implement multi-level nesting

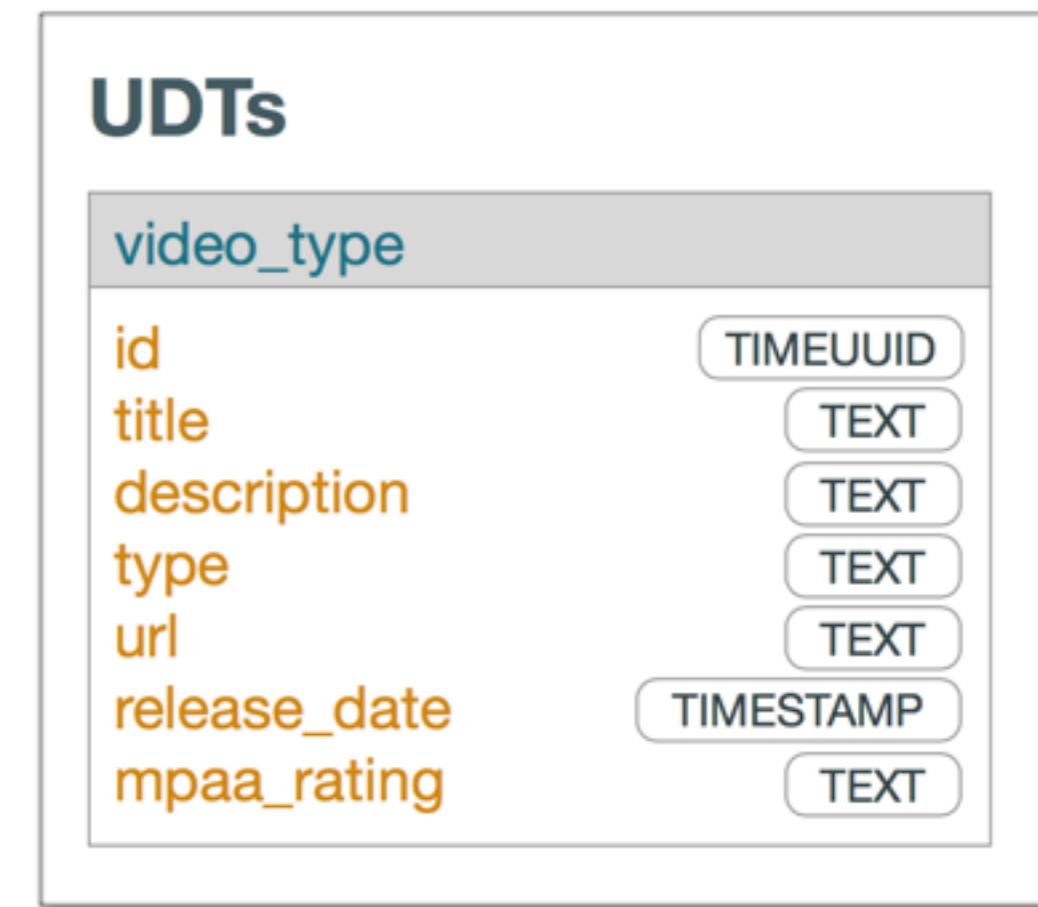
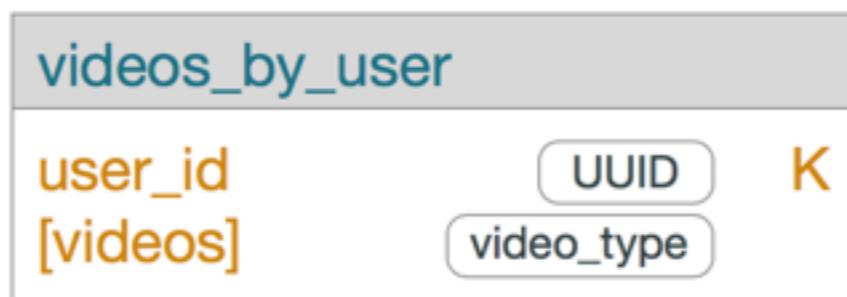
videos	
video_id	K
uploaded_timestamp	
user_id	
title	
description	
type	
{tags}	
<preview_thumbnails>	
{genres}	

actors_by_video	
video_id	K
actor_name	C↑
character_name	C↑

NEST DATA

User-defined type—secondary data nesting mechanism

- Represents one-to-one relationship, but can use in conjunction with collections
- Easier than working with multiple collection columns



DUPLICATE DATA

Better to duplicate than to join data

- Partition per query and data nesting may result in data duplication
 - Query results are pre-computed and materialized
 - Data can be duplicated across tables, partitions, and / or rows

videos_by_actor	
actor	K
release_date	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

videos_by_genre	
genre	K
release_date	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

videos_by_tag	
tag	K
release_date	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

DUPLICATE DATA

Data duplication can scale, joins cannot

videos_by_actor		
actor	K	
video_id	C↑	
character_name	C↑	

Q1 →

videos		
video_id	K	
uploaded_timestamp	C↓	
title		
description		
type		
release_date		
{tags}		
<preview_thumbnails>		
{genres}		

Q2 ↓

videos_by_actor		
actor	K	
video_id	C↑	
character_name	C↑	
uploaded_timestamp		
title		
description		
type		
release_date		
{tags}		
<preview_thumbnails>		
{genres}		

Q1 →

VS.

MAPPING RULES

For the query-driven methodology

- Mapping rules ensure that a logical data model is correct
- Each query has a corresponding table
- Tables are designed to allow queries to execute properly
- Tables return data in the correct order

MAPPING RULES

- Mapping Rule 1: Entities and relationships
- Mapping Rule 2: Equality search attributes
- Mapping Rule 3: Inequality search attributes
- Mapping Rule 4: Ordering attributes
- Mapping Rule 5: Key attributes

MRI: ENTITIES AND RELATIONSHIPS

- Entity and relationship types map to tables
- Entities and relationships map to partitions or rows
- Partition may have data about one or more entities and relationships
- Attributes are represented by columns

Conceptual Model



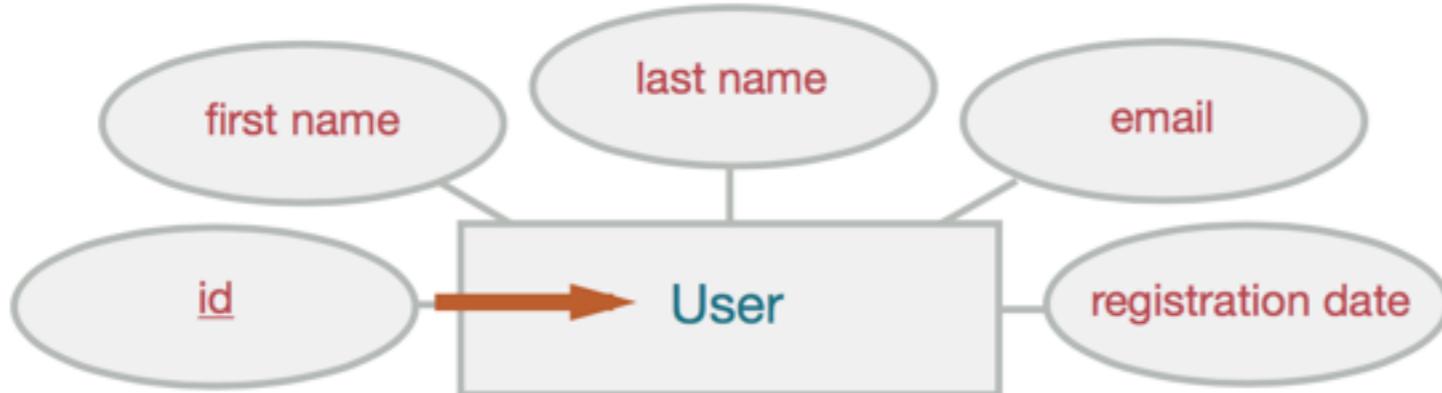
Logical Model

users	K
user_id	
email	
first_name	
last_name	
registration_date	

MRI: ENTITIES AND RELATIONSHIPS

- Entity and relationship types map to tables
- Entities and relationships map to partitions or rows
- Partition may have data about one or more entities and relationships
- Attributes are represented by columns

Conceptual Model



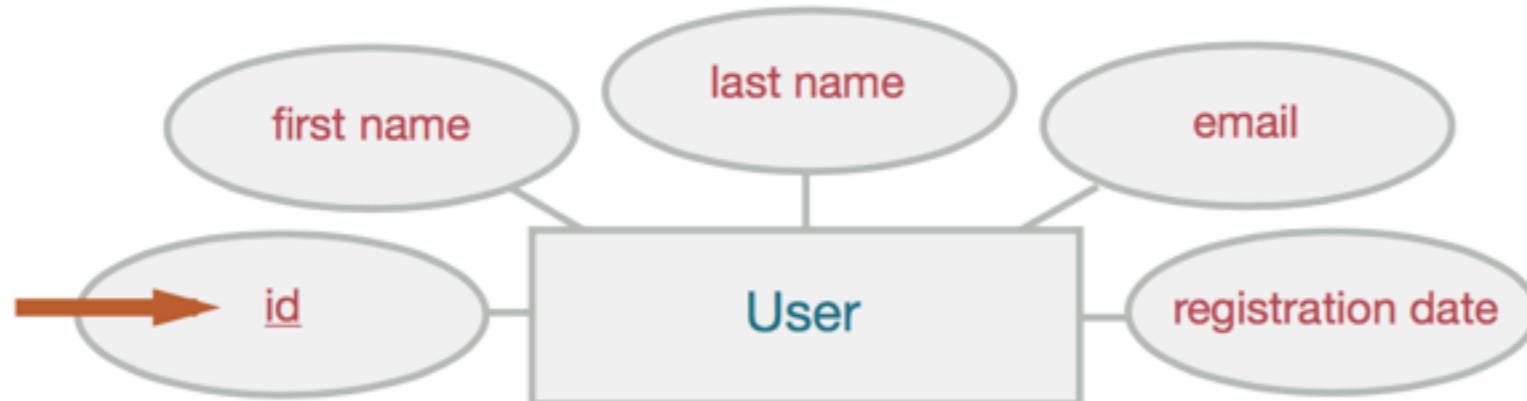
Logical Model

users
user_id
email
first_name
last_name
registration_date

MRI: ENTITIES AND RELATIONSHIPS

- Entity and relationship types map to tables
- Entities and relationships map to partitions or rows
- Partition may have data about one or more entities and relationships
- Attributes are represented by columns

Conceptual Model



Logical Model

users
user_id
email
first_name
last_name
registration_date

MRI: ENTITIES AND RELATIONSHIPS

- Entity and relationship types map to tables
- Entities and relationships map to partitions or rows
- Partition may have data about one or more entities and relationships
- Attributes are represented by columns

Conceptual Model



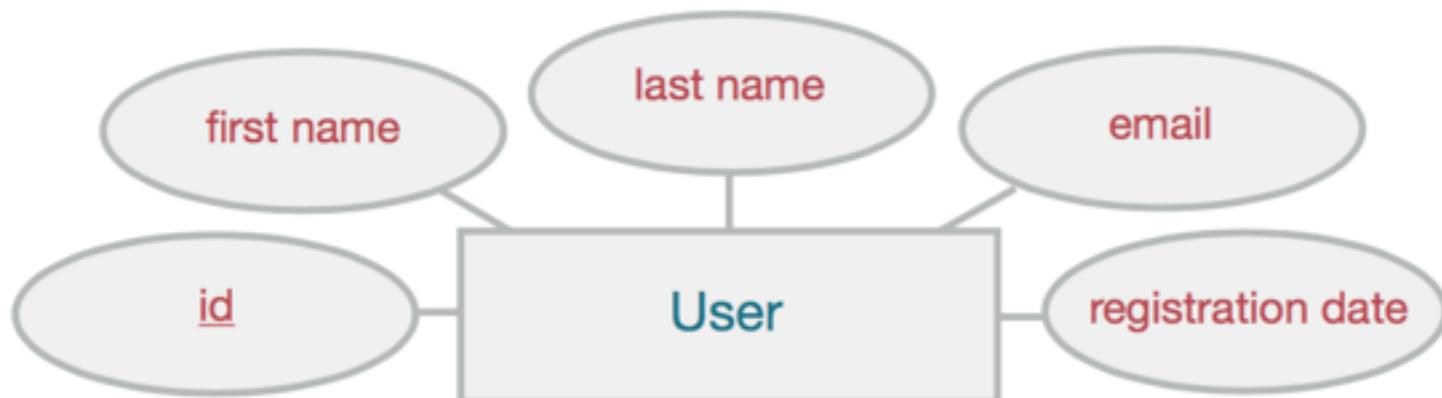
Logical Model

users	K
user_id	
email	
first_name	
last_name	
registration_date	

MRI: ENTITIES AND RELATIONSHIPS

- Entity and relationship types map to tables
- Entities and relationships map to partitions or rows
- Partition may have data about one or more entities and relationships
- Attributes are represented by columns

Conceptual Model



Logical Model

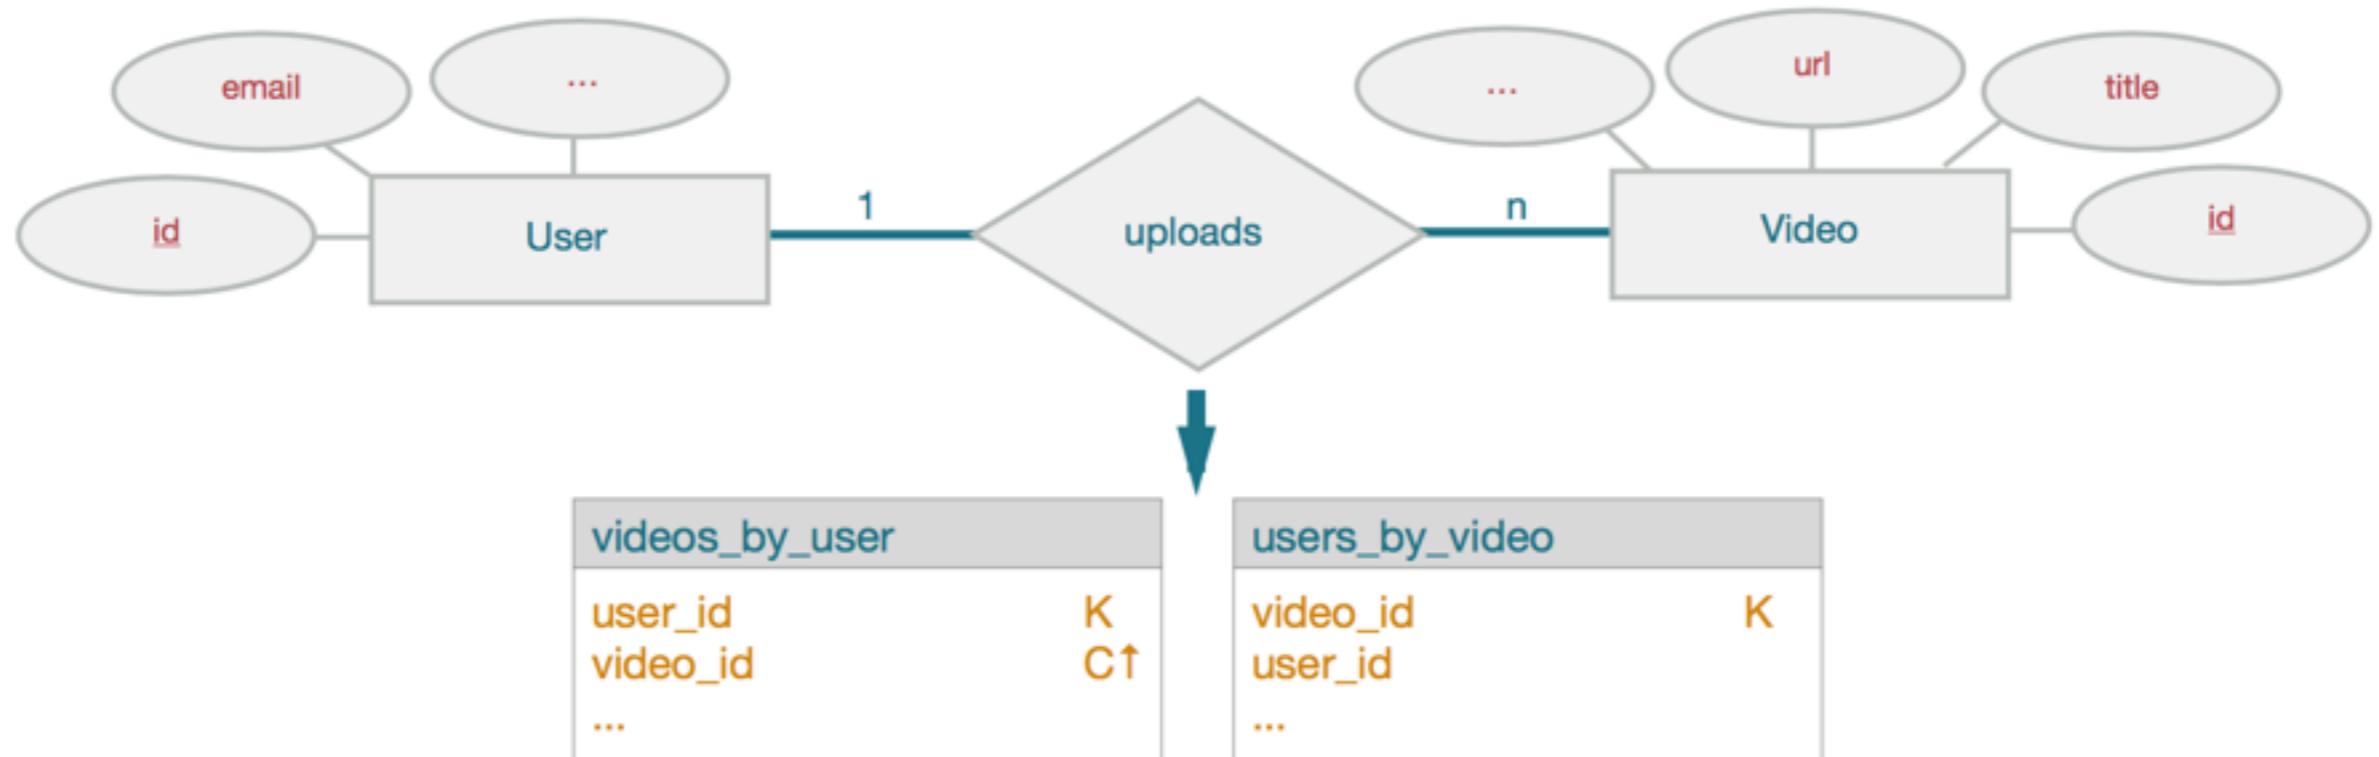
users	K
user_id	
email	
first_name	
last_name	
registration_date	

user_id	email	first_name	last_name	registration_date
a7e78478-0a54-4949-90f3-14ec4cbea40c	jbellis@datastax.com	Jonathan	Ellis	2010-04-01 00:00:00+0000

MRI: ENTITIES AND RELATIONSHIPS

Relationship type example

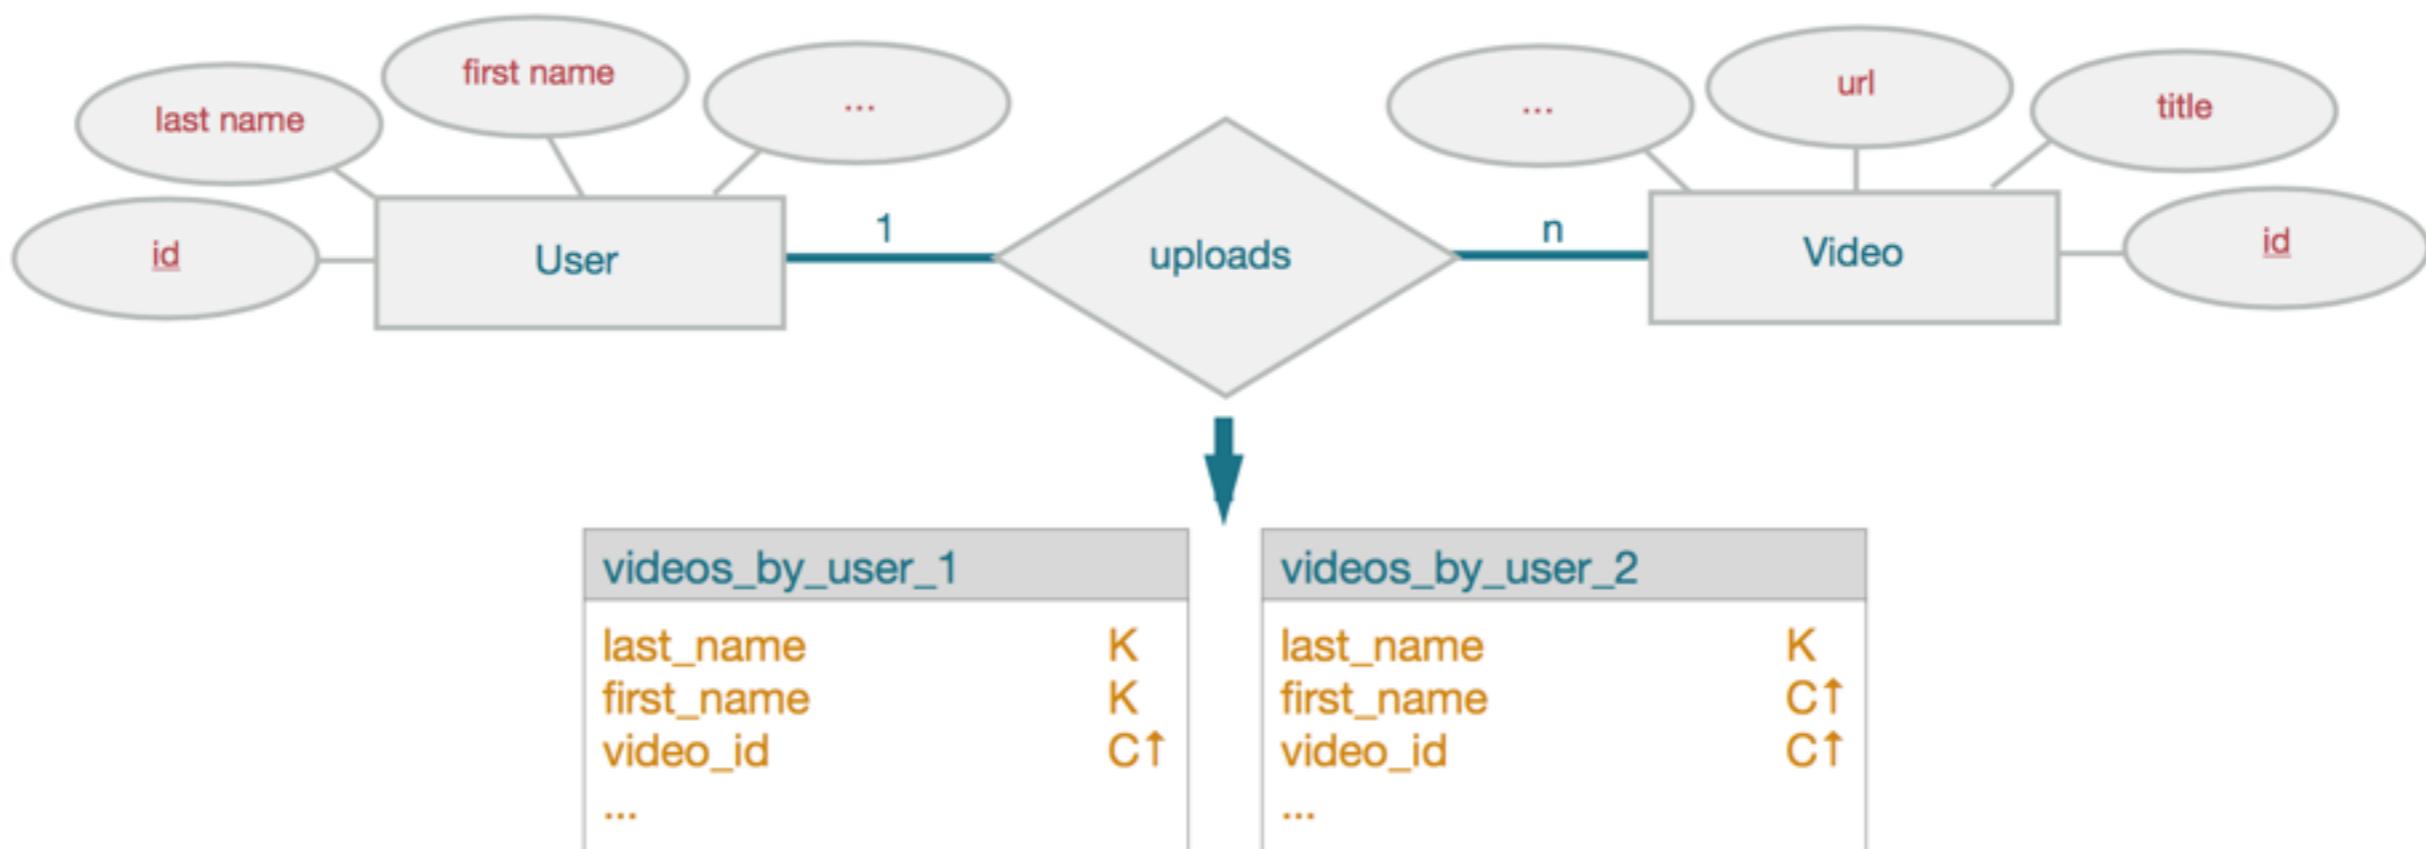
- Each relationship becomes a row in the table
- Relationship type attributes are represented by columns
- Queries and relationship cardinality affects the design of the primary key



MR2: EQUALITY SEARCH ATTRIBUTES

Relationship type example

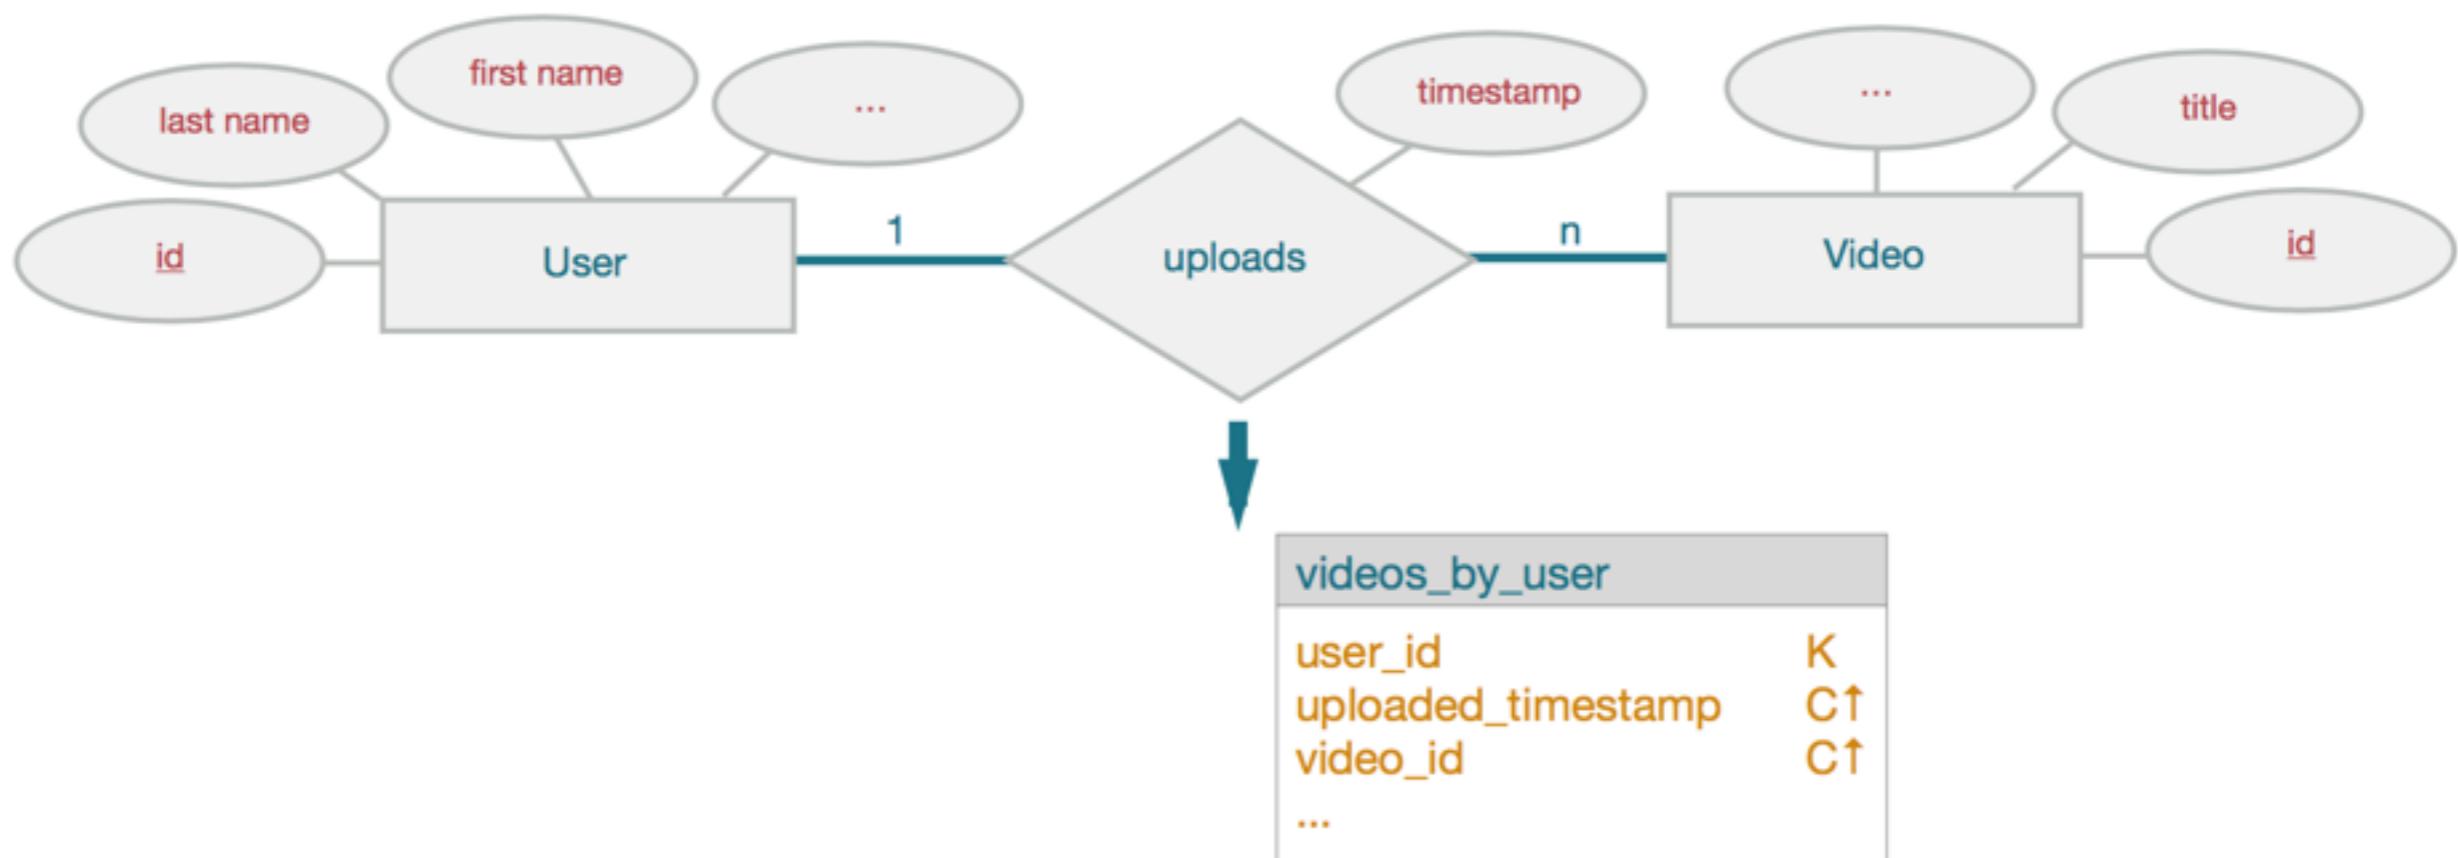
- Equality search attributes become initial columns of a primary key
- Querying on: `last name` and `first name`



MR3: INEQUALITY SEARCH ATTRIBUTES

Relationship type example

- Inequality search attributes become clustering columns
- Querying on: $user_id = ?$ and $uploaded_timestamp > ?$



MR4: ORDERING ATTRIBUTES

Entity type example

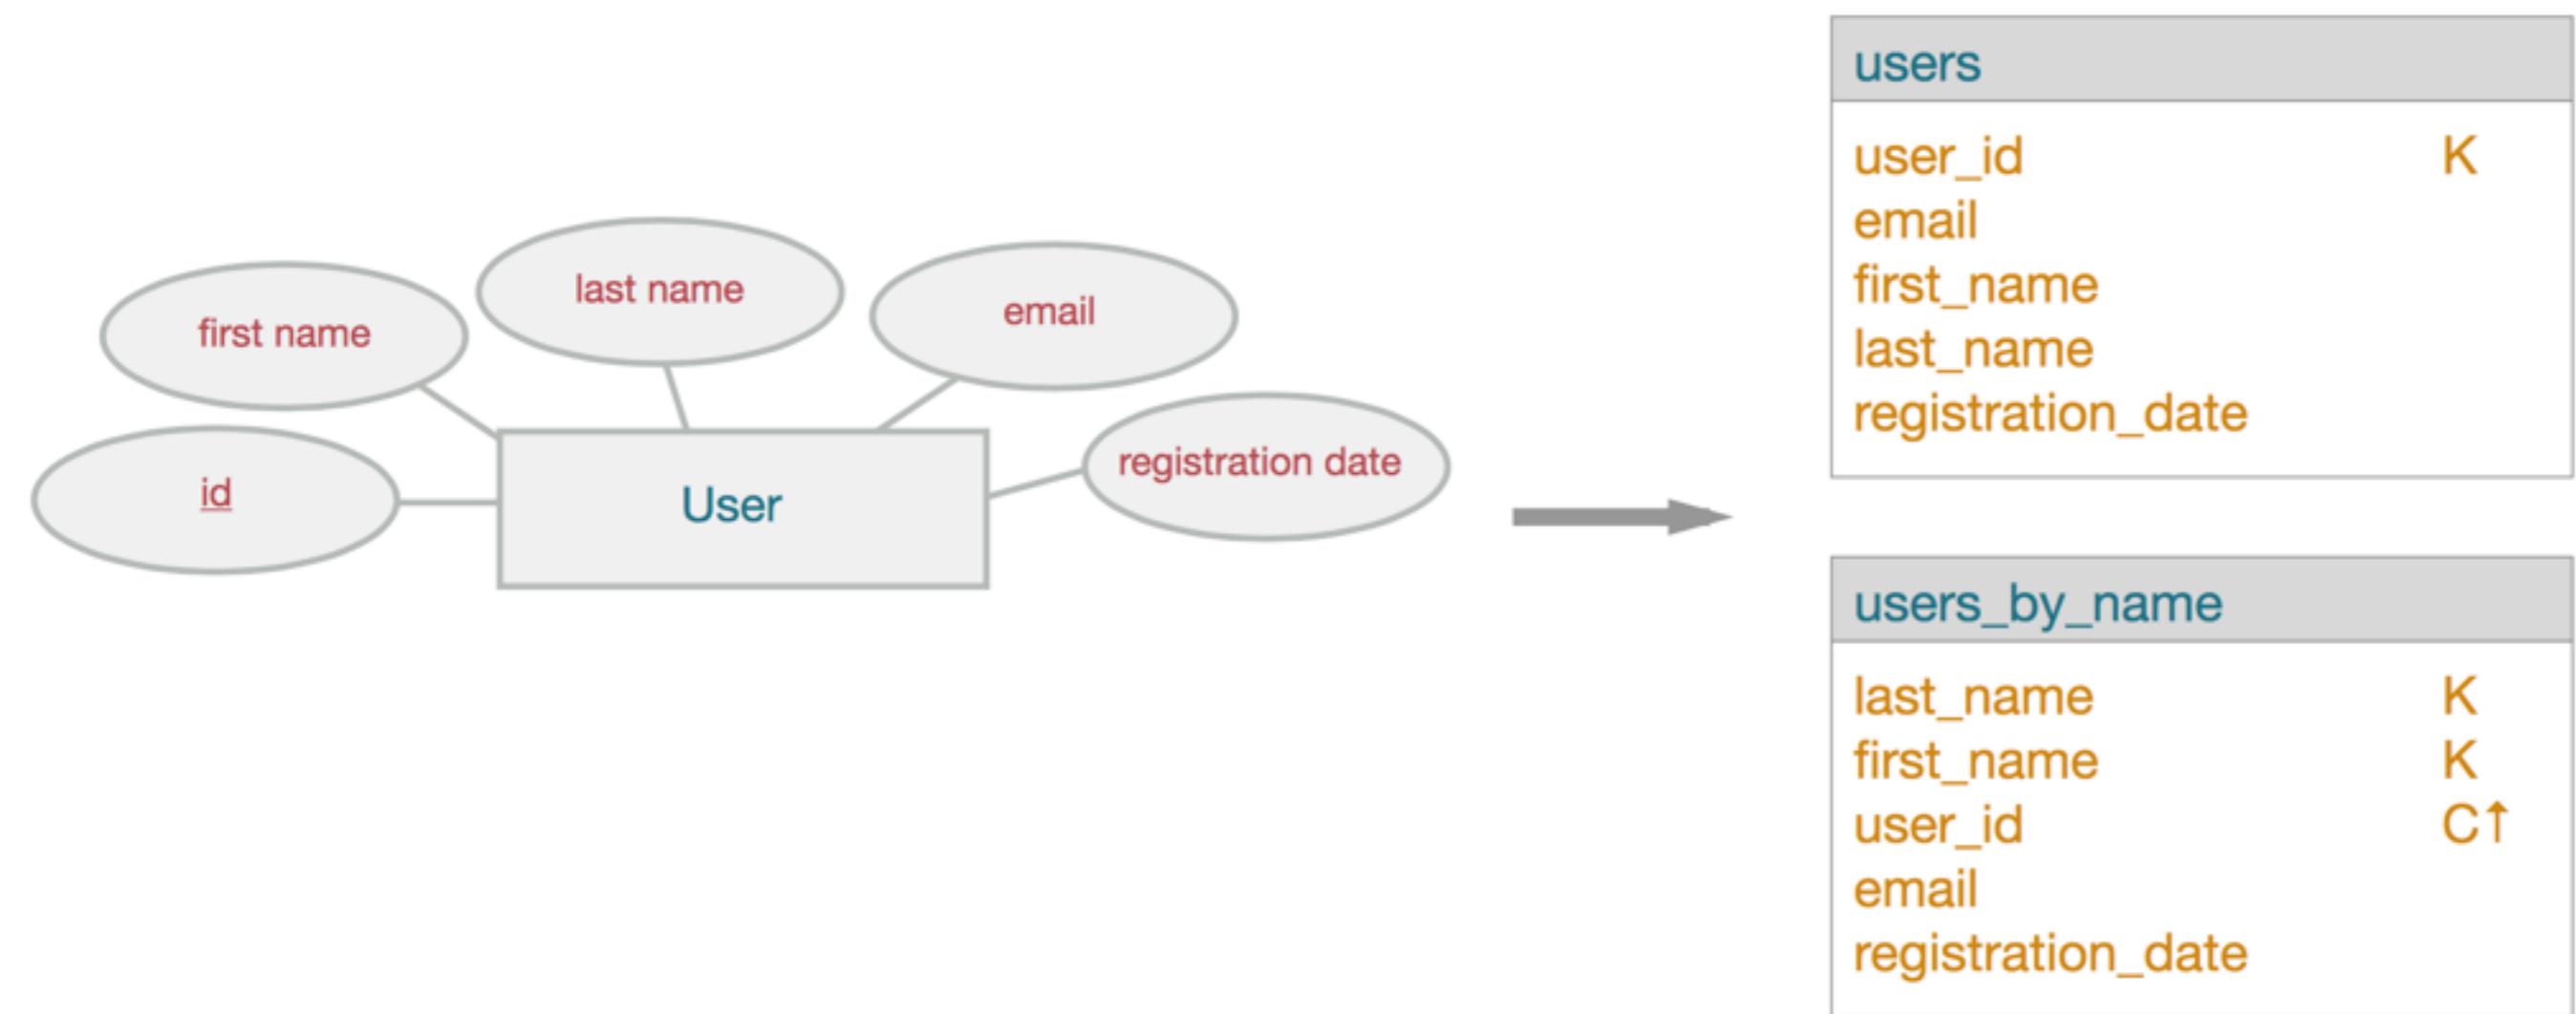
- Ordering attributes become clustering columns
- Querying on: *last_name = ? and registration_date > ?*
- Ordering attributes: *registration_date* (ASC)



MR5: KEY ATTRIBUTES

Entity type example

- Entity type key attributes are included as primary key columns
- Queries also affect primary key design

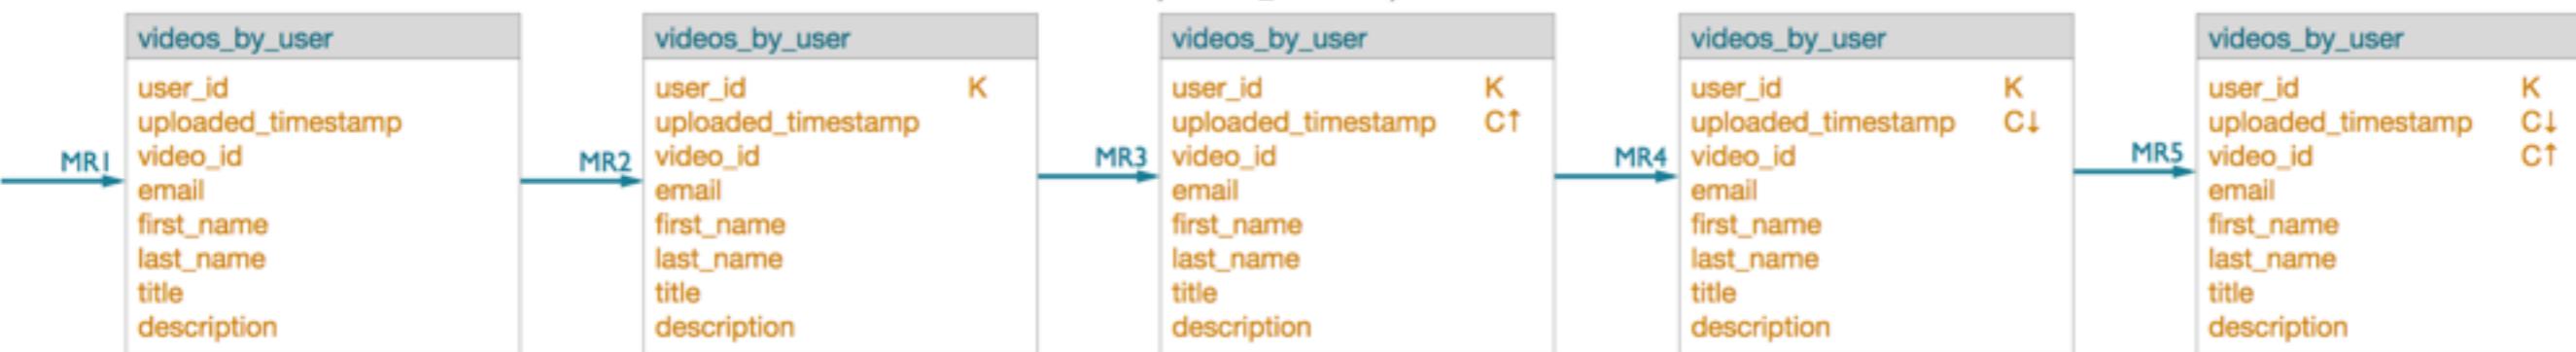


APPLYING MAPPING RULES

- Create a table schema from the conceptual data model and for each query
- Apply the mapping rules in order—MR1—MR2—MR3—MR4—MR5



Query: user_id = ? and uploaded_timestamp > ?
ORDER BY uploaded_timestamp DESC



PHYSICAL DATA MODEL

Query-driven
Methodology

Logical Data
Model

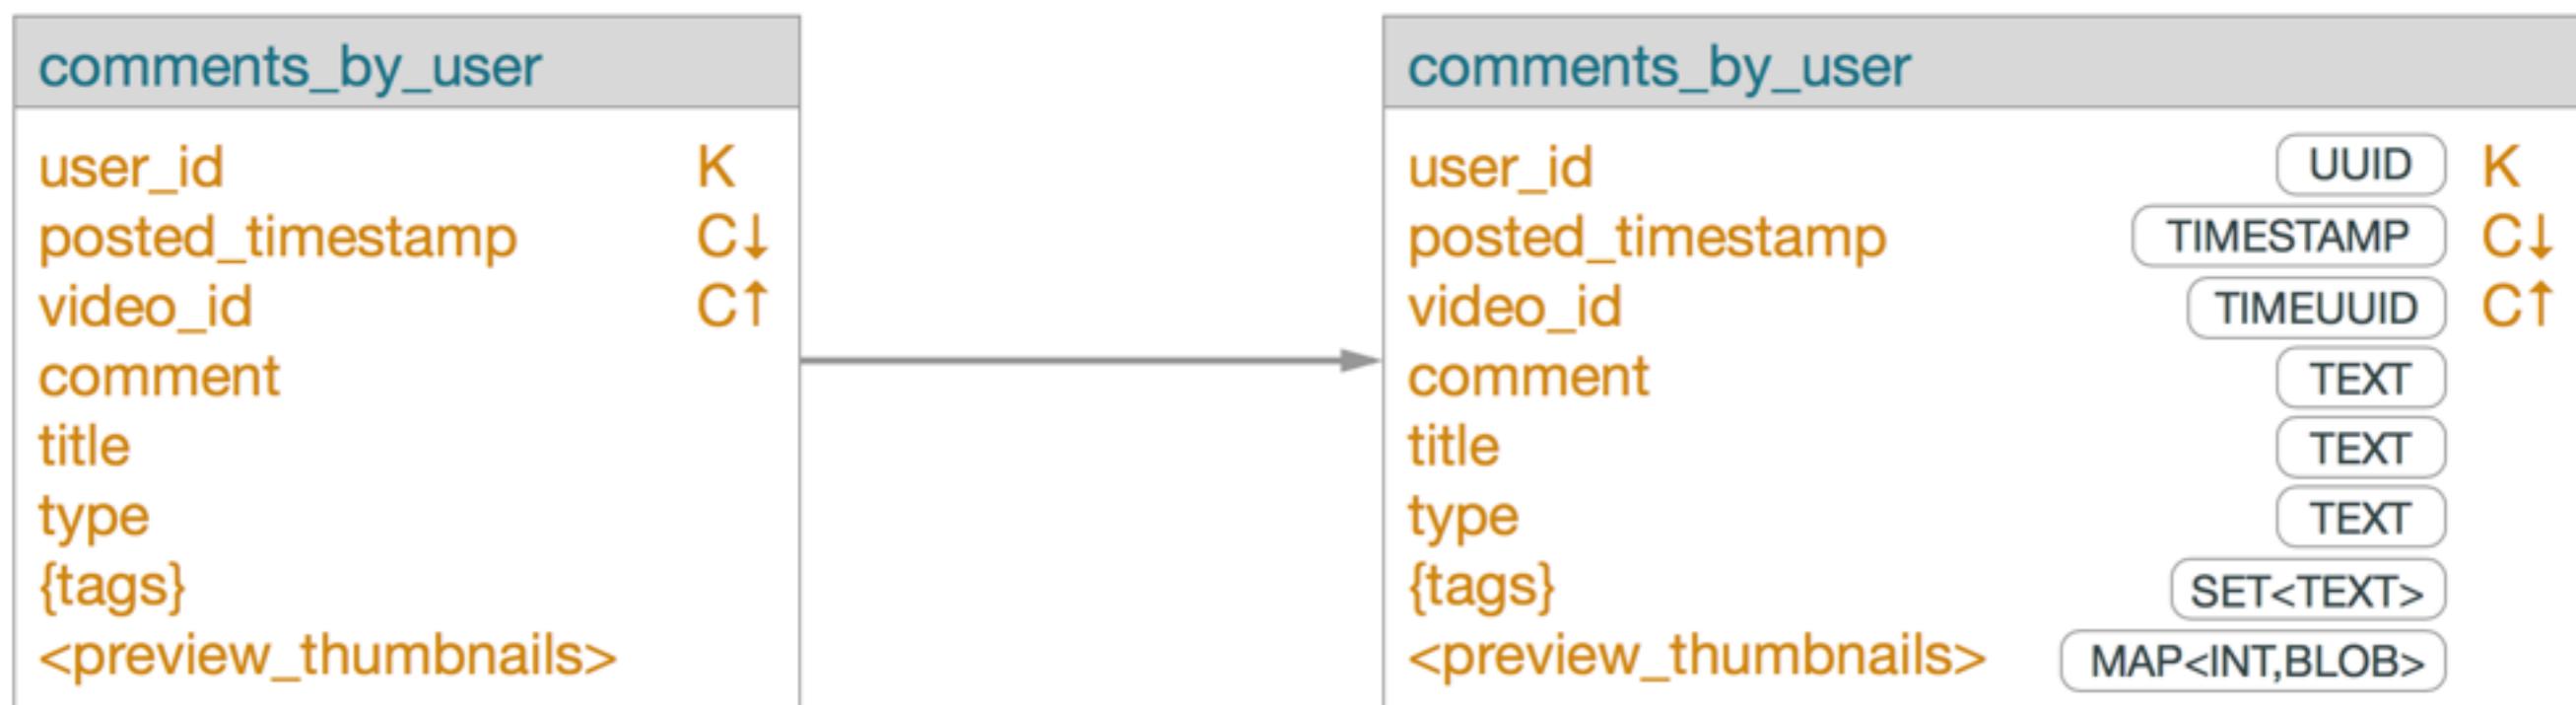
**Analysis and
Validation**

WHAT TO ANALYZE

- Finding the problem
 - Partition size
 - Data redundancy
 - Data consistency
 - Application-side joins
 - Referential integrity constraints
 - Transactions
 - Data aggregation

DESIGNING THE MODEL

Finalizing the physical model



CREATING TABLES

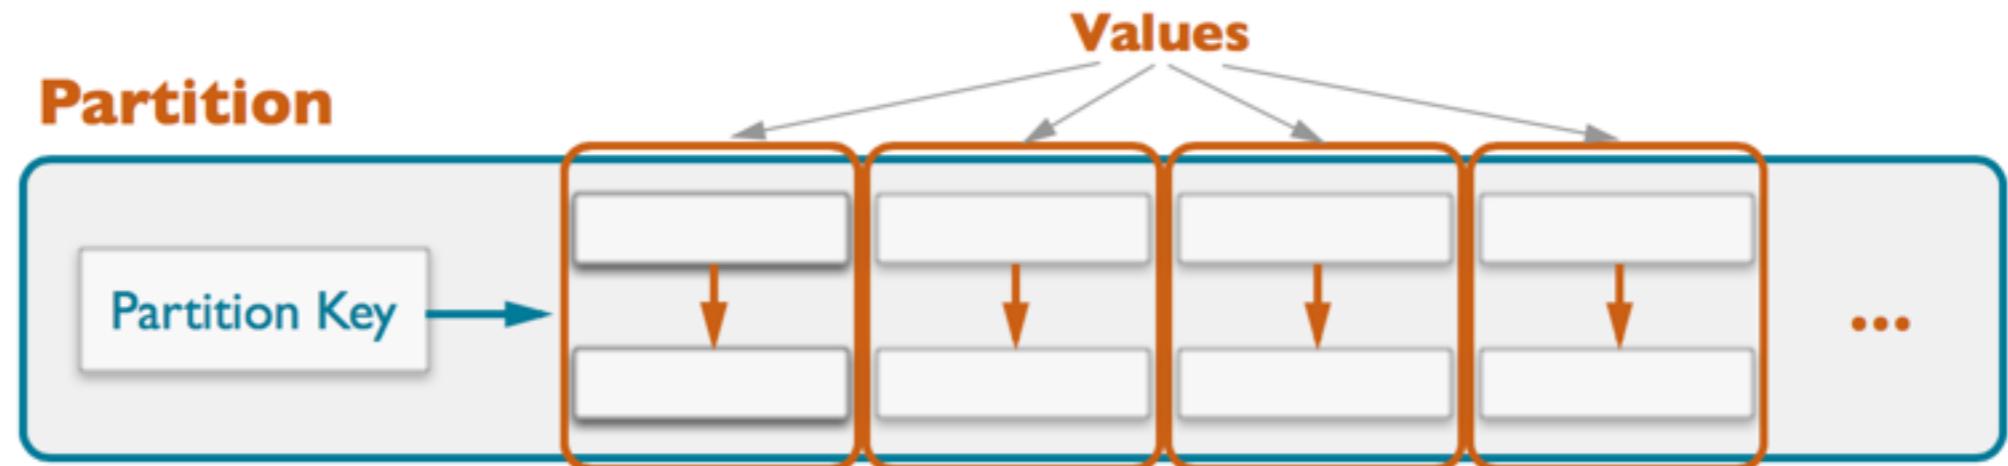
Writing the CQL statements

```
CREATE TABLE comments_by_user (
    user_id UUID,
    posted_timestamp TIMESTAMP,
    video_id TIMEUUID,
    comment TEXT,
    title TEXT,
    type TEXT,
    tags SET<TEXT>,
    preview_thumbnails MAP<INT, BLOB>,
    PRIMARY KEY ((user_id), posted_timestamp, video_id)
) WITH CLUSTERING ORDER BY (posted_timestamp DESC, video_id ASC);
```

PARTITION SIZE

Implementation limit on the size of a partition in Cassandra

- Number of values in a partition – 2 billion



- Size of the partition on disk – must be able to fit entirely on a node

COMPARING CASSANDRA (C*) AND RDBMS

- with RDBMS, a normalized data model is created without considering the exact queries
 - SQL can return almost anything though Joins
- With C*, the data model is designed for specific queries schema is adjusted as new queries introduced
- C*: NO joins, relationships, or foreign keys
 - a separate table is leveraged per query
 - data required by multiple tables is denormalized across those tables

CQL

stands for

Cassandra Query Language

```
cqlsh> help
```

Documented shell commands:

```
=====
```

CAPTURE	CLS	COPY	DESCRIBE	EXPAND	LOGIN	SERIAL	SOURCE	UNICODE
CLEAR	CONSISTENCY	DESC	EXIT	HELP	PAGING	SHOW	TRACING	

CQL help topics:

```
=====
```

AGGREGATES	CREATE_KEYSPACE	DROP_TRIGGER	TEXT
ALTER_KEYSPACE	CREATE_MATERIALIZED_VIEW	DROP_TYPE	TIME
ALTER_MATERIALIZED_VIEW	CREATE_ROLE	DROP_USER	TIMESTAMP
ALTER_TABLE	CREATE_TABLE	FUNCTIONS	TRUNCATE
ALTER_TYPE	CREATE_TRIGGER	GRANT	TYPES
ALTER_USER	CREATE_TYPE	INSERT	UPDATE
APPLY	CREATE_USER	INSERT_JSON	USE
ASCII	DATE	INT	UUID
BATCH	DELETE	JSON	
BEGIN	DROP_AGGREGATE	KEYWORDS	
BLOB	DROP_COLUMNFAMILY	LIST_PERMISSIONS	
BOOLEAN	DROP_FUNCTION	LIST_ROLES	
COUNTER	DROP_INDEX	LIST_USERS	
CREATE_AGGREGATE	DROP_KEYSPACE	PERMISSIONS	
CREATE_COLUMNFAMILY	DROP_MATERIALIZED_VIEW	REVOKE	
CREATE_FUNCTION	DROP_ROLE	SELECT	
CREATE_INDEX	DROP_TABLE	SELECT_JSON	

```
cqlsh> □
```

BASIC DATA TYPES IN CASSANDRA

- Numeric
 - bigint, decimal, double, float, int, varint
- String
 - ascii, text, varchar
- Date
 - timestamp, timeuuid
- Other
 - boolean, uuid, inet, blob

NAMING YOUR KEYSPACES, TABLES AND COLUMNS

- No hyphens: 2016-sales
- No spaces: 2015 sales
- Double quotes required for initial digits: “2016sales”
- Mixed case is lowercased unless surrounded in double quotes: “firstName”

CQL

- creating a keyspace - namespace of tables
 - CREATE KEYSPACE demo WITH replication = {'class': 'SimpleStrategy', replication_factor': 3};
- To use namespace:
 - USE demo;

CREATING TABLES

- creating tables:

```
CREATE TABLE users(
```

```
    email varchar,
```

```
    bio varchar,
```

```
    birthday timestamp,
```

```
    active boolean,
```

```
    PRIMARY KEY (email));
```

```
CREATE TABLE tweets(
```

```
    email varchar,
```

```
    time_posted timestamp,
```

```
    tweet varchar,
```

```
    PRIMARY KEY (email, time_posted));
```

QUERYING TABLES

- SELECT expression reads one or more records from Cassandra column family and returns a result-set of rows
- SELECT * FROM users;
- SELECT email FROM users WHERE active = true;
- SELECT email FROM users WHERE active = true LIMIT 100;

INSERTING/UPDATING DATA

- `INSERT INTO users (email, bio, birthday, active)
VALUES ('john.crost@cnn.com', 'Coach',
646464676600, true);`
- `UPDATE users SET email='john.crost@cnn.com',
bio='Coach', birthday=646464676600,
active=true)`

UPsert

- Under the hood, INSERT and UPDATE are same
- INSERT = UPDATE
- BOTH OPERATIONS REQUIRE PRIMARY KEY

WHEN WAS DATA WRITTEN?

- SELECT email, WRITETIME(bio) FROM users;
 - Unix time (e.g. 1430825689)

DELETING DATA

- Deleting a row
 - `DELETE FROM customers WHERE id = '2829';`
- Deleting a column
 - `DELETE email FROM customers WHERE id = '2829';`
 - `UPDATE customers SET email = null WHERE id = '2829';`
 - `INSERT INTO customers (id, email) VALUES ('2829', null);`

DEMO

- Cassandra query language
- <http://www.planetcassandra.org/try-cassandra/>

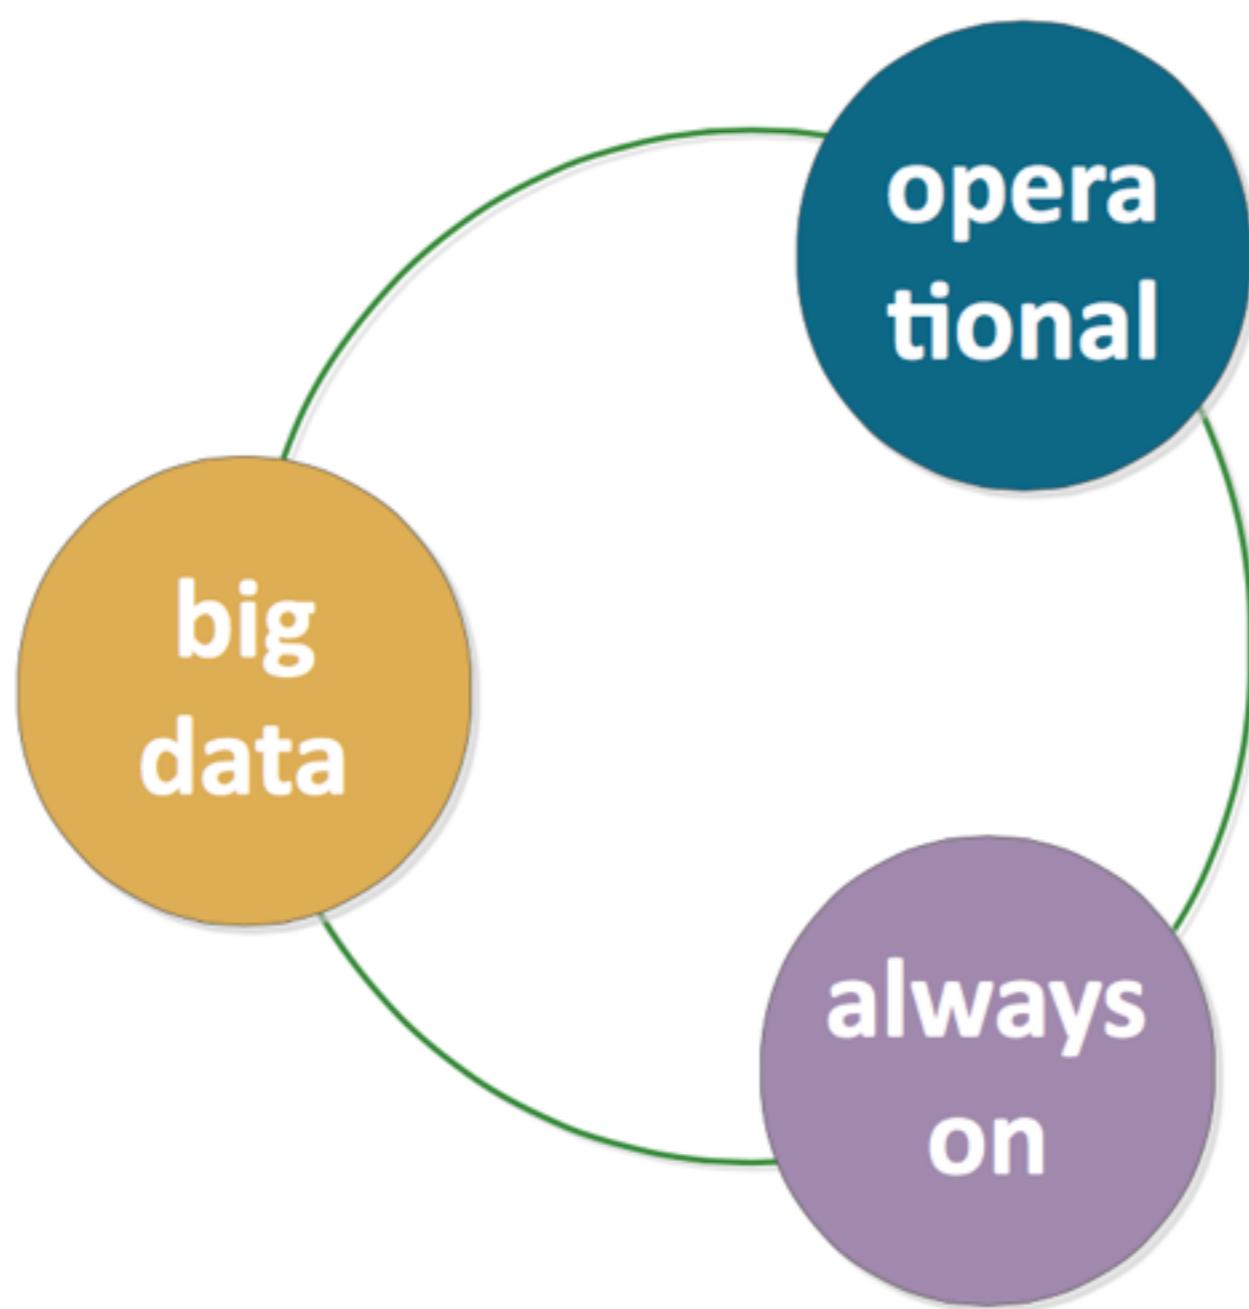
CASSANDRA ADVANTAGES

- Perfect for time-series data
- High performance
- Decentralization
- Nearly linear scalability
- Replication support
- No single points of failure
- MapReduce support

CASSANDRA WEAKNESSES

- No referential integrity
- No concept of JOIN
- Querying options for retrieving data are limited
- Sorting data is a design decision
- No GROUP BY
- No support for atomic operations
- If operation fails, changes can still occur
- First think about queries, then about data model

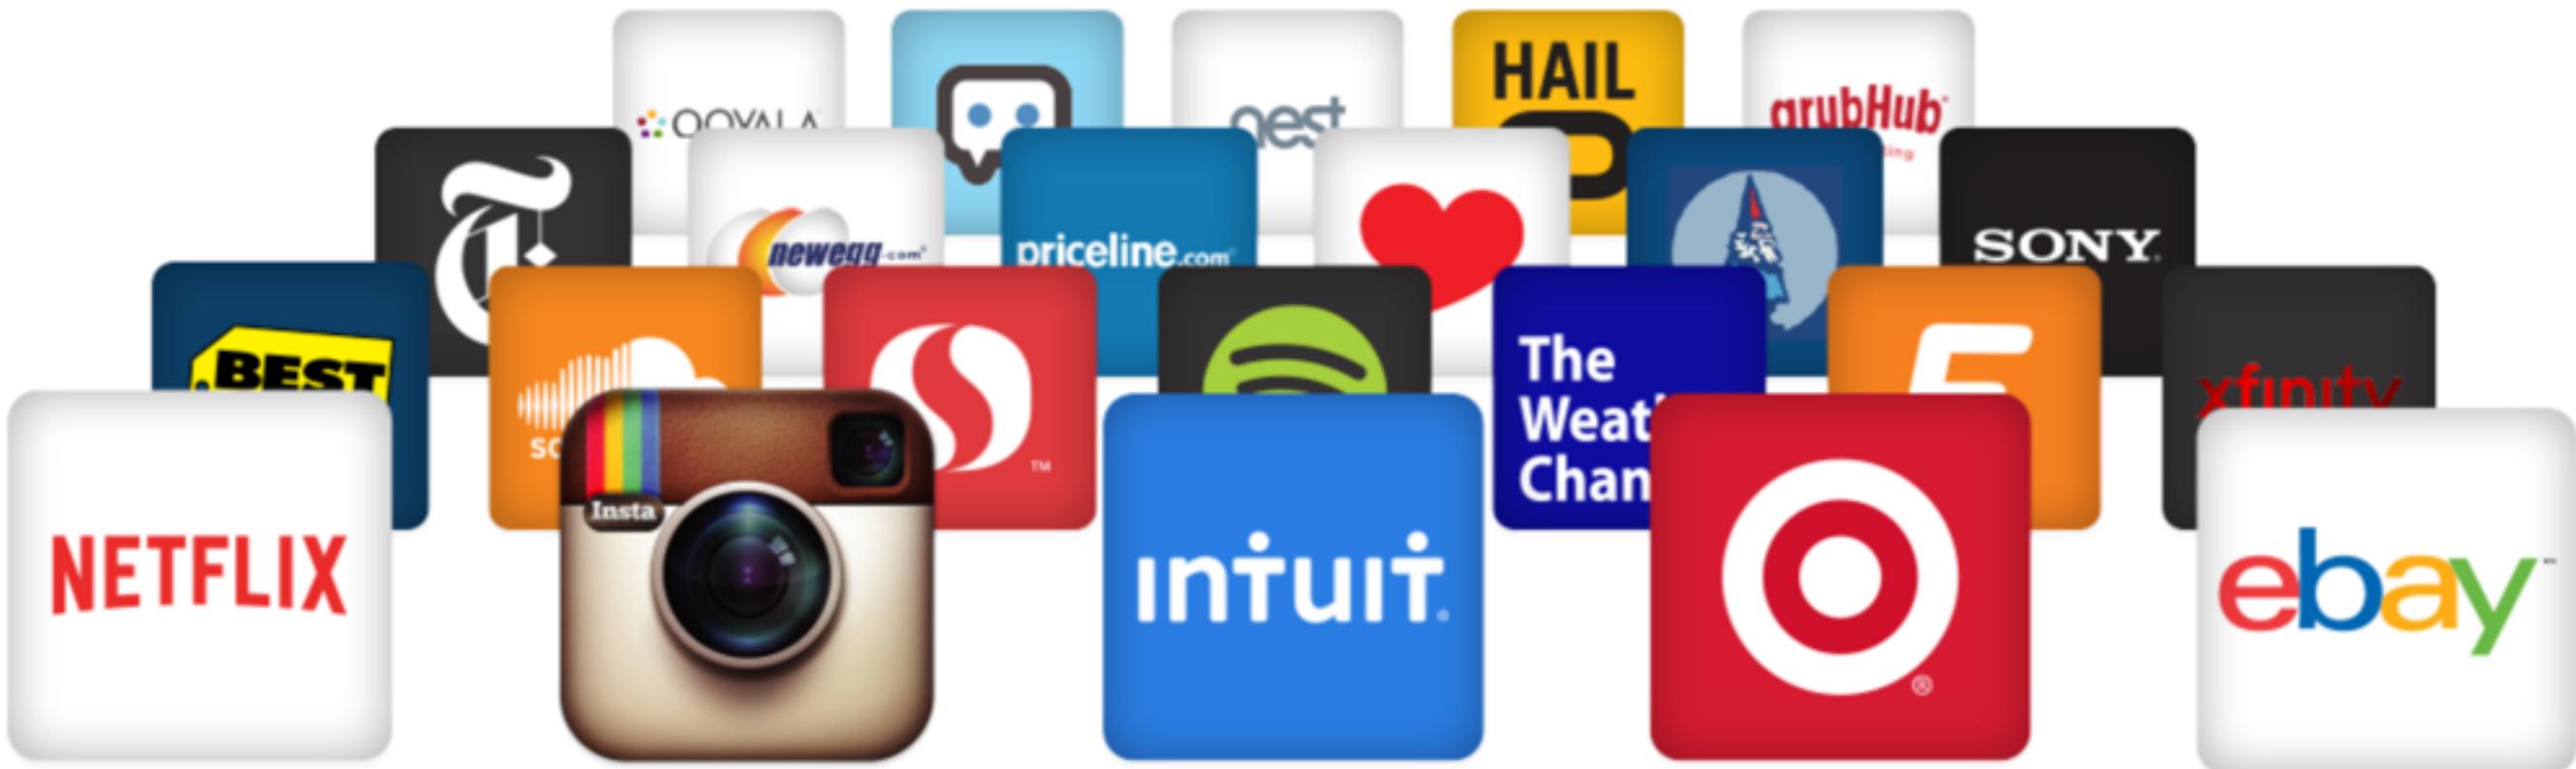
USECASES



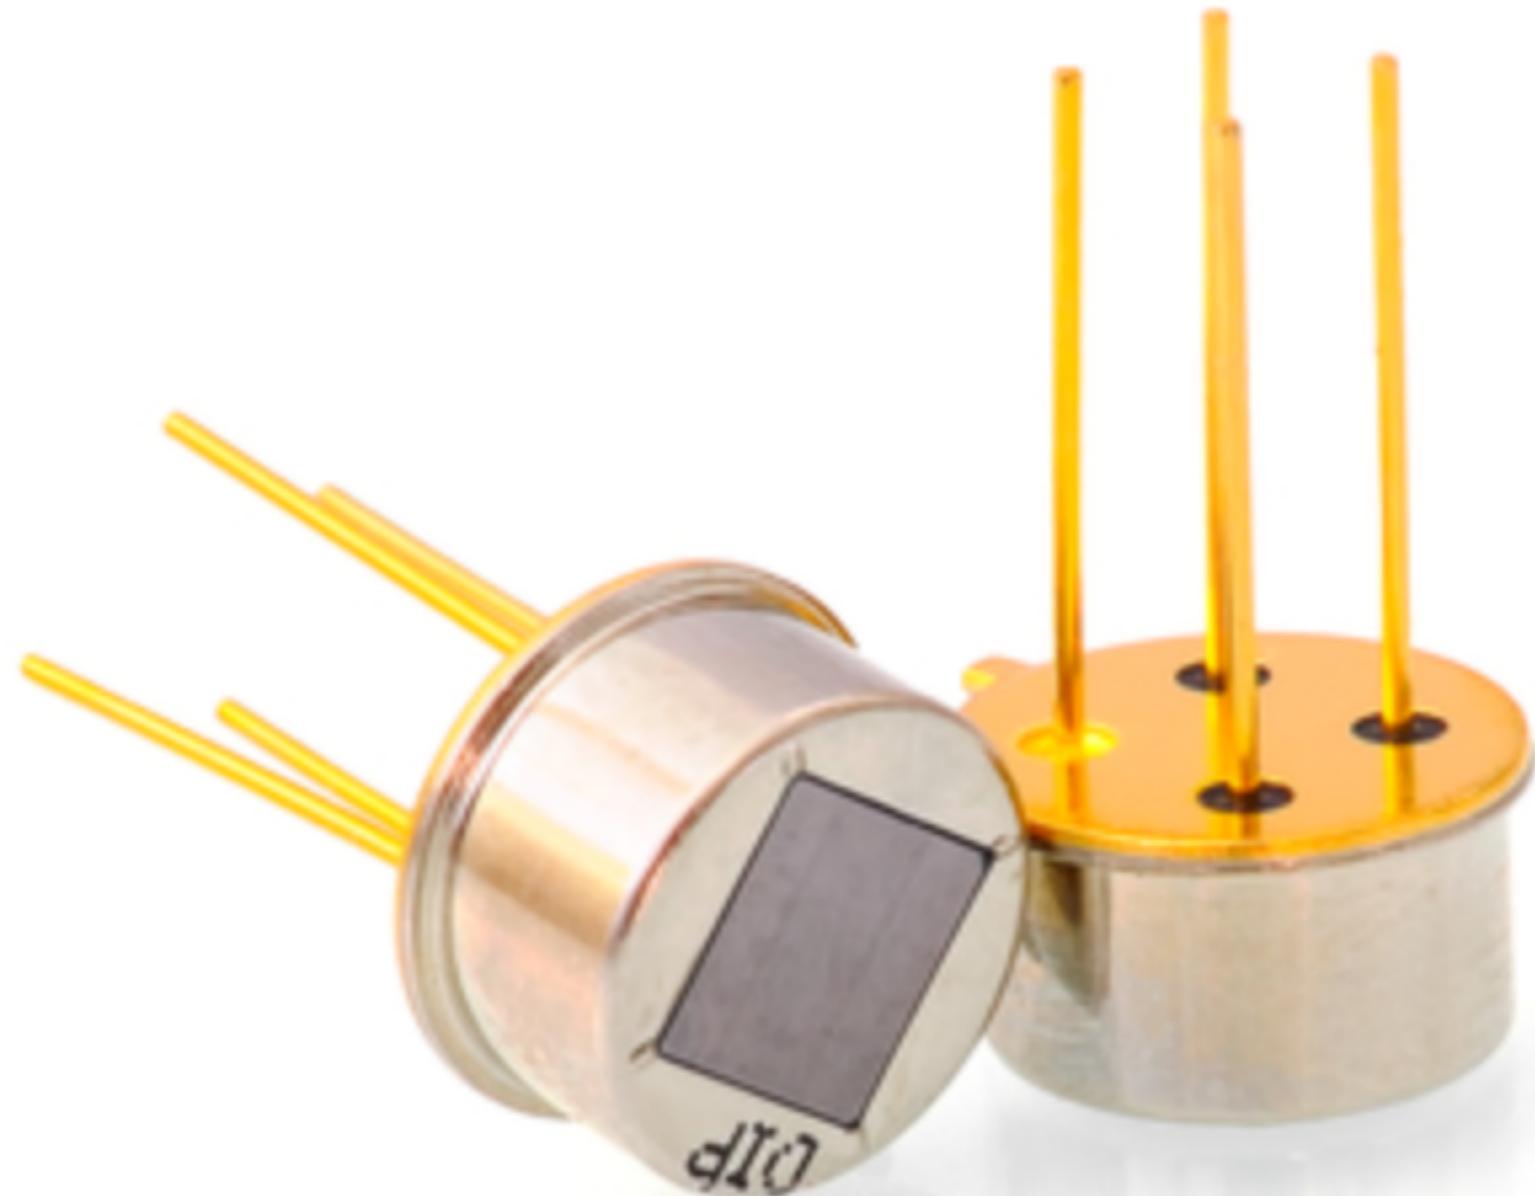
MAIN FUNCTIONAL USE CASES

Use Case	Example
Product catalogs and playlists	Collection of items
Internet of things	Sensor data
Messaging	Emails, instant messages, alerts, comments
Recommendation and personalization	Trend information
Fraud detection	Data patterns

WHAT COMPANIES?



SENSOR DATA USECASE



REFERENCES

- http://docs.datastax.com/en/landing_page/doc/landing_page/current.html
- <http://spark.apache.org/>
- <http://cassandra.apache.org/>
- <http://www.planetcassandra.org/>
- <http://demo.gethue.com/home>